

```

1 import matplotlib.pyplot as plt
2 import tensorflow as tf
3 import numpy as np
4 from tensorflow.keras import datasets, layers, models
5

```

```

1
2 mnist = tf.keras.datasets.cifar10
3

```

```

1 print(x_train.shape)
2

```

```
→ (50000, 32, 32, 3)
```

```

1 (train_images, train_labels), (test_images, test_labels) = datasets.cifar10.load_data()
2 train_images, test_images = train_images / 255.0, test_images / 255.0
3 train_labels = tf.keras.utils.to_categorical(train_labels, 10)
4 test_labels = tf.keras.utils.to_categorical(test_labels, 10)
5

```

```

1 test_labels[0]

```

```
→ array([0., 0., 0., 1., 0., 0., 0., 0., 0.], dtype=float32)
```

```

1 from tensorflow.keras import models, layers
2 def build_lenet():
3     model = models.Sequential()
4     model.add(layers.Conv2D(6, (5, 5), activation='relu', input_shape=(32, 32, 3)))
5     model.add(layers.AveragePooling2D((2, 2)))
6     model.add(layers.Conv2D(16, (5, 5), activation='relu'))
7     model.add(layers.AveragePooling2D((2, 2)))
8     model.add(layers.Flatten())
9     model.add(layers.Dense(120, activation='relu'))
10    model.add(layers.Dense(84, activation='relu'))
11    model.add(layers.Dense(10, activation='softmax'))
12
13    return model

```

```

1
2 # Compile the model
3 model = build_lenet()
4 model.compile(optimizer='adam',
5               loss='categorical_crossentropy',
6               metrics=['accuracy'])
7 model.summary()

```

```
→ Model: "sequential_2"
```

Layer (type)	Output Shape	Param #
=====		
conv2d_4 (Conv2D)	(None, 28, 28, 6)	456
average_pooling2d_4 (AveragePooling2D)	(None, 14, 14, 6)	0
conv2d_5 (Conv2D)	(None, 10, 10, 16)	2416
average_pooling2d_5 (AveragePooling2D)	(None, 5, 5, 16)	0
flatten_2 (Flatten)	(None, 400)	0
dense_6 (Dense)	(None, 120)	48120
dense_7 (Dense)	(None, 84)	10164
dense_8 (Dense)	(None, 10)	850
=====		
Total params: 62006 (242.21 KB)		
Trainable params: 62006 (242.21 KB)		
Non-trainable params: 0 (0.00 Byte)		

```

1 history = model.fit(train_images, train_labels, epochs=20,
2                       validation_data=(test_images, test_labels))
3

```

```

Epoch 1/20
1563/1563 [=====] - 7s 5ms/step - loss: 1.3661 - accuracy: 0.5085 - val_loss: 1.3219 - val_accuracy: 0.5295
Epoch 2/20
1563/1563 [=====] - 8s 5ms/step - loss: 1.2638 - accuracy: 0.5446 - val_loss: 1.2660 - val_accuracy: 0.5501
Epoch 3/20
1563/1563 [=====] - 7s 5ms/step - loss: 1.1945 - accuracy: 0.5719 - val_loss: 1.2253 - val_accuracy: 0.5653
Epoch 4/20
1563/1563 [=====] - 8s 5ms/step - loss: 1.1383 - accuracy: 0.5950 - val_loss: 1.1864 - val_accuracy: 0.5784
Epoch 5/20
1563/1563 [=====] - 9s 6ms/step - loss: 1.0904 - accuracy: 0.6122 - val_loss: 1.1525 - val_accuracy: 0.5918
Epoch 6/20
1563/1563 [=====] - 7s 5ms/step - loss: 1.0490 - accuracy: 0.6263 - val_loss: 1.2124 - val_accuracy: 0.5718
Epoch 7/20
1563/1563 [=====] - 9s 5ms/step - loss: 1.0152 - accuracy: 0.6398 - val_loss: 1.1433 - val_accuracy: 0.6020
Epoch 8/20
1563/1563 [=====] - 8s 5ms/step - loss: 0.9810 - accuracy: 0.6521 - val_loss: 1.1361 - val_accuracy: 0.6062
Epoch 9/20
1563/1563 [=====] - 9s 5ms/step - loss: 0.9425 - accuracy: 0.6659 - val_loss: 1.1388 - val_accuracy: 0.6031
Epoch 10/20
1563/1563 [=====] - 9s 6ms/step - loss: 0.9142 - accuracy: 0.6731 - val_loss: 1.1581 - val_accuracy: 0.6039
Epoch 11/20
1563/1563 [=====] - 8s 5ms/step - loss: 0.8835 - accuracy: 0.6879 - val_loss: 1.1564 - val_accuracy: 0.6048
Epoch 12/20
1563/1563 [=====] - 8s 5ms/step - loss: 0.8603 - accuracy: 0.6945 - val_loss: 1.1776 - val_accuracy: 0.5976
Epoch 13/20
1563/1563 [=====] - 9s 6ms/step - loss: 0.8267 - accuracy: 0.7044 - val_loss: 1.1961 - val_accuracy: 0.5939
Epoch 14/20
1563/1563 [=====] - 8s 5ms/step - loss: 0.8078 - accuracy: 0.7143 - val_loss: 1.1769 - val_accuracy: 0.6078
Epoch 15/20
1563/1563 [=====] - 9s 6ms/step - loss: 0.7854 - accuracy: 0.7199 - val_loss: 1.2016 - val_accuracy: 0.6017
Epoch 16/20
1563/1563 [=====] - 8s 5ms/step - loss: 0.7603 - accuracy: 0.7284 - val_loss: 1.2341 - val_accuracy: 0.6005
Epoch 17/20
1563/1563 [=====] - 8s 5ms/step - loss: 0.7405 - accuracy: 0.7374 - val_loss: 1.2414 - val_accuracy: 0.5986
Epoch 18/20
1563/1563 [=====] - 9s 6ms/step - loss: 0.7170 - accuracy: 0.7432 - val_loss: 1.2346 - val_accuracy: 0.5995
Epoch 19/20
1563/1563 [=====] - 9s 6ms/step - loss: 0.6998 - accuracy: 0.7507 - val_loss: 1.2699 - val_accuracy: 0.5986
Epoch 20/20
1563/1563 [=====] - 8s 5ms/step - loss: 0.6771 - accuracy: 0.7594 - val_loss: 1.3213 - val_accuracy: 0.5957

```

```

1 test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
2 print(f"Test accuracy: {test_acc}")

```

```

313/313 - 1s - loss: 1.3213 - accuracy: 0.5957 - 1s/epoch - 4ms/step
Test accuracy: 0.5957000255584717


```

```


1 predictions = model.predict(test_images[:5])
2
3 # Display the first 5 test images and their predicted labels
4 for i in range(5):
5     plt.imshow(test_images[i])
6     plt.show()
7     print(f"Predicted label: {np.argmax(predictions[i])}")
8

```


↩ 1/1 [=====] - 0s 158ms/step



Predicted label: 3




Predicted label: 8




Predicted label: 8





Predicted label: 0



Predicted label: 6

```

1 import tensorflow as tf
2 from tensorflow.keras import datasets, layers, models
3 import numpy as np
4 import matplotlib.pyplot as plt
5
6 # Load and preprocess dataset
7 (train_images, train_labels), (test_images, test_labels) = datasets.cifar10.load_data()
8 train_images, test_images = train_images / 255.0, test_images / 255.0
9 train_labels = tf.keras.utils.to_categorical(train_labels, 10)
10 test_labels = tf.keras.utils.to_categorical(test_labels, 10)
11
12 # Define the LeNet model
13 def create_lenet_model():
14     model = models.Sequential()
15     model.add(layers.Conv2D(6, (5, 5), activation='relu', input_shape=(32, 32, 3)))
16     model.add(layers.AveragePooling2D((2, 2)))
17     model.add(layers.Conv2D(16, (5, 5), activation='relu'))
18     model.add(layers.AveragePooling2D((2, 2)))
19     model.add(layers.Flatten())
20     model.add(layers.Dense(120, activation='relu'))
21     model.add(layers.Dense(84, activation='relu'))
22     model.add(layers.Dense(10, activation='softmax'))
23     return model
24
25 # Compile the model
26 model = create_lenet_model()
27 model.compile(optimizer='adam',
28               loss='categorical_crossentropy',
29               metrics=['accuracy'])
30
31 # Train the model
32 history = model.fit(train_images, train_labels, epochs=10,
33                     validation_data=(test_images, test_labels))
34
35 # Evaluate the model
36 test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
37 print(f"Test accuracy: {test_acc}")
38
39 # Predict the first 5 images in the test set
40 predictions = model.predict(test_images[:5])
41
42 # Display the first 5 test images and their predicted labels
43 for i in range(5):
44     plt.imshow(test_images[i])
45     plt.show()
46     print(f"Predicted label: {np.argmax(predictions[i])}")
47


```

```


1
2 # Train the model
3 history = model.fit(train_images, train_labels, epochs=10,
4                     validation_data=(test_images, test_labels))
5
6 # Evaluate the model
7 test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
8 print(f"Test accuracy: {test_acc}")
9
10 # Predict the first 5 images in the test set
11 predictions = model.predict(test_images[:5])
12
13 # Display the first 5 test images and their predicted labels
14 for i in range(5):
15     plt.imshow(test_images[i])
16     plt.show()
17     print(f"Predicted label: {np.argmax(predictions[i])}")
18

```

Epoch 1/10
1563/1563 [=====] - 41s 25ms/step - loss: 1.6807 - accuracy: 0.
Epoch 2/10
1563/1563 [=====] - 36s 23ms/step - loss: 1.4188 - accuracy: 0.
Epoch 3/10
1563/1563 [=====] - 38s 25ms/step - loss: 1.3149 - accuracy: 0.
Epoch 4/10
1563/1563 [=====] - 40s 26ms/step - loss: 1.2460 - accuracy: 0.
Epoch 5/10
1563/1563 [=====] - 43s 28ms/step - loss: 1.1851 - accuracy: 0.
Epoch 6/10
1563/1563 [=====] - 40s 26ms/step - loss: 1.1360 - accuracy: 0.
Epoch 7/10
1563/1563 [=====] - 45s 29ms/step - loss: 1.0897 - accuracy: 0.
Epoch 8/10
1563/1563 [=====] - 39s 25ms/step - loss: 1.0458 - accuracy: 0.
Epoch 9/10
1563/1563 [=====] - 38s 24ms/step - loss: 1.0073 - accuracy: 0.
Epoch 10/10
1563/1563 [=====] - 36s 23ms/step - loss: 0.9732 - accuracy: 0.
313/313 - 3s - loss: 1.1774 - accuracy: 0.5928 - 3s/epoch - 10ms/step
Test accuracy: 0.5928000211715698
1/1 [=====] - 0s 113ms/step




Predicted label: 3




Predicted label: 8






Predicted label: 8



Predicted label: 0



Predicted label: 3

```

1 def create_vgg16_model():
2     vgg16_base = tf.keras.applications.VGG16(weights='imagenet', include_top=False, input_shape=(32, 32, 3))
3     model = models.Sequential()
4     model.add(vgg16_base)
5     model.add(layers.Flatten())
6     model.add(layers.Dense(512, activation='relu'))
7     model.add(layers.Dense(10, activation='softmax'))
8
9     # Freeze the convolutional base
10    vgg16_base.trainable = False
11
12    return model
13
14 vgg16_model = create_vgg16_model()
15 vgg16_model.compile(optimizer='adam',
16                     loss='categorical_crossentropy',
17                     metrics=['accuracy'])
18
19
20

```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.58889256/58889256 [=====] - 4s 0us/step

```
1 vgg16_model.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 1, 1, 512)	14714688
flatten_3 (Flatten)	(None, 512)	0
dense_9 (Dense)	(None, 512)	262656
dense_10 (Dense)	(None, 10)	5130
Total params: 14982474 (57.15 MB)		
Trainable params: 267786 (1.02 MB)		
Non-trainable params: 14714688 (56.13 MB)		

```

1 vgg16_history = vgg16_model.fit(train_images, train_labels, epochs=20,
2                                validation_data=(test_images, test_labels))

```

```

Epoch 1/20
1563/1563 [=====] - 17s 11ms/step - loss: 0.8470 - accuracy: 0.7009 - val_loss: 1.1105 - val_accuracy: 0.6215
Epoch 2/20
1563/1563 [=====] - 17s 11ms/step - loss: 0.8074 - accuracy: 0.7148 - val_loss: 1.1205 - val_accuracy: 0.6171
Epoch 3/20
1563/1563 [=====] - 17s 11ms/step - loss: 0.7709 - accuracy: 0.7295 - val_loss: 1.1459 - val_accuracy: 0.6173
Epoch 4/20
1563/1563 [=====] - 17s 11ms/step - loss: 0.7349 - accuracy: 0.7429 - val_loss: 1.1610 - val_accuracy: 0.6181
Epoch 5/20
1563/1563 [=====] - 17s 11ms/step - loss: 0.7006 - accuracy: 0.7552 - val_loss: 1.1781 - val_accuracy: 0.6237
Epoch 6/20
1563/1563 [=====] - 17s 11ms/step - loss: 0.6655 - accuracy: 0.7659 - val_loss: 1.1836 - val_accuracy: 0.6165
Epoch 7/20
1563/1563 [=====] - 17s 11ms/step - loss: 0.6315 - accuracy: 0.7775 - val_loss: 1.2129 - val_accuracy: 0.6191
Epoch 8/20
1563/1563 [=====] - 17s 11ms/step - loss: 0.6006 - accuracy: 0.7900 - val_loss: 1.2379 - val_accuracy: 0.6198
Epoch 9/20
1563/1563 [=====] - 17s 11ms/step - loss: 0.5701 - accuracy: 0.7995 - val_loss: 1.2941 - val_accuracy: 0.6079
Epoch 10/20
1563/1563 [=====] - 16s 11ms/step - loss: 0.5411 - accuracy: 0.8094 - val_loss: 1.2959 - val_accuracy: 0.6151
Epoch 11/20
1563/1563 [=====] - 17s 11ms/step - loss: 0.5131 - accuracy: 0.8201 - val_loss: 1.3391 - val_accuracy: 0.6194
Epoch 12/20
1563/1563 [=====] - 17s 11ms/step - loss: 0.4849 - accuracy: 0.8309 - val_loss: 1.3618 - val_accuracy: 0.6101
Epoch 13/20
1563/1563 [=====] - 17s 11ms/step - loss: 0.4601 - accuracy: 0.8395 - val_loss: 1.3743 - val_accuracy: 0.6154
Epoch 14/20
1563/1563 [=====] - 17s 11ms/step - loss: 0.4355 - accuracy: 0.8494 - val_loss: 1.4235 - val_accuracy: 0.6069
Epoch 15/20
1563/1563 [=====] - 17s 11ms/step - loss: 0.4101 - accuracy: 0.8571 - val_loss: 1.4775 - val_accuracy: 0.6078
Epoch 16/20
1563/1563 [=====] - 17s 11ms/step - loss: 0.3896 - accuracy: 0.8664 - val_loss: 1.5127 - val_accuracy: 0.6114

```



```

Epoch 17/20
1563/1563 [=====] - 17s 11ms/step - loss: 0.3687 - accuracy: 0.8724 - val_loss: 1.5741 - val_accuracy: 0.6076
Epoch 18/20
1563/1563 [=====] - 16s 11ms/step - loss: 0.3518 - accuracy: 0.8776 - val_loss: 1.6029 - val_accuracy: 0.6035
Epoch 19/20
1563/1563 [=====] - 17s 11ms/step - loss: 0.3338 - accuracy: 0.8869 - val_loss: 1.6592 - val_accuracy: 0.5965
Epoch 20/20
1563/1563 [=====] - 17s 11ms/step - loss: 0.3155 - accuracy: 0.8915 - val_loss: 1.6957 - val_accuracy: 0.6075

```

```

1 test_loss, vgg16_model = model.evaluate(test_images, test_labels, verbose=2)
2 print(f"Test accuracy: {test_acc}")
3

```

```



➡ 313/313 - 1s - loss: 1.3213 - accuracy: 0.5957 - 704ms/epoch - 2ms/step
   Test accuracy: 0.5957000255584717

```

```

1 import matplotlib.pyplot as plt
2
3 # Function to plot training history
4 def plot_history(history, model_name):
5     plt.plot(history.history['accuracy'], label='Training Accuracy')
6     plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
7     plt.xlabel('Epoch')
8     plt.ylabel('Accuracy')
9     plt.ylim([0, 1])
10    plt.title(f'{model_name} Training and Validation Accuracy')
11    plt.legend(loc='lower right')
12    plt.show()
13
14 # Plot the training history for LeNet-5
15 plot_history(history, "LeNet-5")
16
17 # Plot the training history for VGG16
18 plot_history(vgg16_history, "VGG16")
19

```



```
1 predictions = vgg16_model.predict(test_images[:5])
2
3 # Display the first 5 test images and their predicted labels
4 for i in range(5):
5     plt.imshow(test_images[i])
6     plt.show()
7     print(f"Predicted label: {np.argmax(predictions[i])}")
```



```
-----
AttributeError                                Traceback (most recent call last)
<ipython-input-37-ca0b7faaae12> in <cell line: 1>()
----> 1 predictions = vgg16_model.predict(test_images[:5])
      2
      3 # Display the first 5 test images and their predicted labels
      4 for i in range(5):
      5     plt.imshow(test_images[i])

AttributeError: 'float' object has no attribute 'predict'
```

Next steps: [Explain error](#)

```
1 predictions = model.predict(test_images[:5])
2
3 # Display the first 5 test images and their predicted labels
4 for i in range(5):
5     plt.imshow(test_images[i])
```