

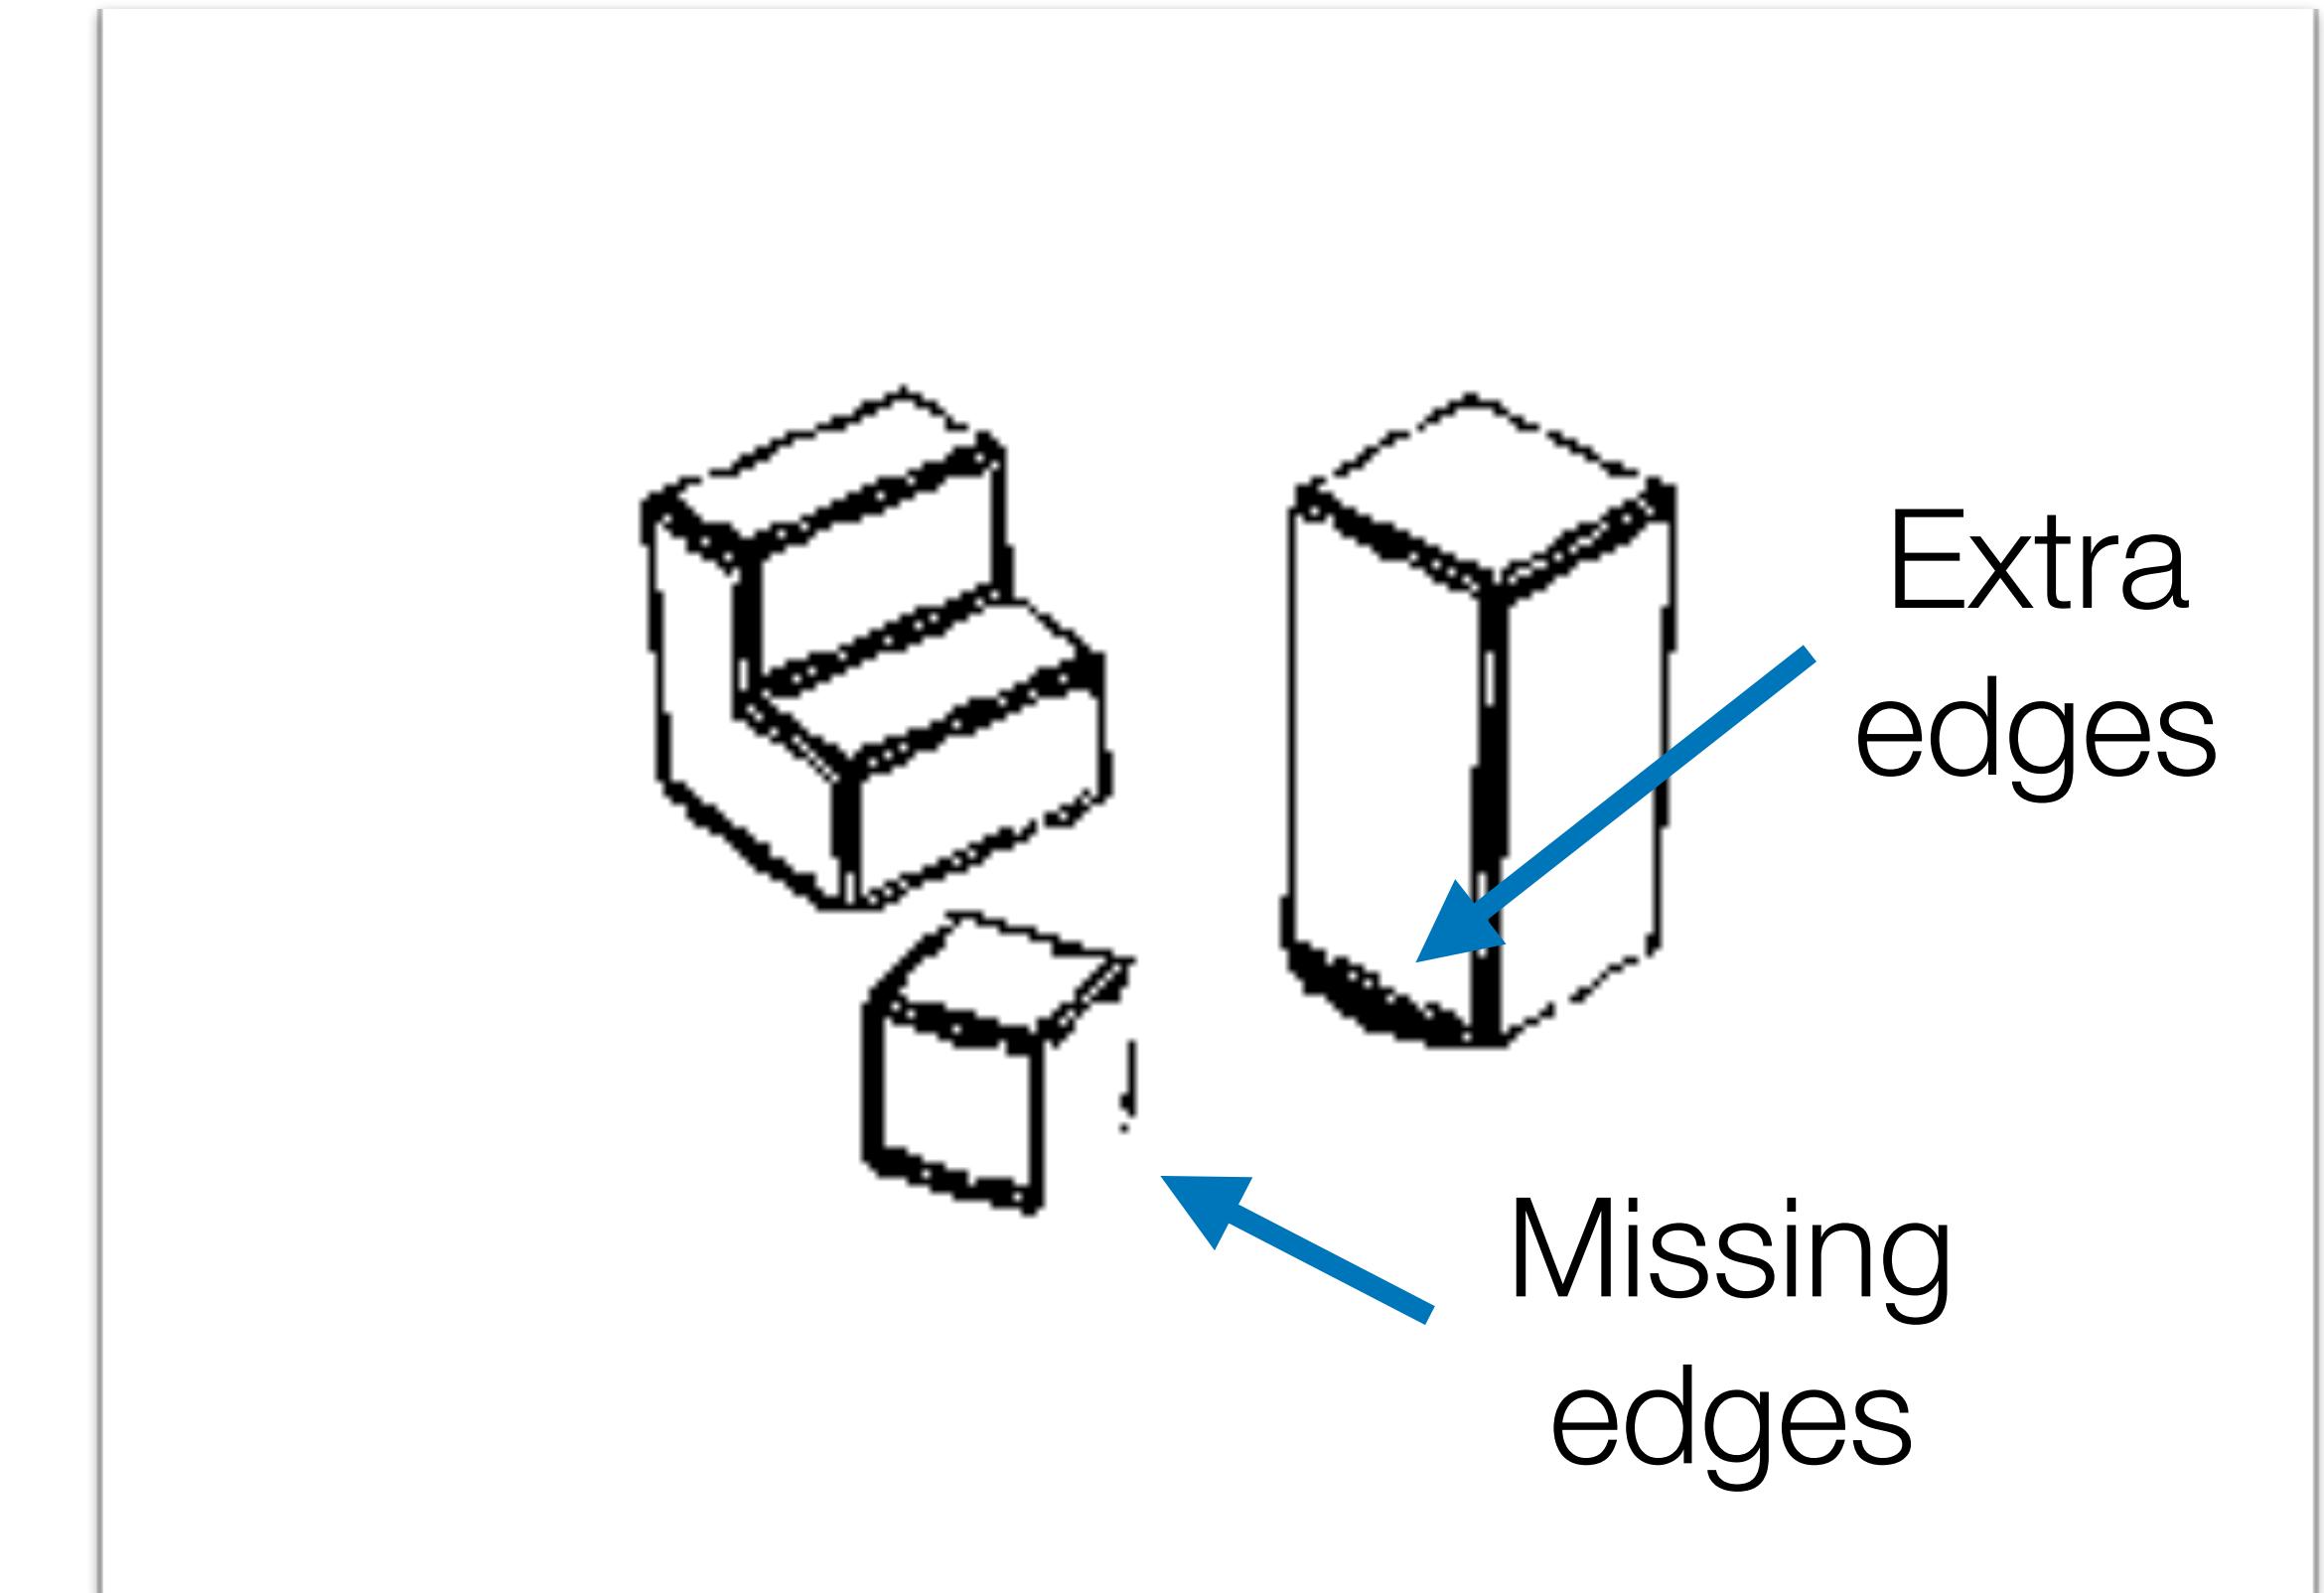
Lecture 2: Image filtering

- PS1 due next Tuesday
- Updated office hours next week, due to holiday. New times will be on Piazza.
- Questions?

Recall last week...



Input image



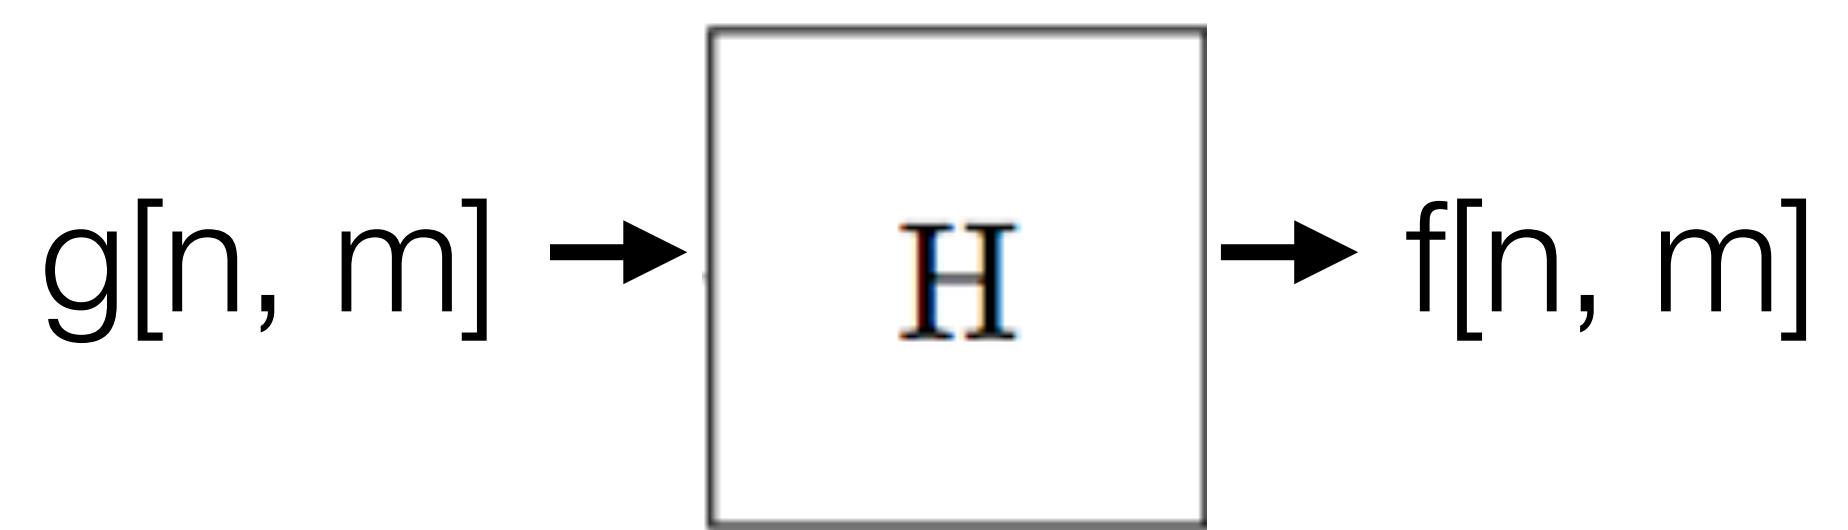
Edges

In this lecture



What *other* transformations can we do?

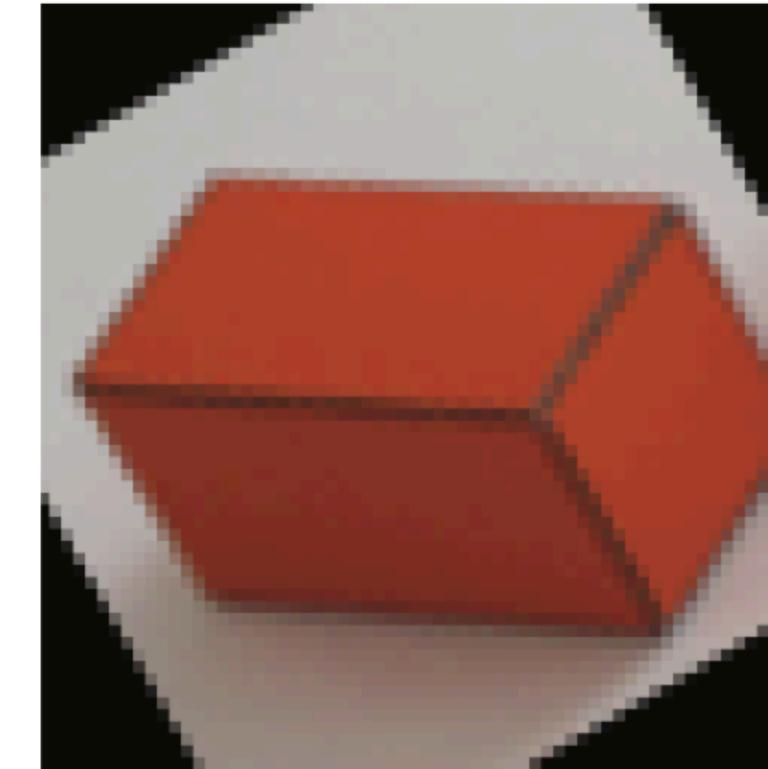
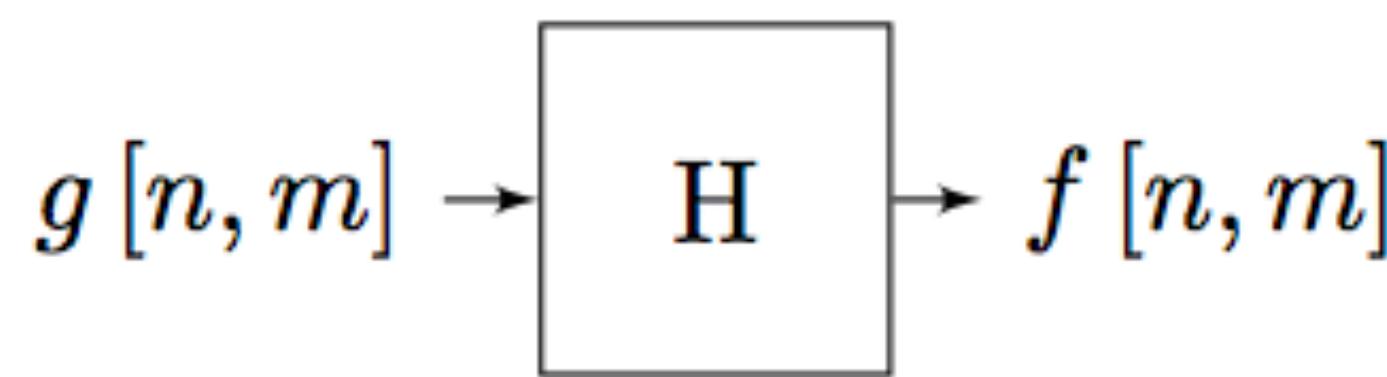
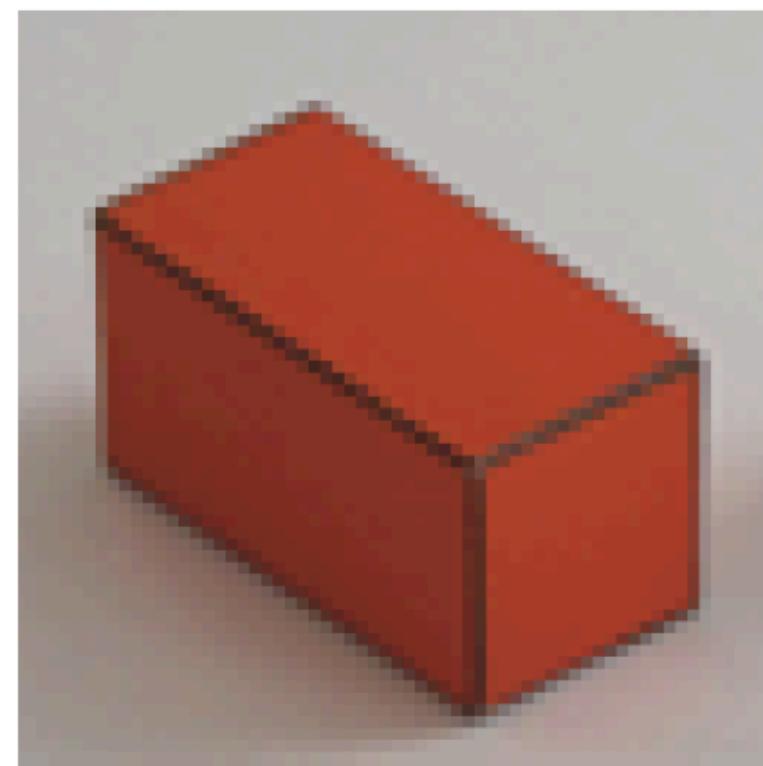
Filtering



Our goal: remove unwanted sources of variation, and keep the information relevant for whatever task we need to solve.



Linear filtering

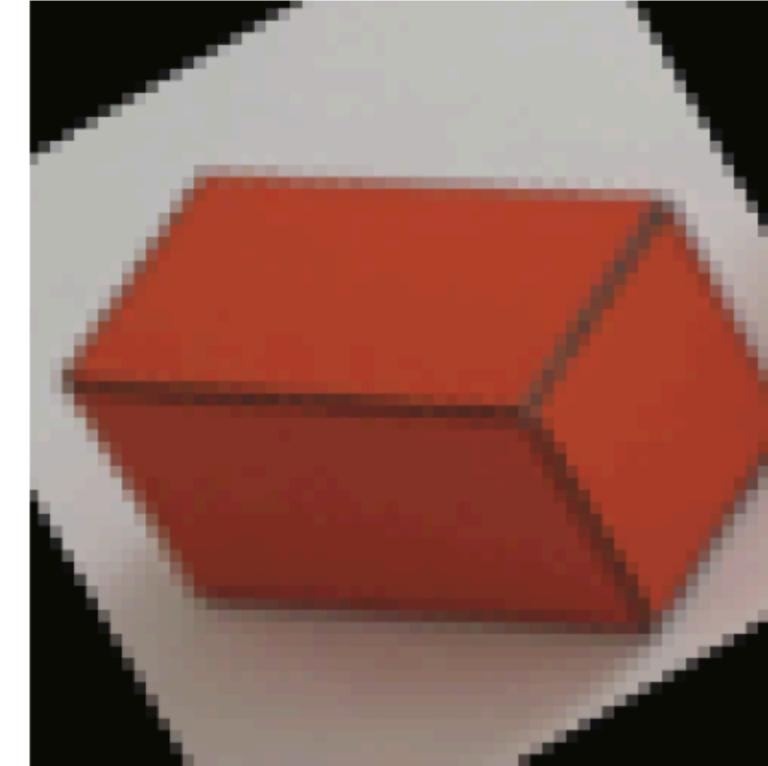
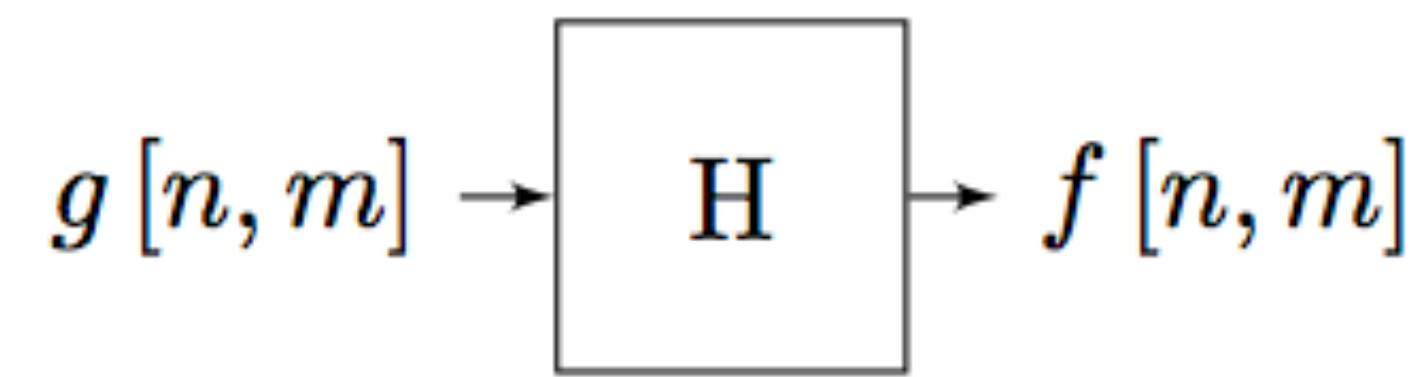
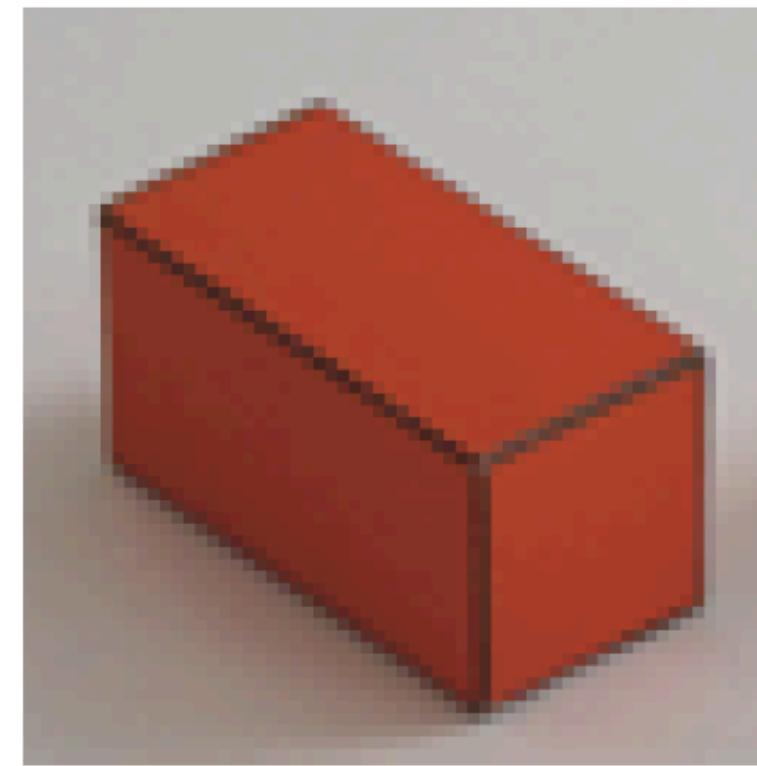


Very general! For a filter, H , to be linear, it has to satisfy:

$$H(a[m, n] + b[m, n]) = H(a[m, n]) + H(b[m, n])$$

$$H(Ca[m, n]) = CH(a[m, n])$$

Linear filtering



A linear filter in its most general form can be written as (for a 1D signal of length N):

$$f[n] = \sum_{k=0}^{N-1} h[n, k] g[k]$$

In matrix form:

$$\begin{bmatrix} f[0] \\ f[1] \\ \vdots \\ f[M-1] \end{bmatrix} = \begin{bmatrix} h[0, 0] & h[0, 1] & \dots & h[0, N-1] \\ h[1, 0] & h[1, 1] & \dots & h[1, N-1] \\ \vdots & \vdots & \vdots & \vdots \\ h[M-1, 0] & h[M-1, 1] & \dots & h[M-1, N-1] \end{bmatrix} \begin{bmatrix} g[0] \\ g[1] \\ \vdots \\ g[N-1] \end{bmatrix}$$

$$\begin{bmatrix} f[0] \\ f[1] \\ \vdots \\ f[M-1] \end{bmatrix} = \begin{bmatrix} h[0,0] & h[0,1] & \dots & h[0,N-1] \\ h[1,0] & h[1,1] & \dots & h[1,N-1] \\ \vdots & \vdots & \vdots & \vdots \\ h[M-1,0] & h[M-1,1] & \dots & h[M-1,N-1] \end{bmatrix} \begin{bmatrix} g[0] \\ g[1] \\ \vdots \\ g[N-1] \end{bmatrix}$$

Why handle each spatial position differently?

Want translation invariance!



Image denoising



Moving average

- Let's replace each pixel with a weighted average of its neighborhood
- The weights are called the **filter kernel**
- What are the weights for the average of a 3x3 neighborhood?

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

“box filter”

Moving average

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Filter kernel

0	0	0	0	0	0	0	0
0	90	90	90	90	90	0	0
0	90	90	90	90	90	0	0
0	90	90	90	90	90	0	0
0	90	0	90	90	0	0	0
0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0

Input

Output

Moving average

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

Filter kernel

0	0	0	0	0	0	0	0
0	90	90	90	90	90	0	0
0	90	90	90	90	90	0	0
0	90	90	90	90	90	0	0
0	90	0	90	90	0	0	0
0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0

Input

Output

Moving average

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

Filter kernel

0	0	0	0	0	0	0	0
0	90	90	90	90	90	0	0
0	90	90	90	90	90	0	0
0	90	90	90	90	90	0	0
0	90	0	90	90	90	0	0
0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0

Input

Output

Moving average

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

Filter kernel

0	0	0	0	0	0	0	0
0	90	90	90	90	90	0	0
0	90	90	90	90	90	0	0
0	90	90	90	90	90	0	0
0	90	0	90	90	90	0	0
0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0

Input

40							

Output

Moving average

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

Filter kernel

0	0	0	0	0	0	0
0	90	90	90	90	0	0
0	90	90	90	90	0	0
0	90	90	90	90	0	0
0	90	0	90	90	0	0
0	90	90	90	90	0	0
0	0	0	0	0	0	0

Input

		40	60			

Output

Moving average

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

Filter kernel

0	0	0	0	0	0	0
0	90	90	90	90	0	0
0	90	90	90	90	0	0
0	90	90	90	90	0	0
0	90	0	90	90	0	0
0	90	90	90	90	0	0
0	0	0	0	0	0	0

Input

	40	60				

Output

Moving average

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

Filter kernel

0	0	0	0	0	0	0
0	90	90	90	90	0	0
0	90	90	90	90	0	0
0	90	90	90	90	0	0
0	90	0	90	90	0	0
0	90	90	90	90	0	0
0	0	0	0	0	0	0

Input

	40	60	60	40	20	
	60	90	60	40	20	
	50	80	80	60	30	
	50	80	80	60	30	
	30	50	50	40	20	

Output

Moving average

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

Filter kernel

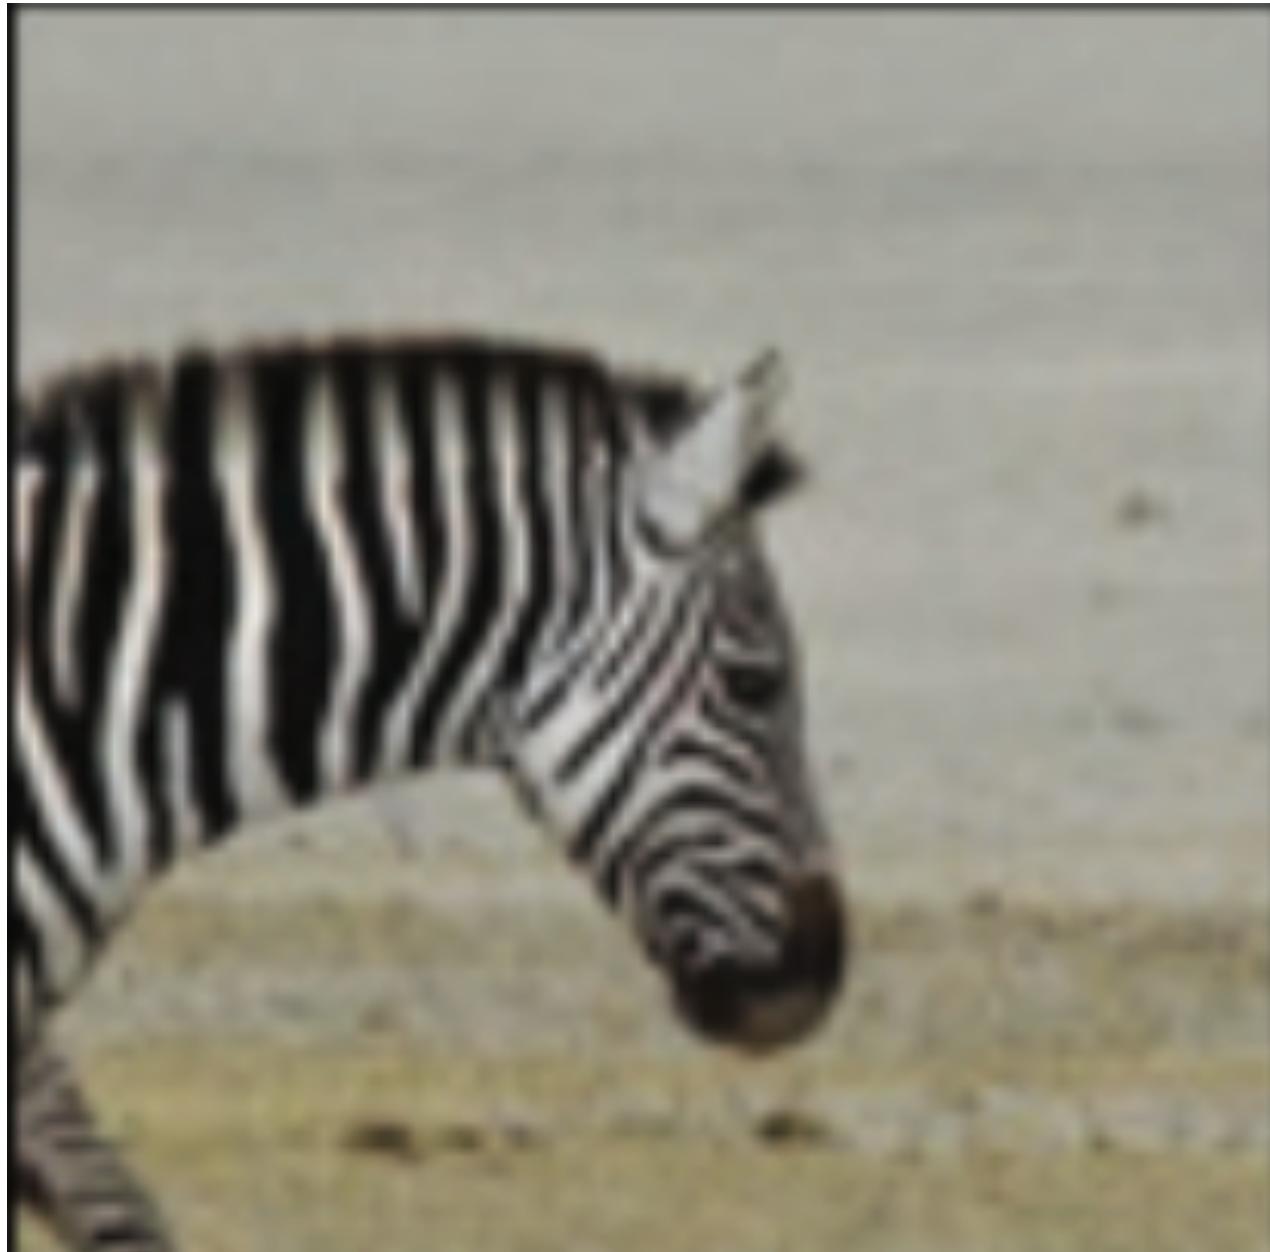
0	0	0	0	0	0	0
0	90	90	90	90	90	0
0	90	90	90	90	90	0
0	90	90	90	90	90	0
0	90	0	90	90	0	0
0	90	90	90	90	90	0
0	0	0	0	0	0	0

Input

	40	60	60	40	20	
	60	90	60	40	20	
	50	80	80	60	30	
?	50	80	80	60	30	
	30	50	50	40	20	

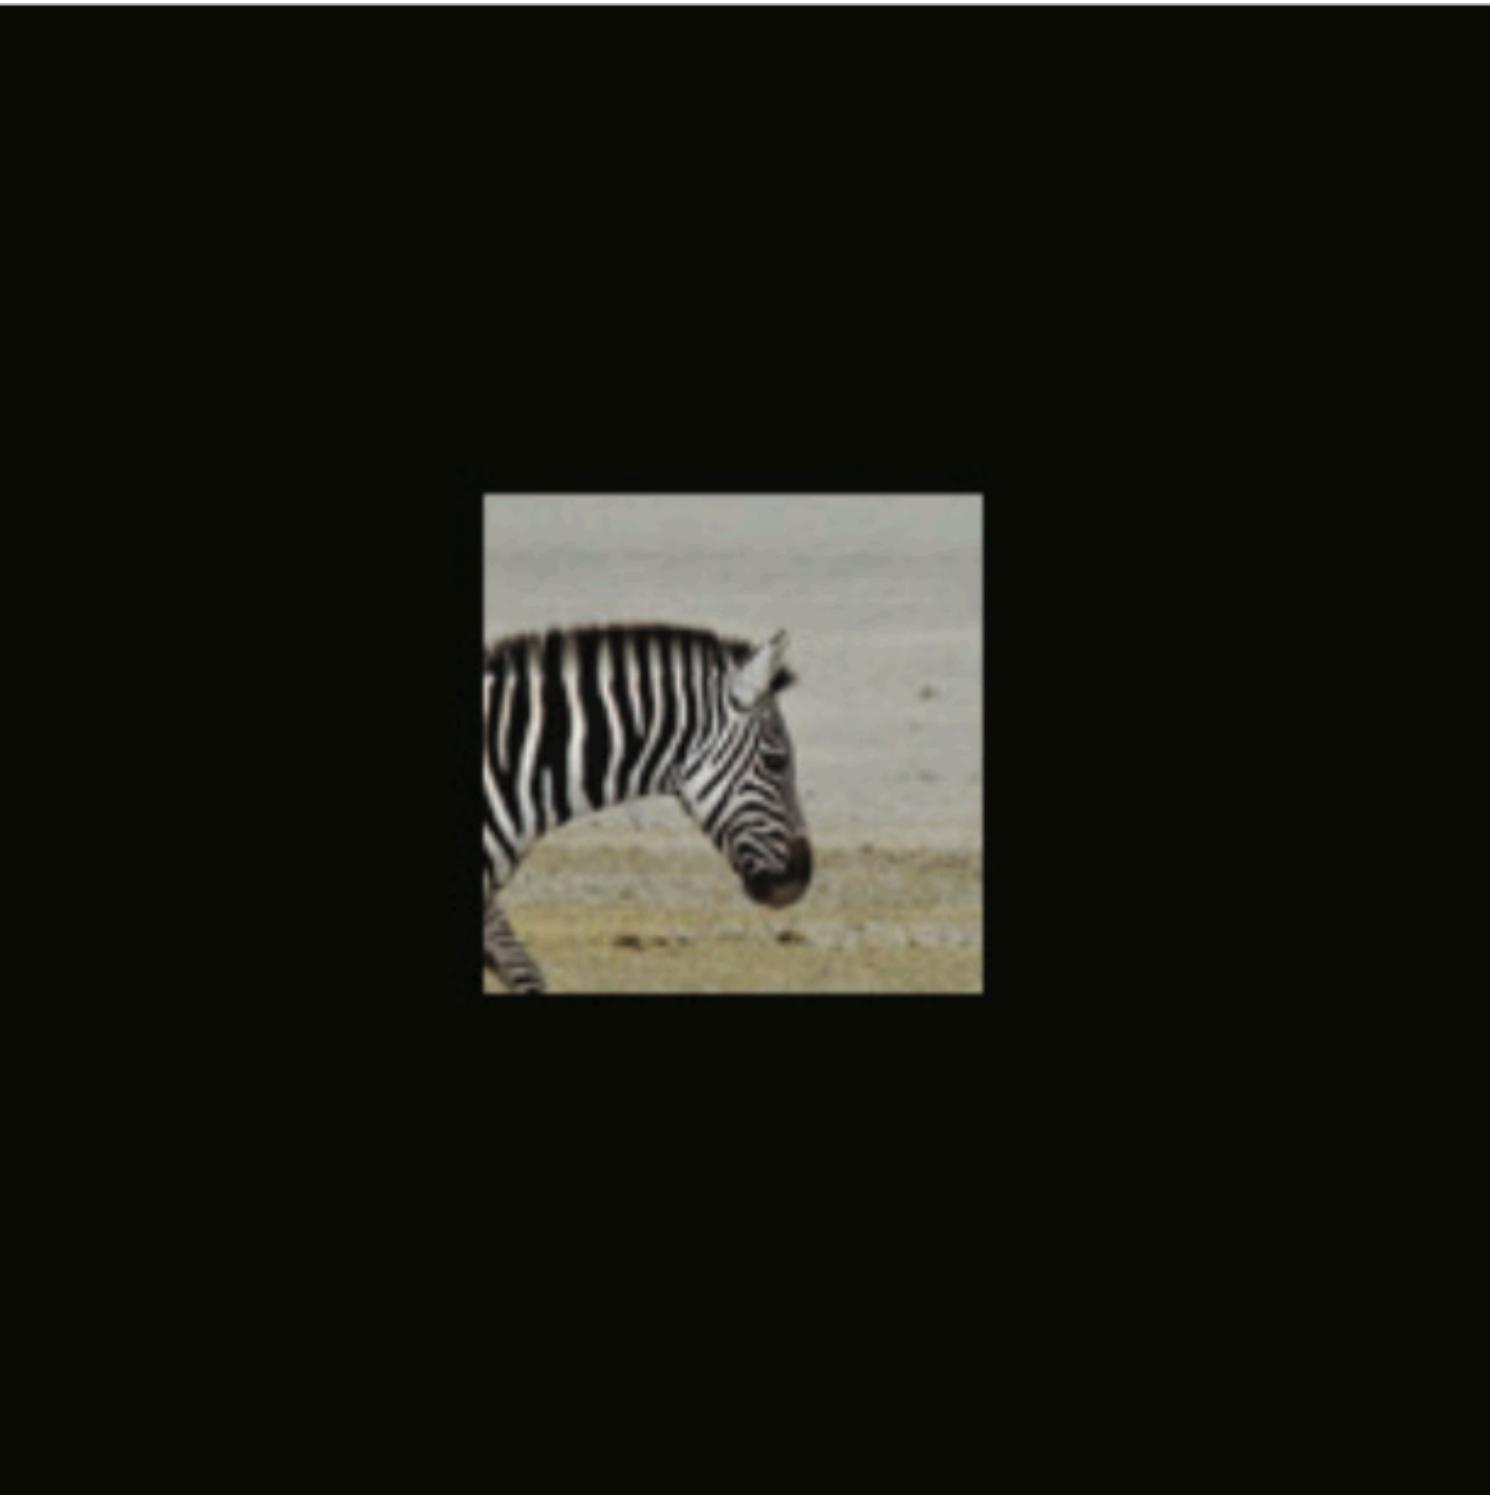
Output

Handling boundaries

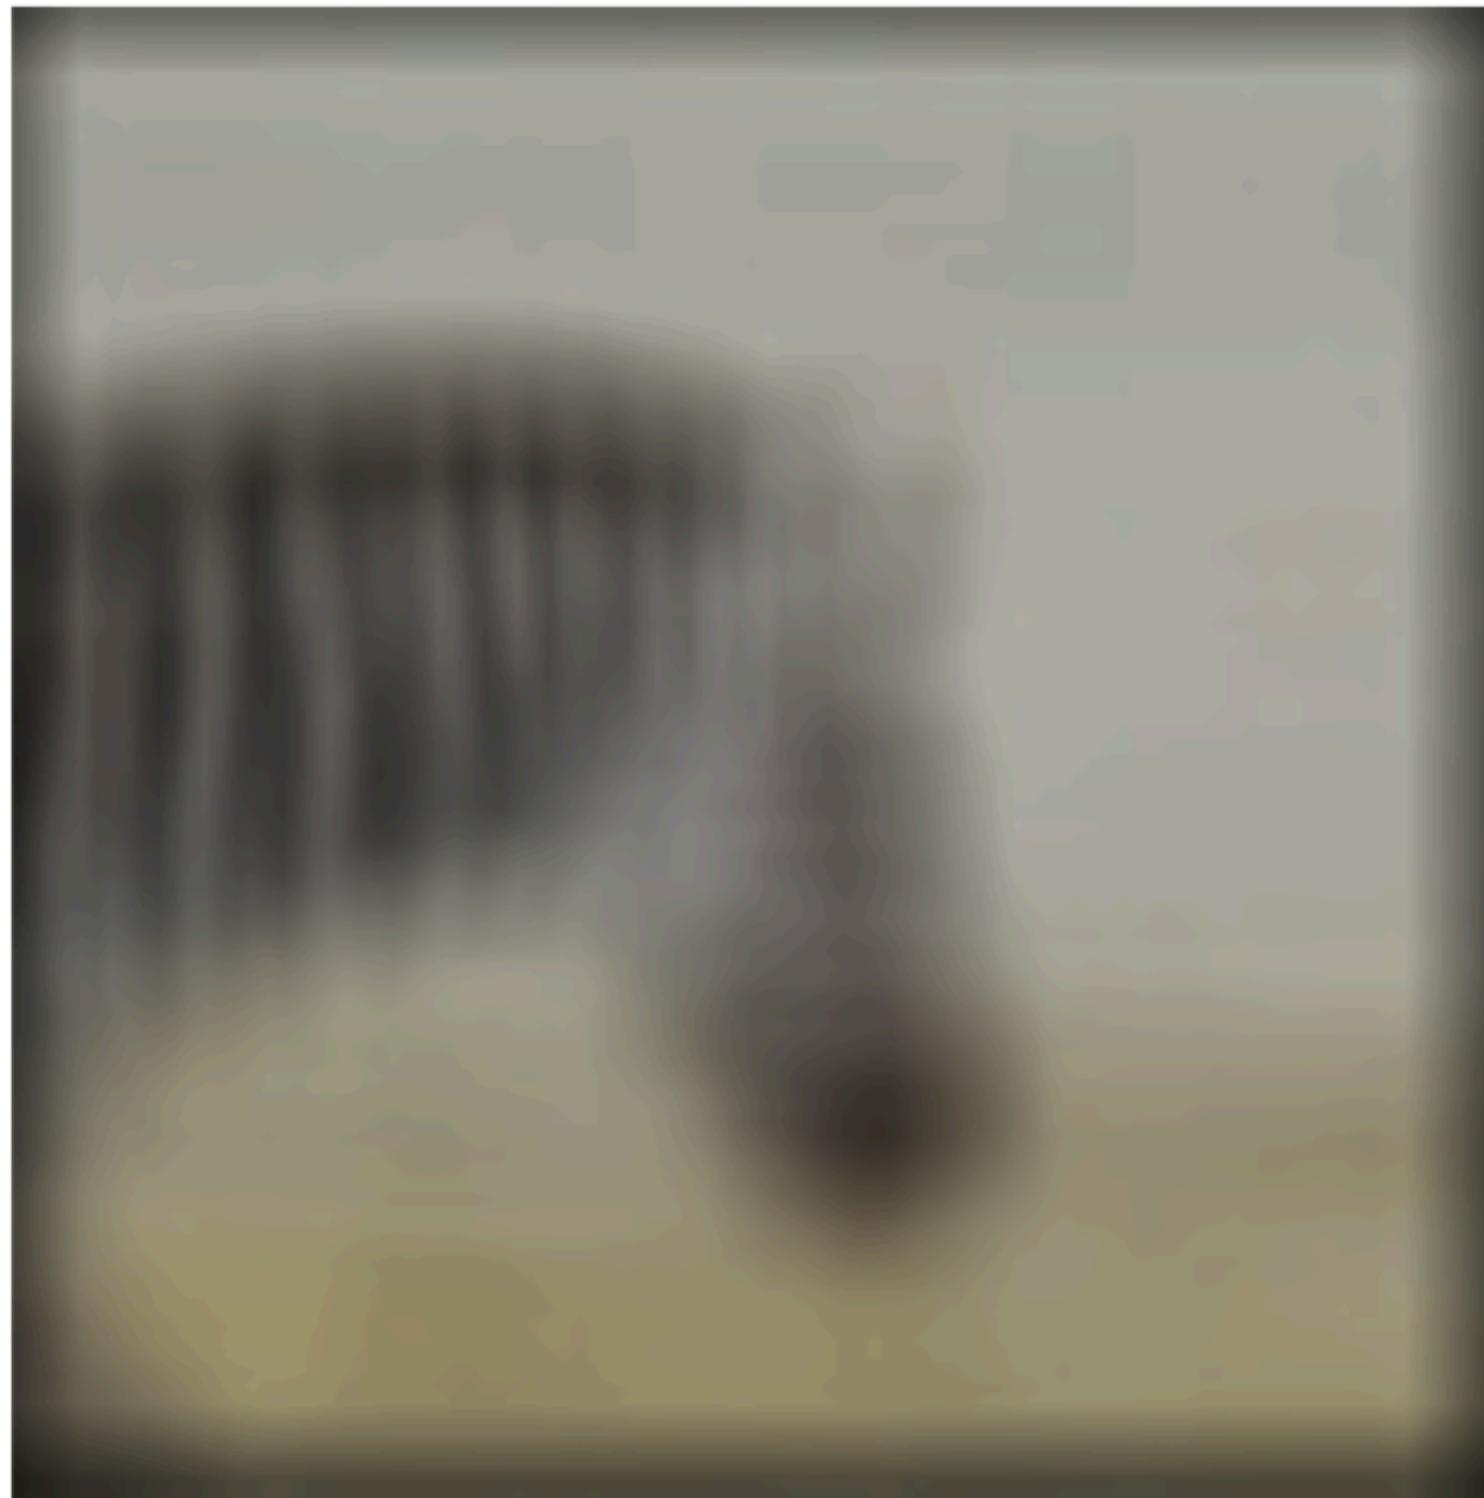


Handling boundaries

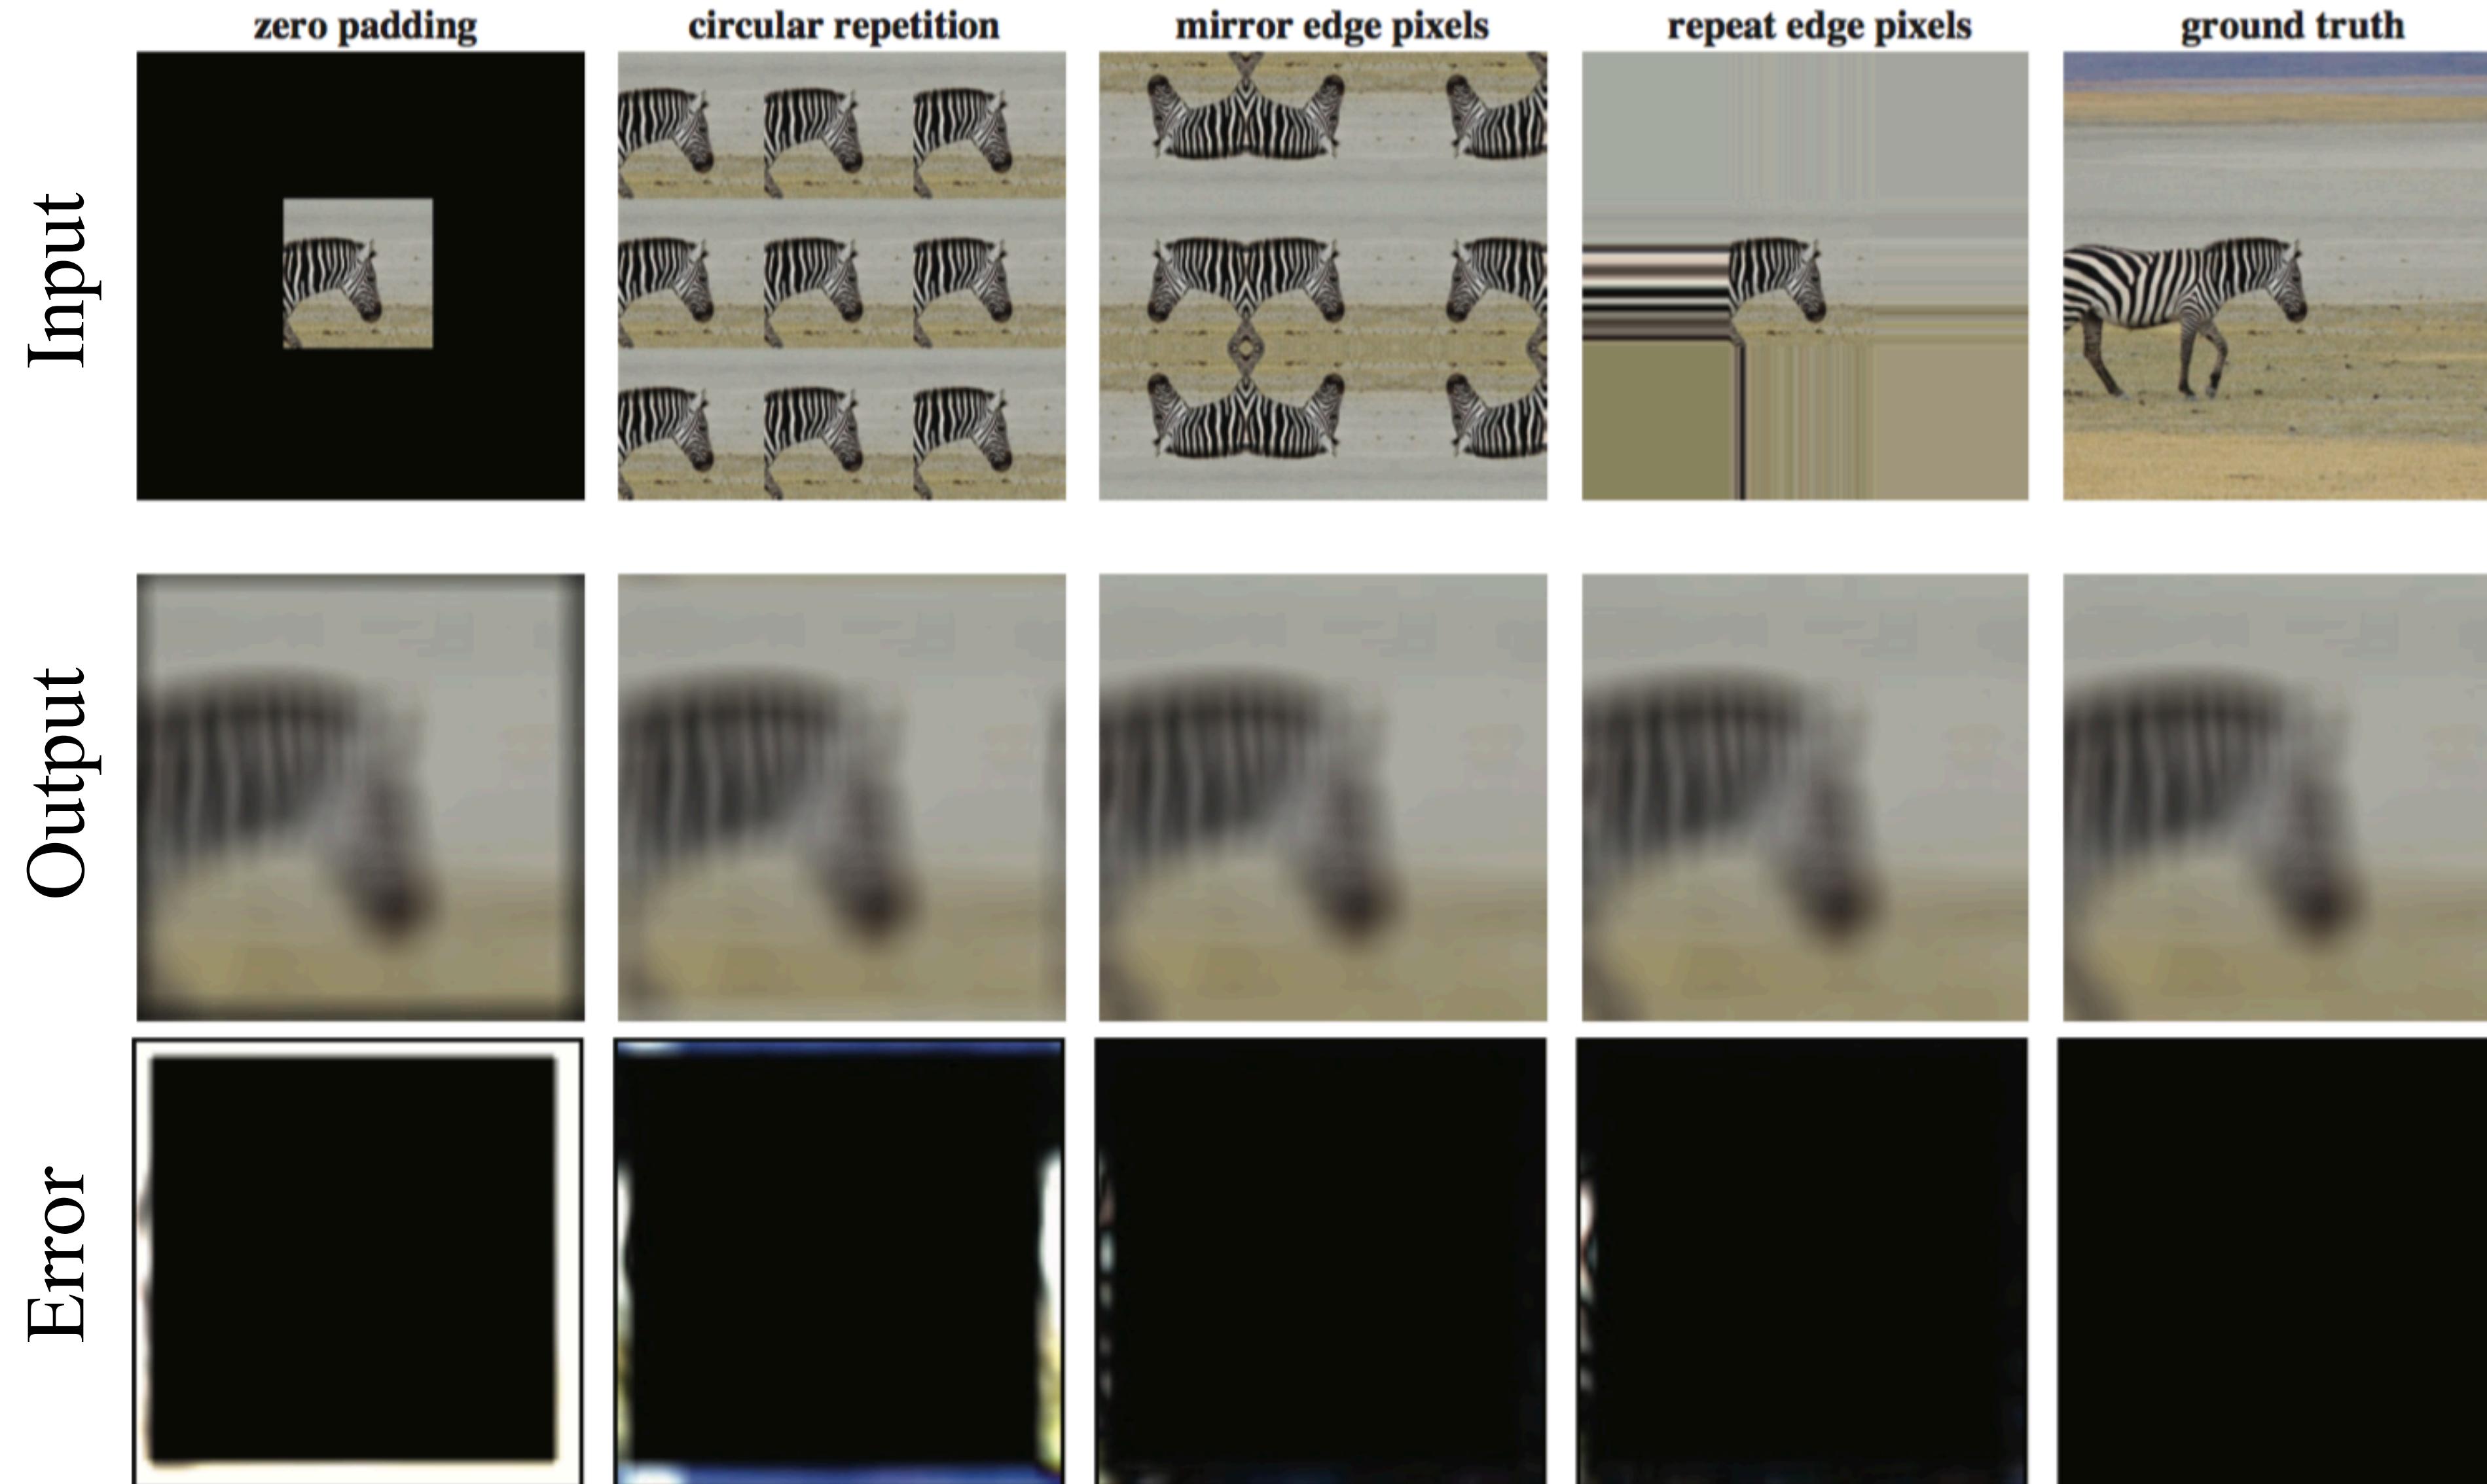
Zero padding



○  =
↑
11x11 box



Handling boundaries



Moving average

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

Filter kernel

0	0	0	0	0	0	0	0
0	0	90	90	90	90	0	0
0	0	90	90	90	90	0	0
0	0	90	90	90	90	0	0
0	0	90	90	90	90	0	0
0	0	90	0	90	90	0	0
0	0	90	90	90	90	0	0
0	0	0	0	0	0	0	0

Input

		40	60	60	40	20
		60	90	60	40	20
		50	80	80	60	30
		30	50	80	80	60
		30	50	50	40	20

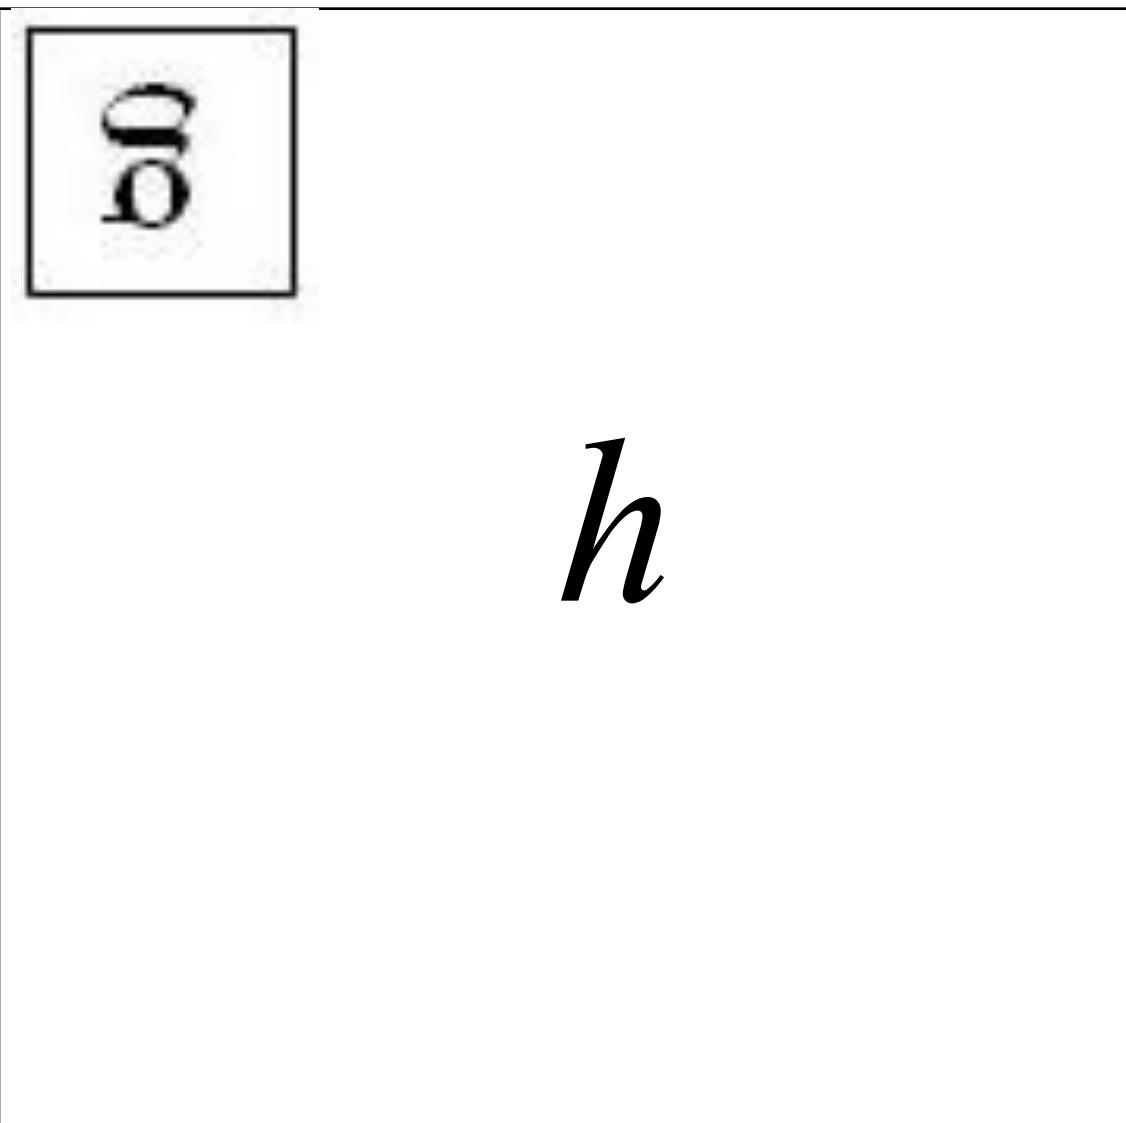
Output

Convolution

- Let h be the image and g be the kernel. The output of convolving h with g is:

$$f[m, n] = h \circ g = \sum_{k,l} h[m - k, n - l] g[k, l]$$

Convention:
kernel is “flipped”



Properties of the convolution

Commutative

$$h[n] \circ g[n] = g[n] \circ h[n]$$

Associative

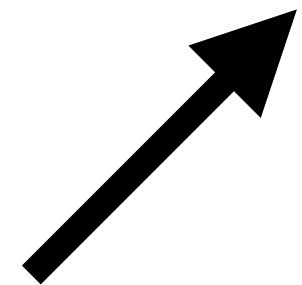
$$h[n] \circ g[n] \circ q[n] = h[n] \circ (g[n] \circ q[n]) = (h[n] \circ g[n]) \circ q[n]$$

Distributive with respect to the sum

$$h[n] \circ (f[n] + g[n]) = h[n] \circ f[n] + h[n] \circ g[n]$$

Why flip the kernel?

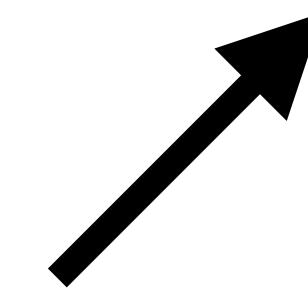
$$f[m, n] = h \circ g = \sum_{k,l} h[m - k, n - l] g[k, l]$$



Indexes go backward!

Cross correlation

$$f[m, n] = h * g = \sum_{k,l} h[m + k, n + l] g[k, l]$$

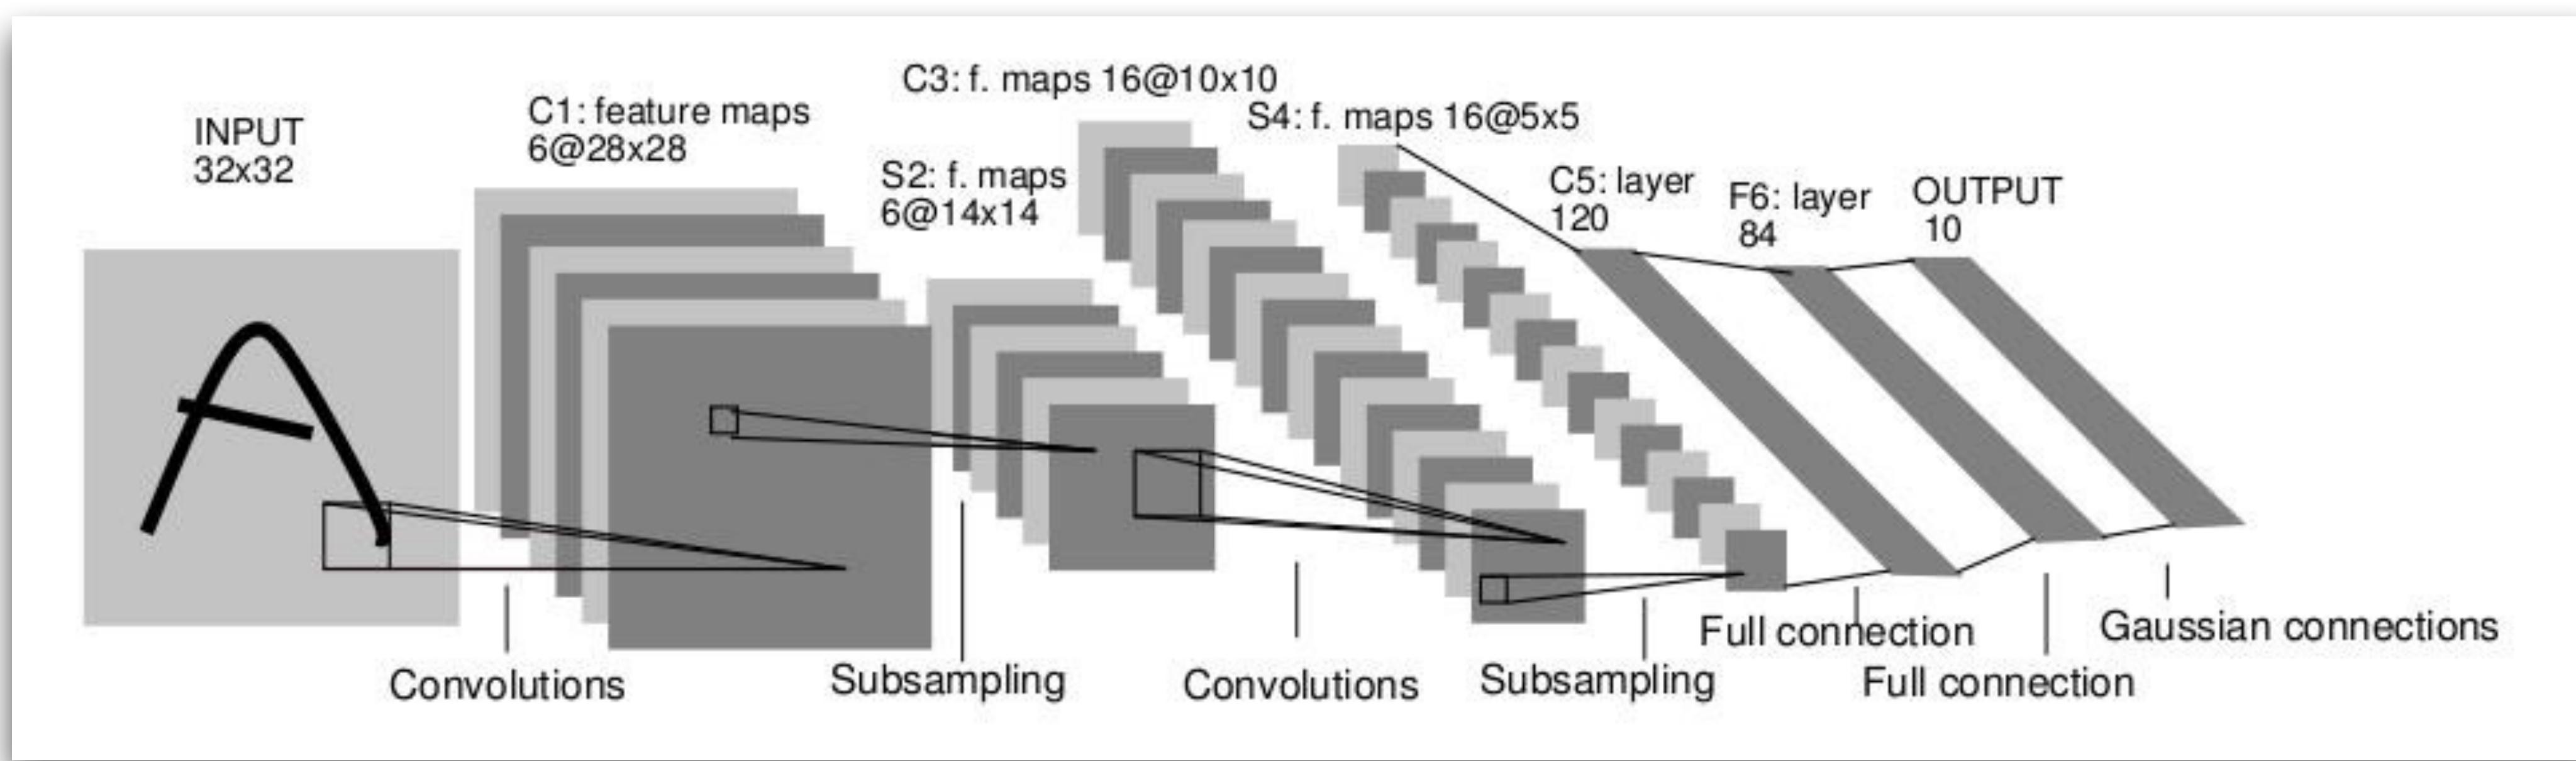


No flipping!

- Sometimes called just **correlation**
- Neither associative nor commutative
- In the literature, people often just call both “convolution”
- Filters often symmetric, so won’t matter

Convolutional neural networks

- Neural network with specialized connectivity structure
- Mostly just convolutions!

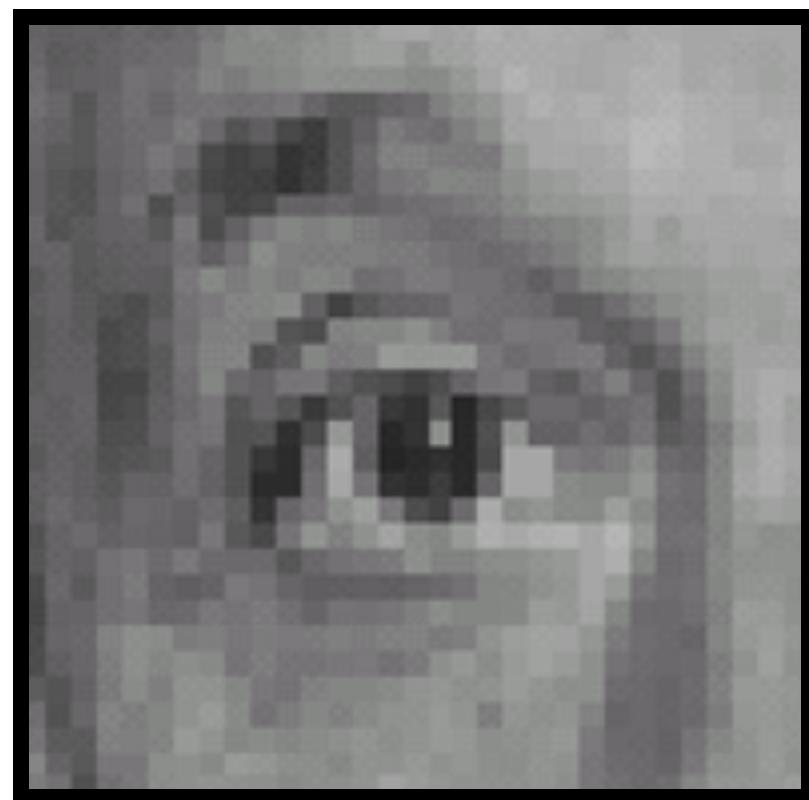


(LeCun et al. 1989)

Source: Torralba, Freeman, Isola

Filtering examples

Practice with linear filters

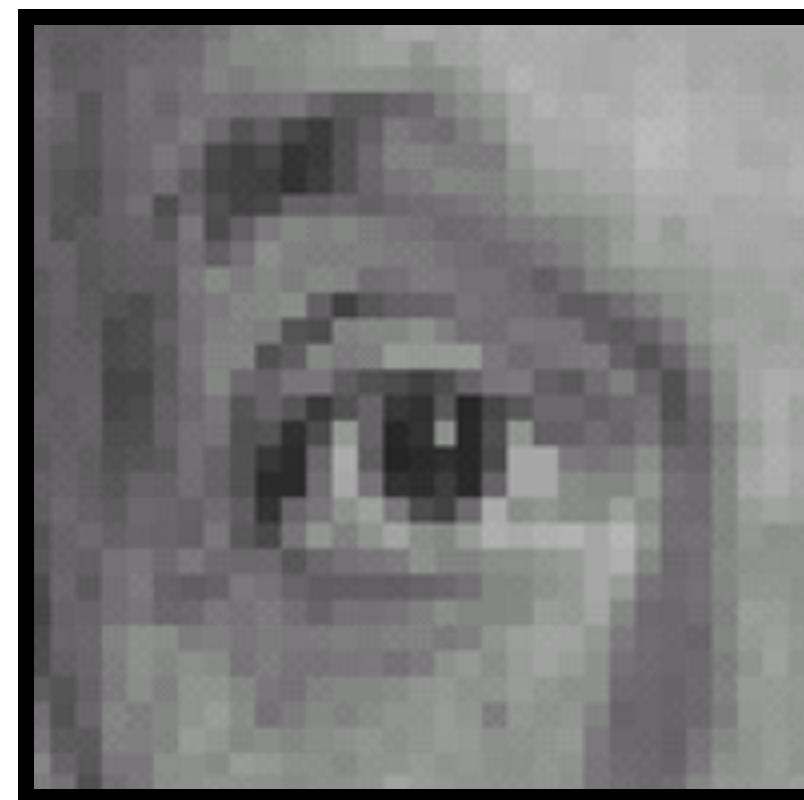


Original

0	0	0
0	1	0
0	0	0

?

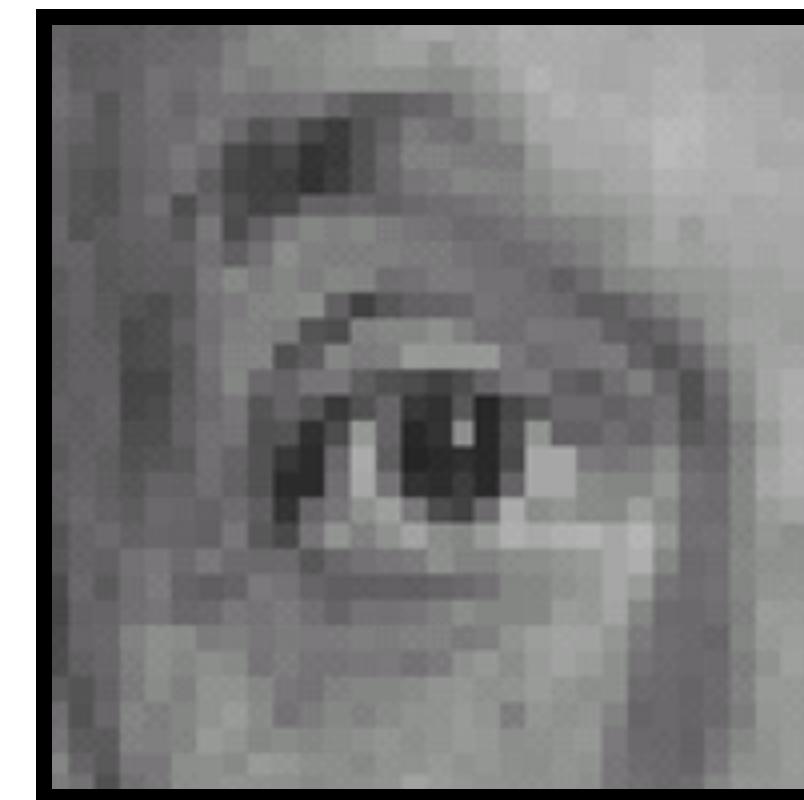
Practice with linear filters



Original

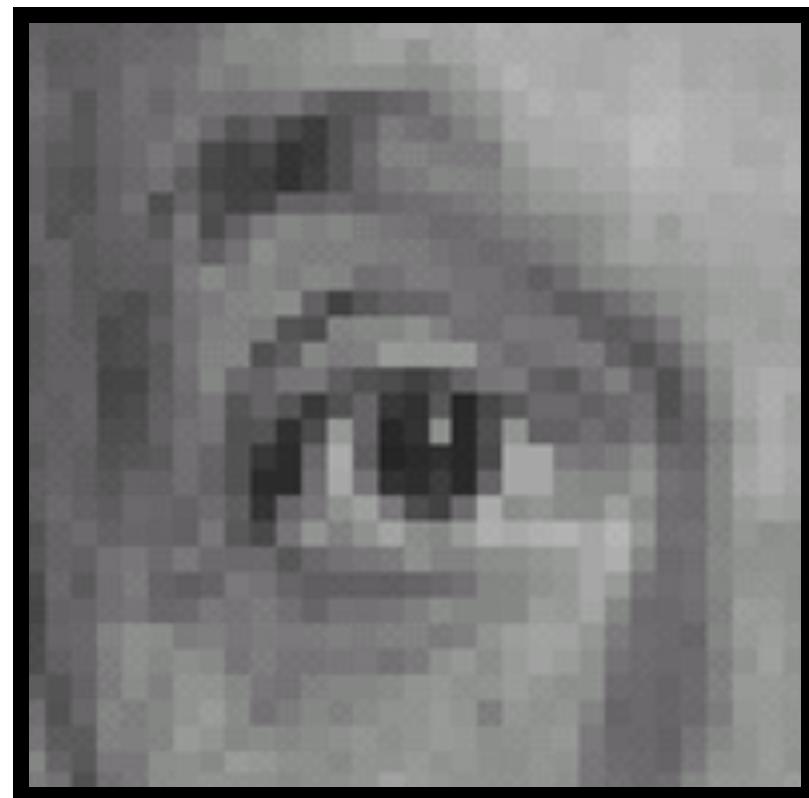
0	0	0
0	1	0
0	0	0

“Impulse”



Filtered
(no change)

Practice with linear filters



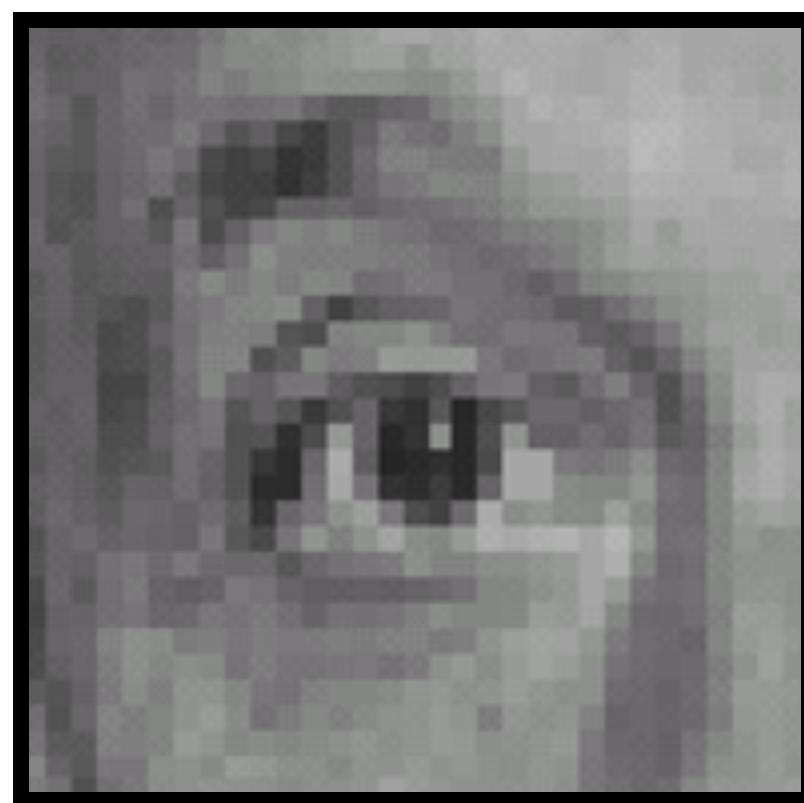
Original

0	0	0
0	0	1
0	0	0

?

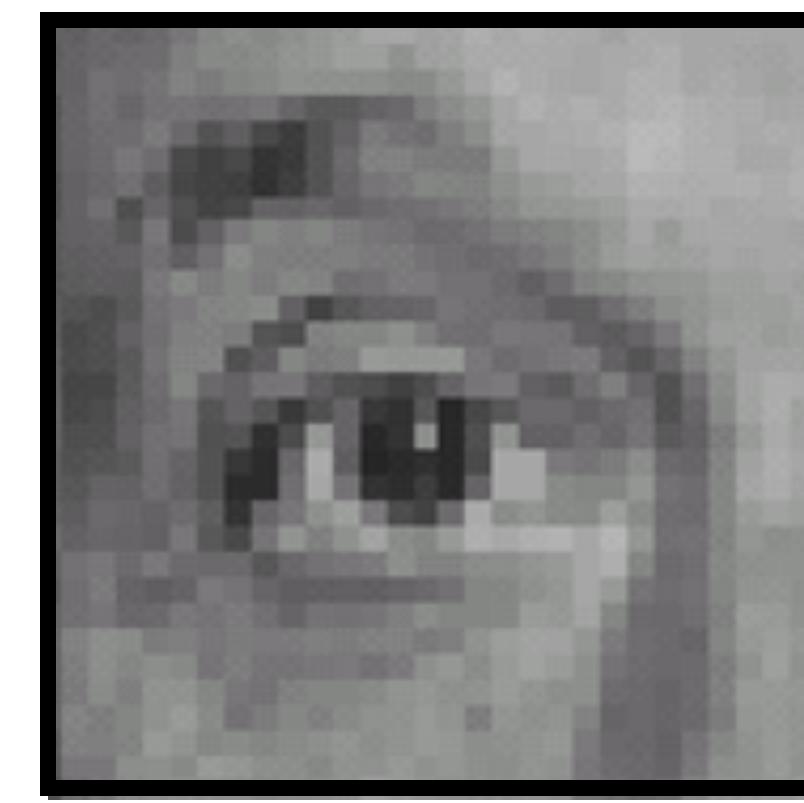
“Translated
Impulse”

Practice with linear filters



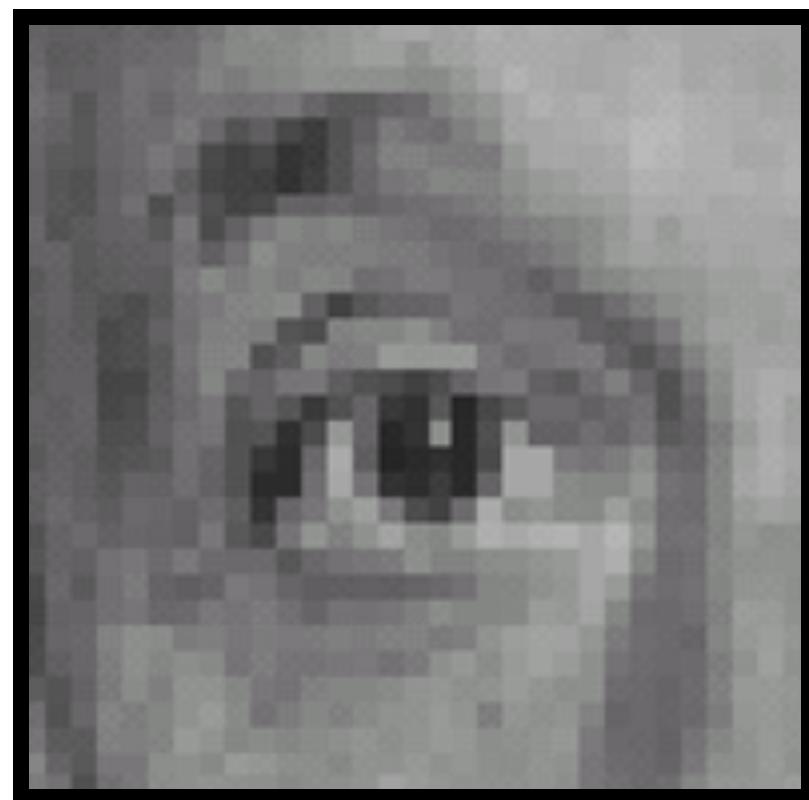
Original

0	0	0
0	0	1
0	0	0



Shifted *left*
By 1 pixel

Practice with linear filters

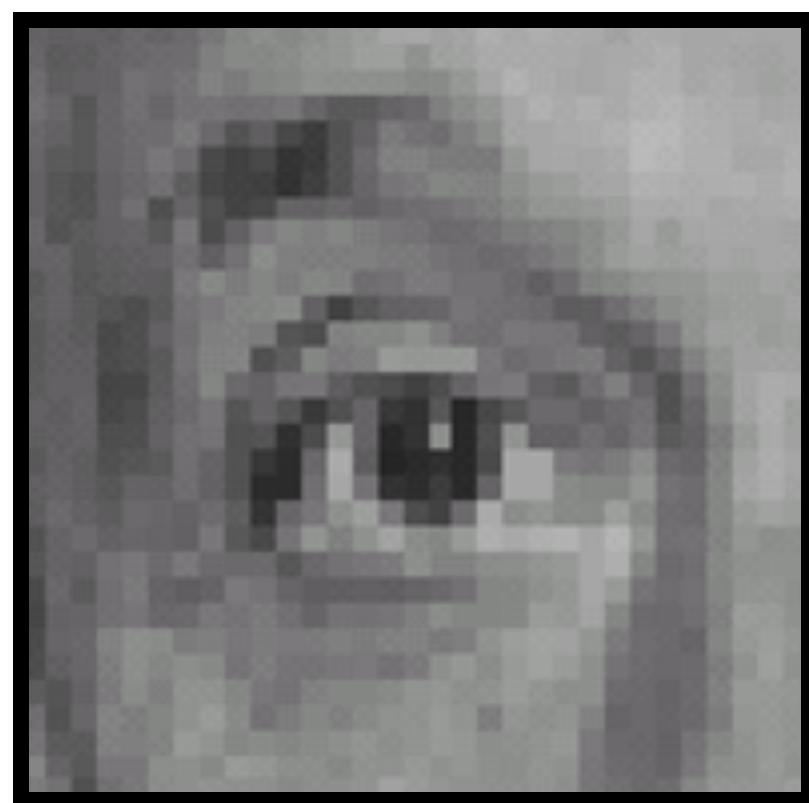


$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

?

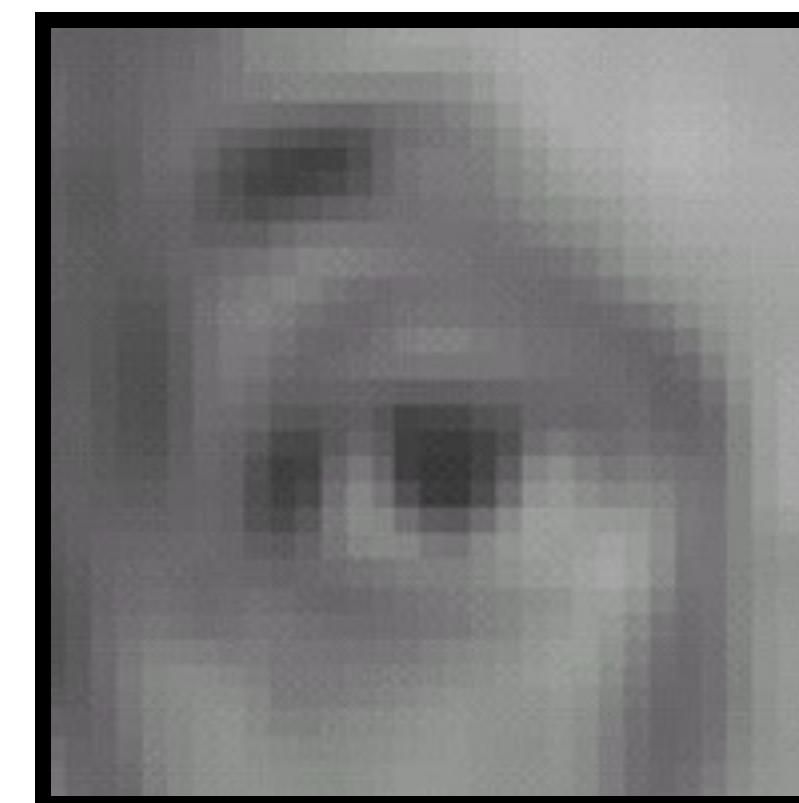
Original

Practice with linear filters



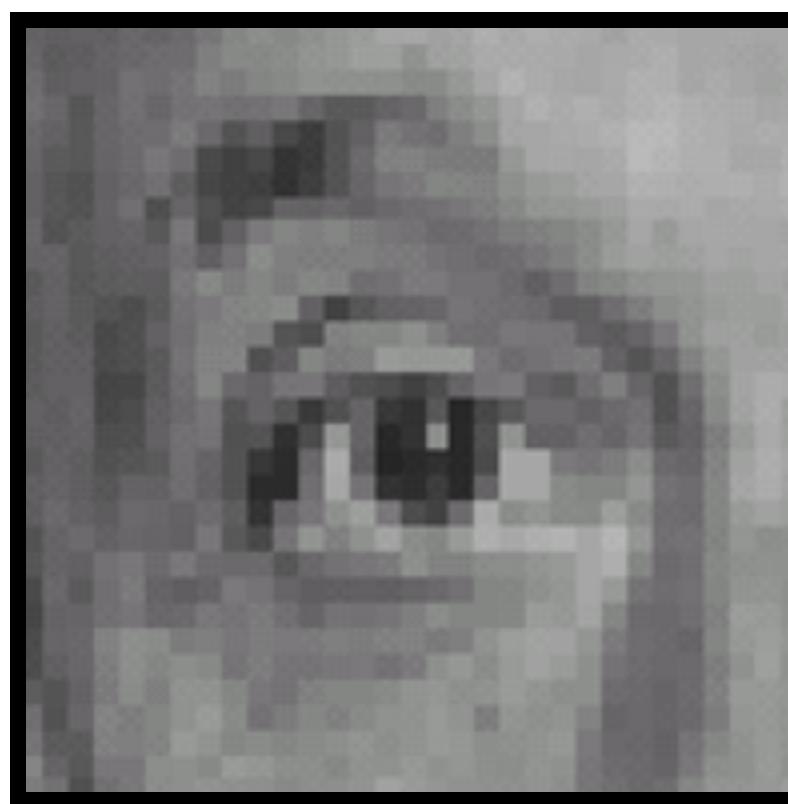
Original

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$



Blur (with a
box filter)

Practice with linear filters



Original

$$\begin{matrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{matrix}$$

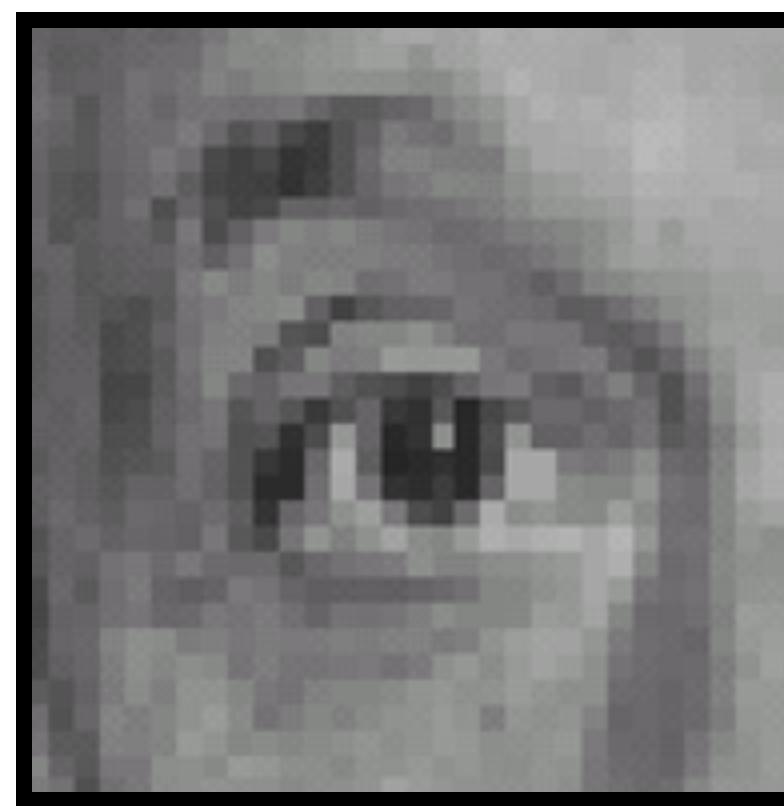
-

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

?

(Note that filter sums to 1)

Practice with linear filters



$$\begin{matrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{matrix}$$

-

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

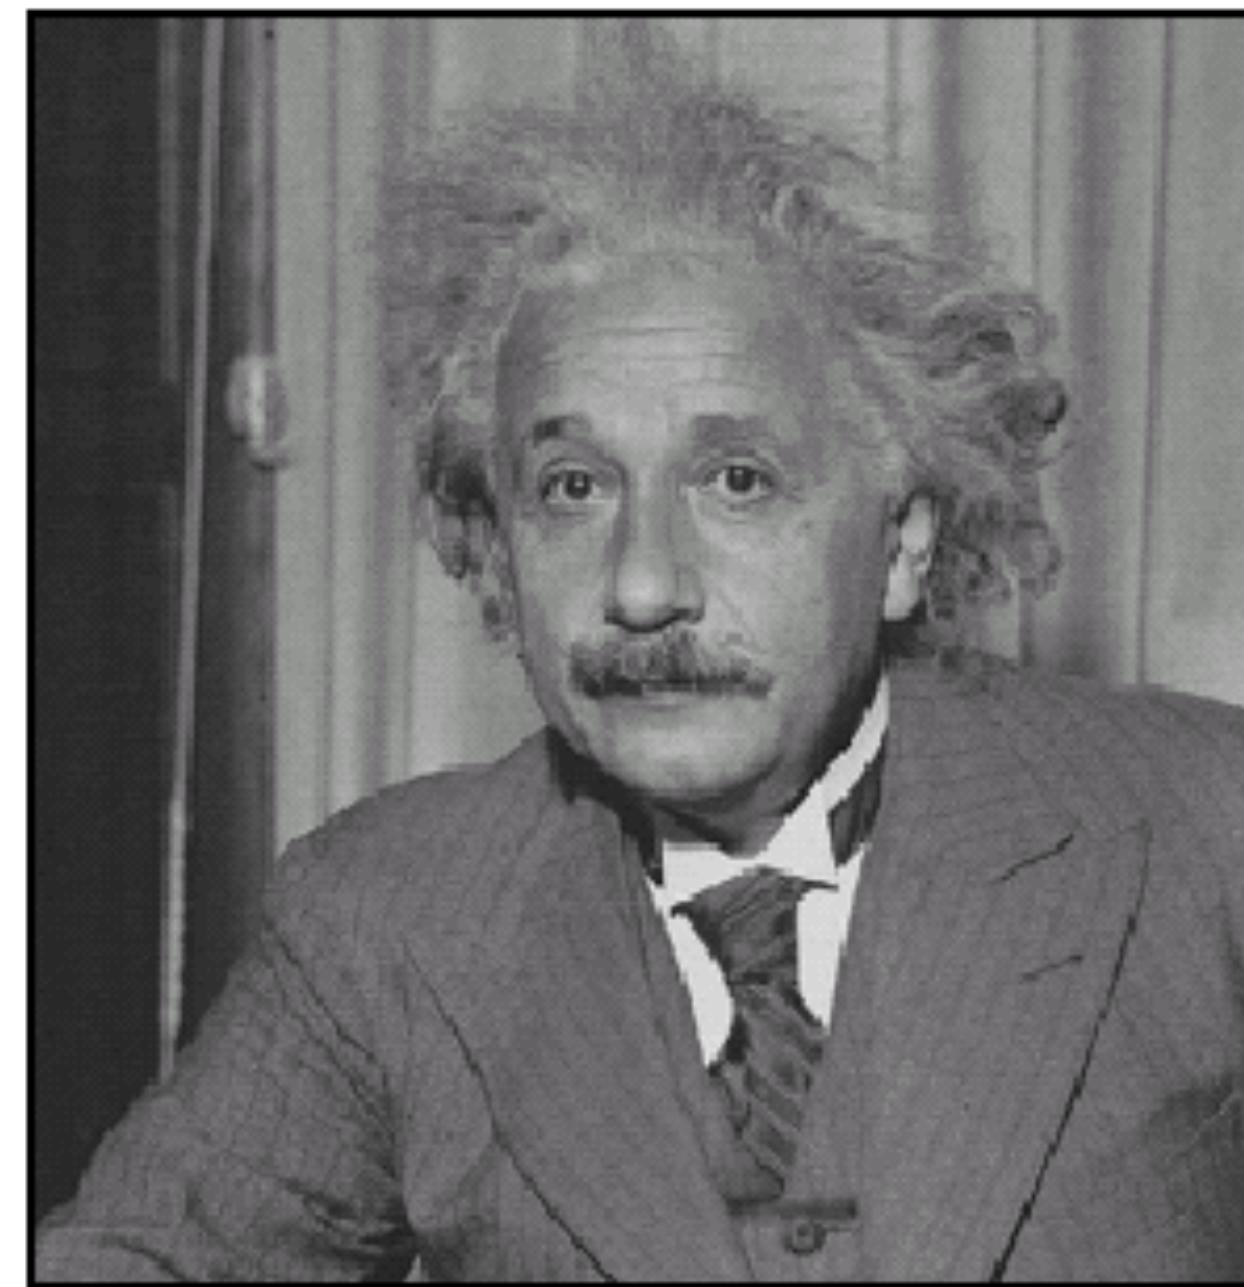


Original

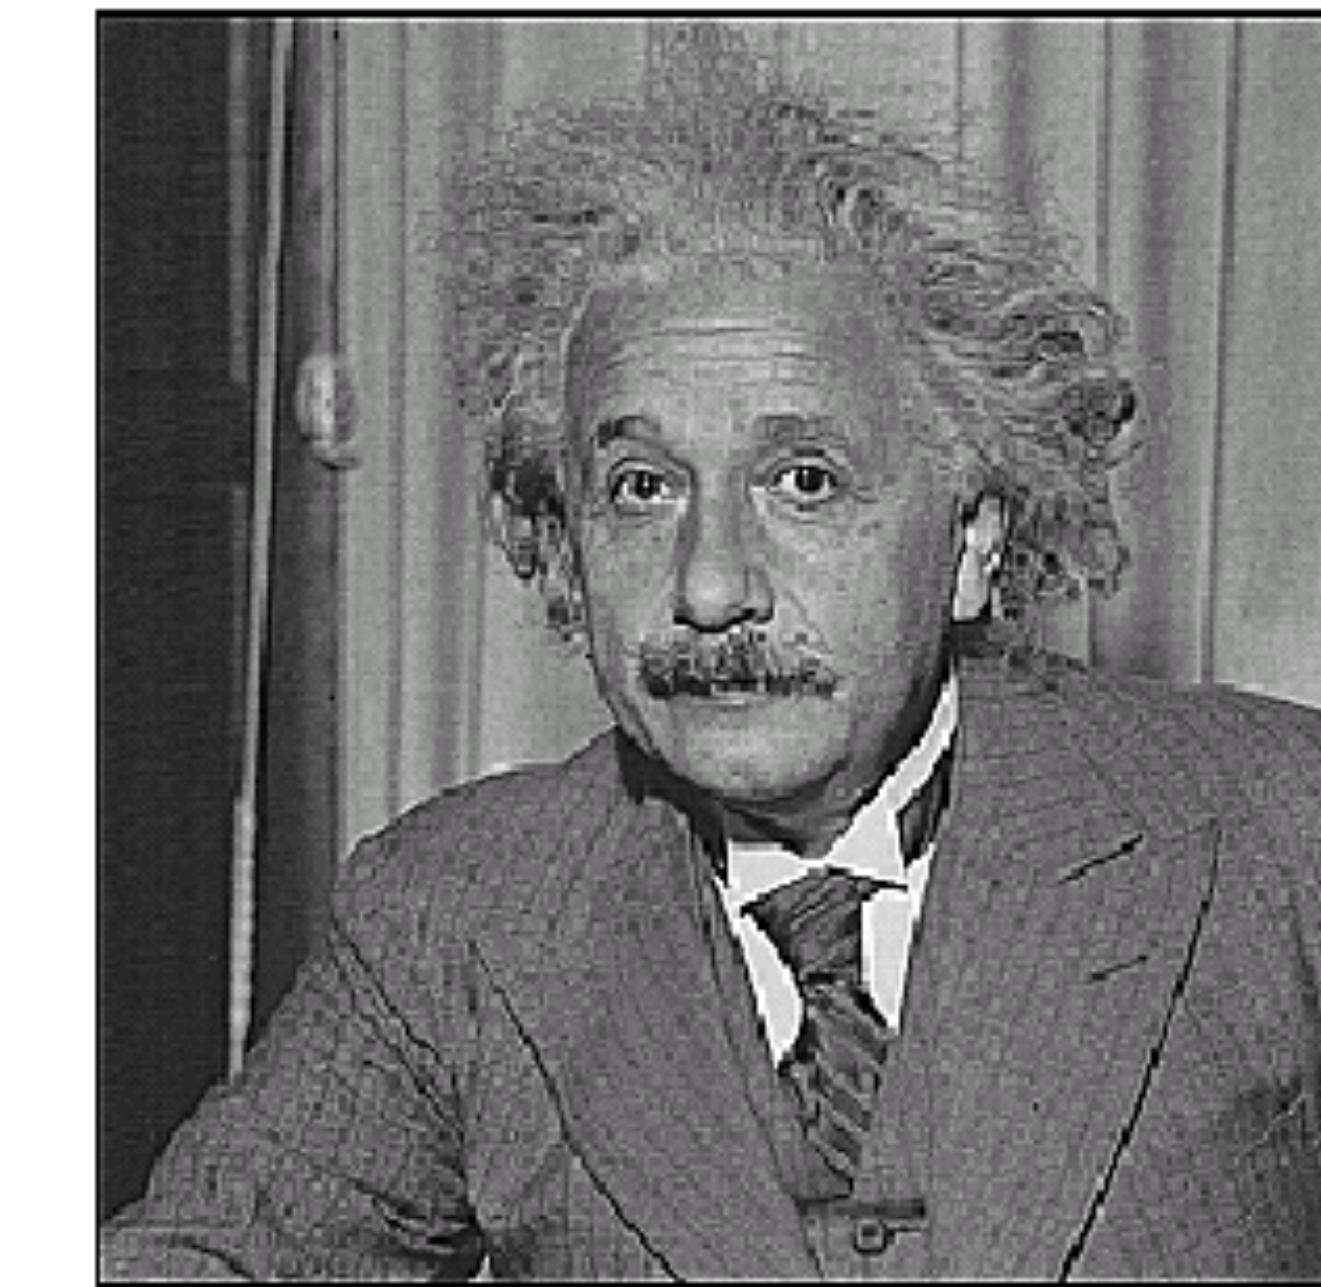
Sharpening filter

- Accentuates differences with local average

Sharpening

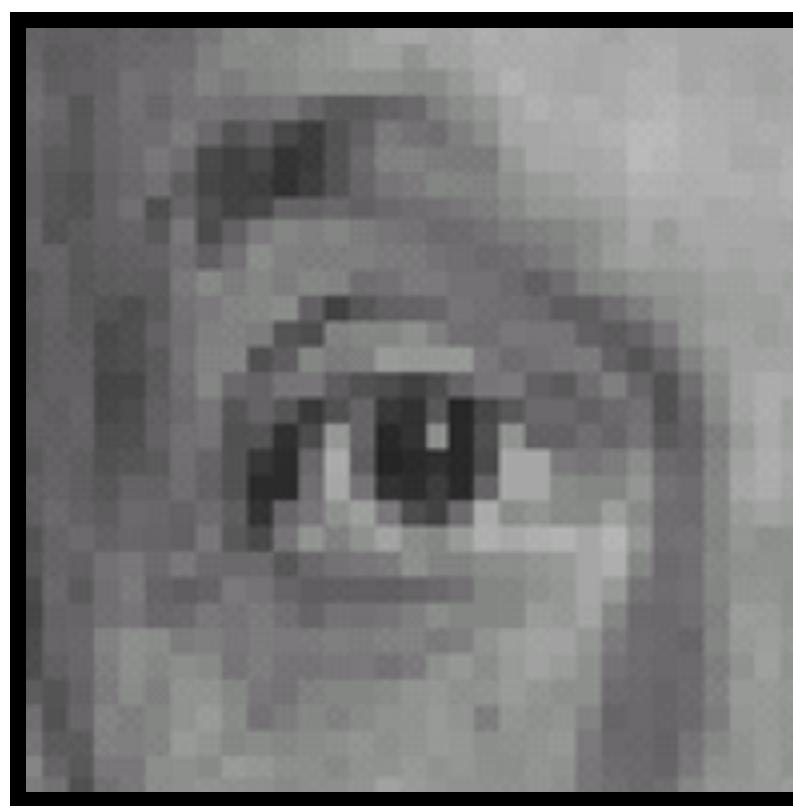


before

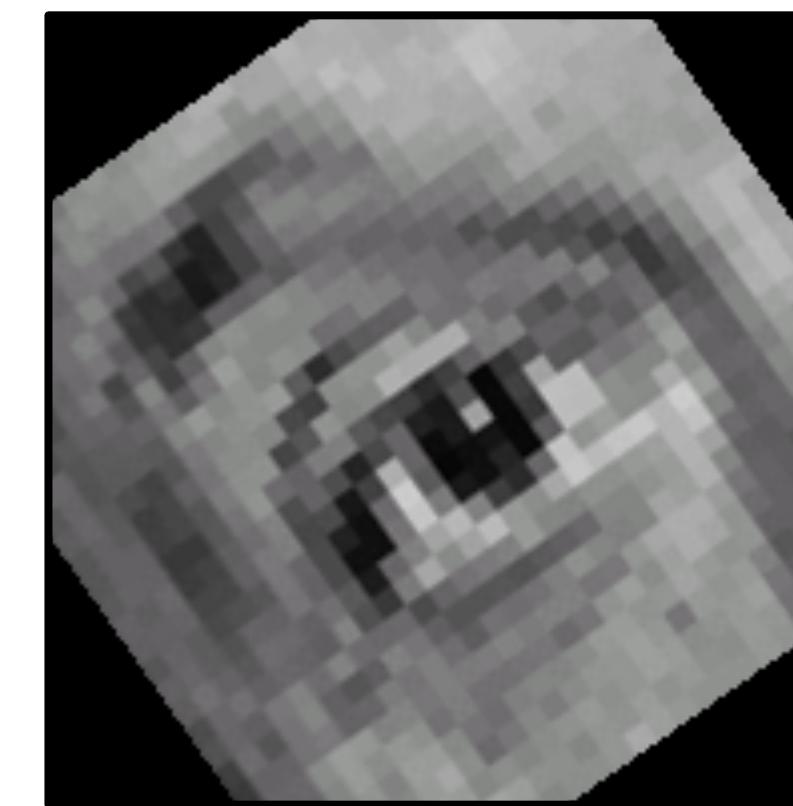
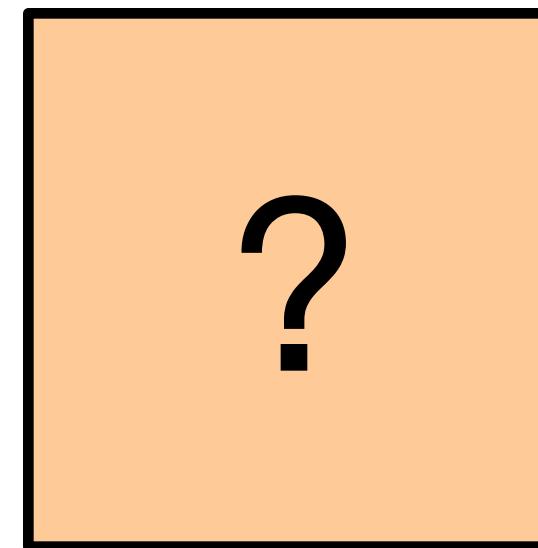


after

Practice with linear filters



Original



Can you do this?

Rectangular filter



$g[m,n]$



\otimes =

$h[m,n]$



$f[m,n]$

Rectangular filter



$g[m,n]$

\otimes

$h[m,n]$

=

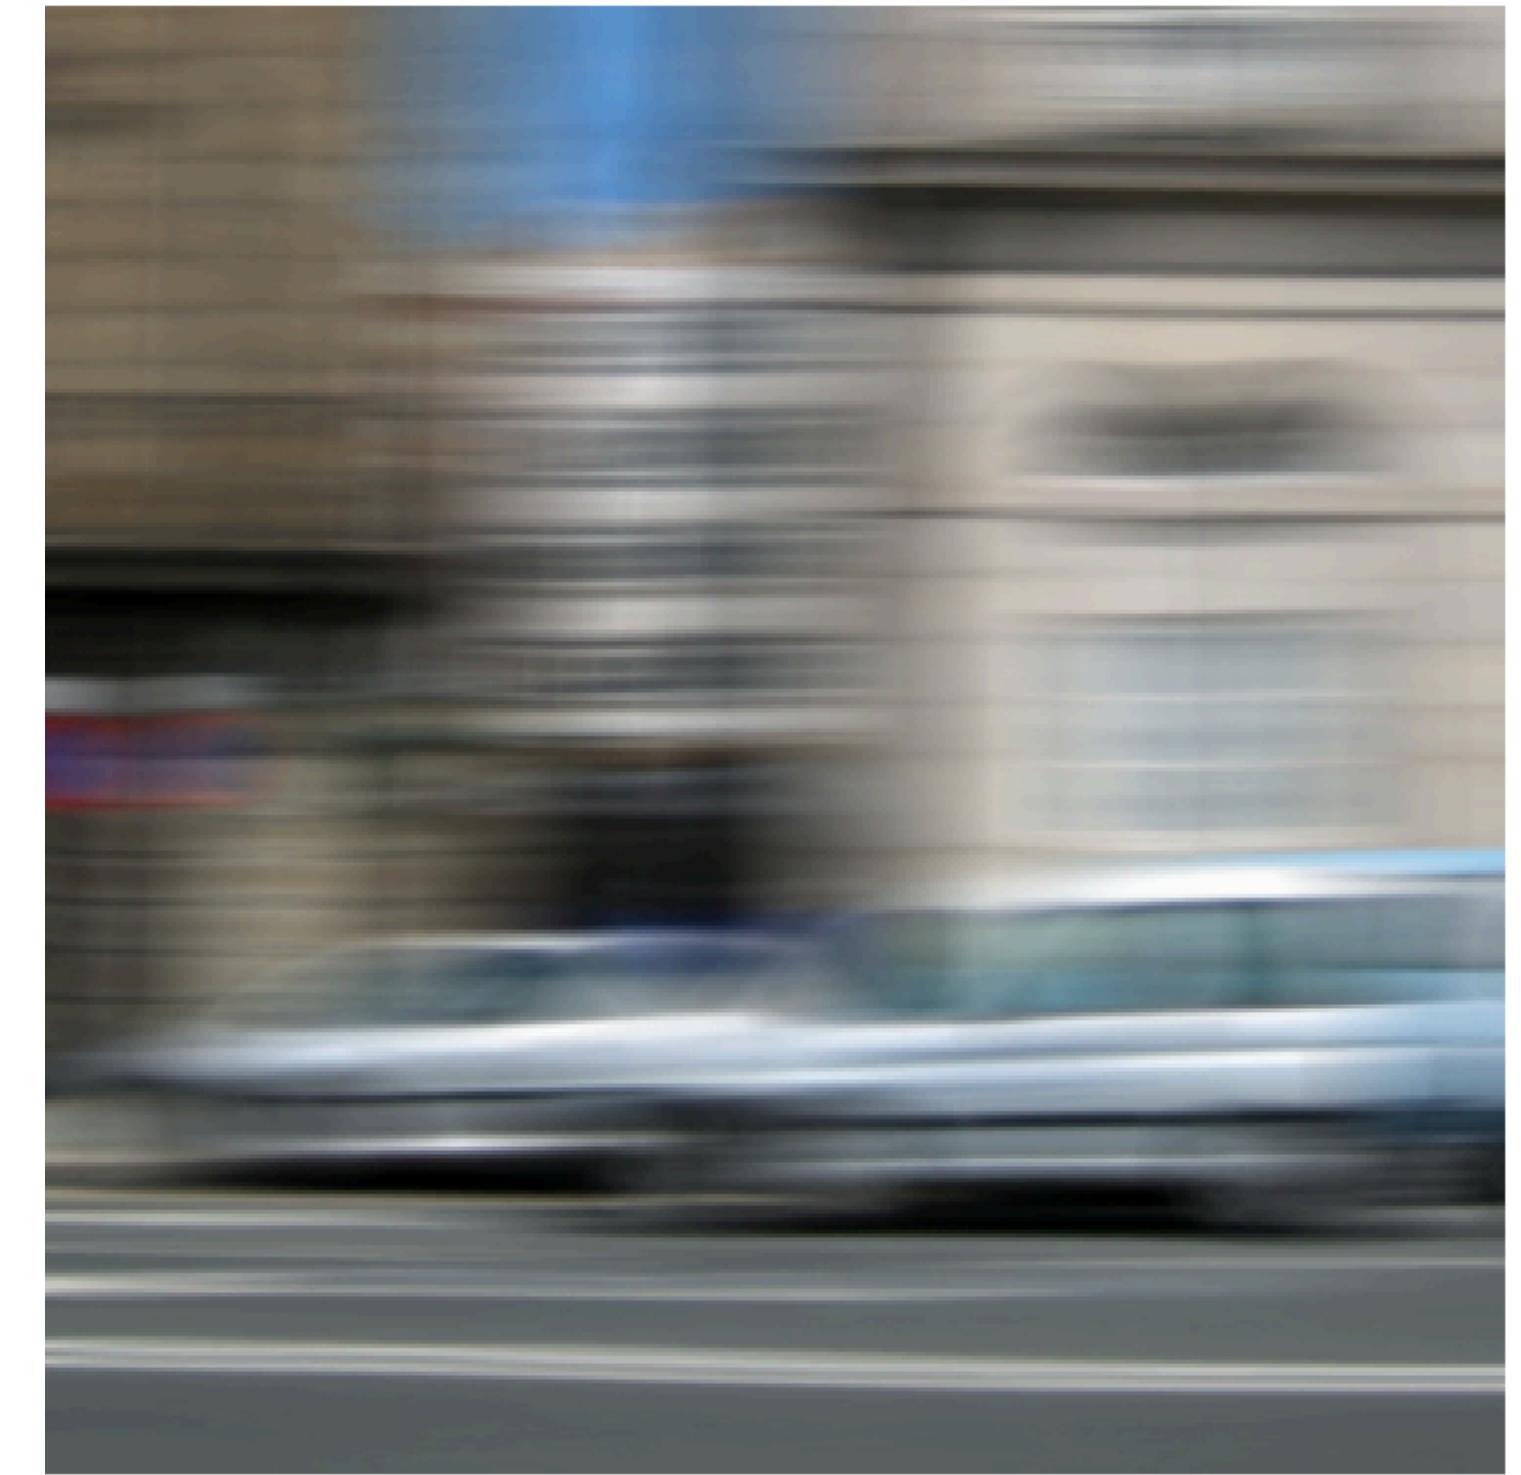


$f[m,n]$

“Naturally” occurring filters



Input image

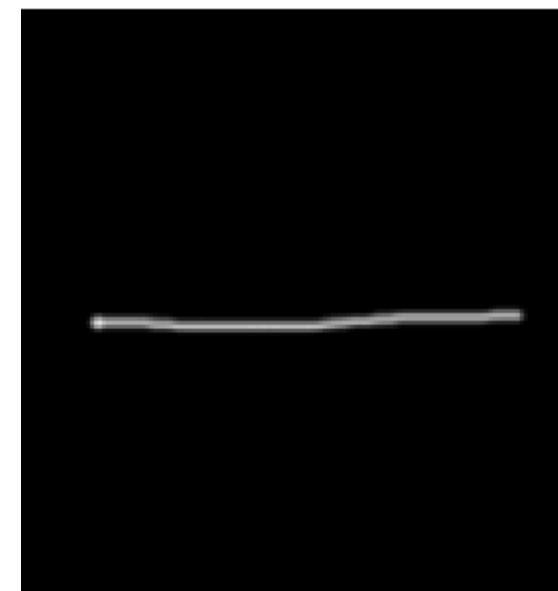


Motion blur

“Naturally” occurring filters



Input image

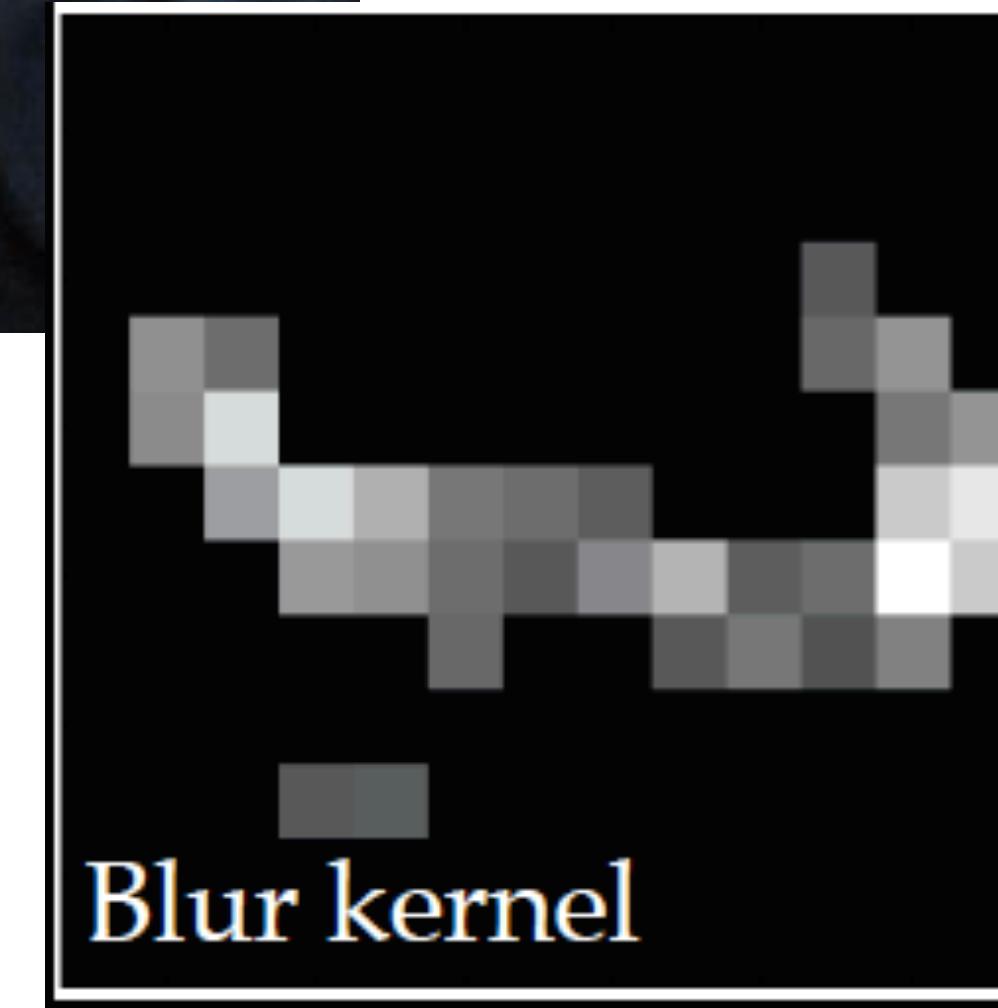
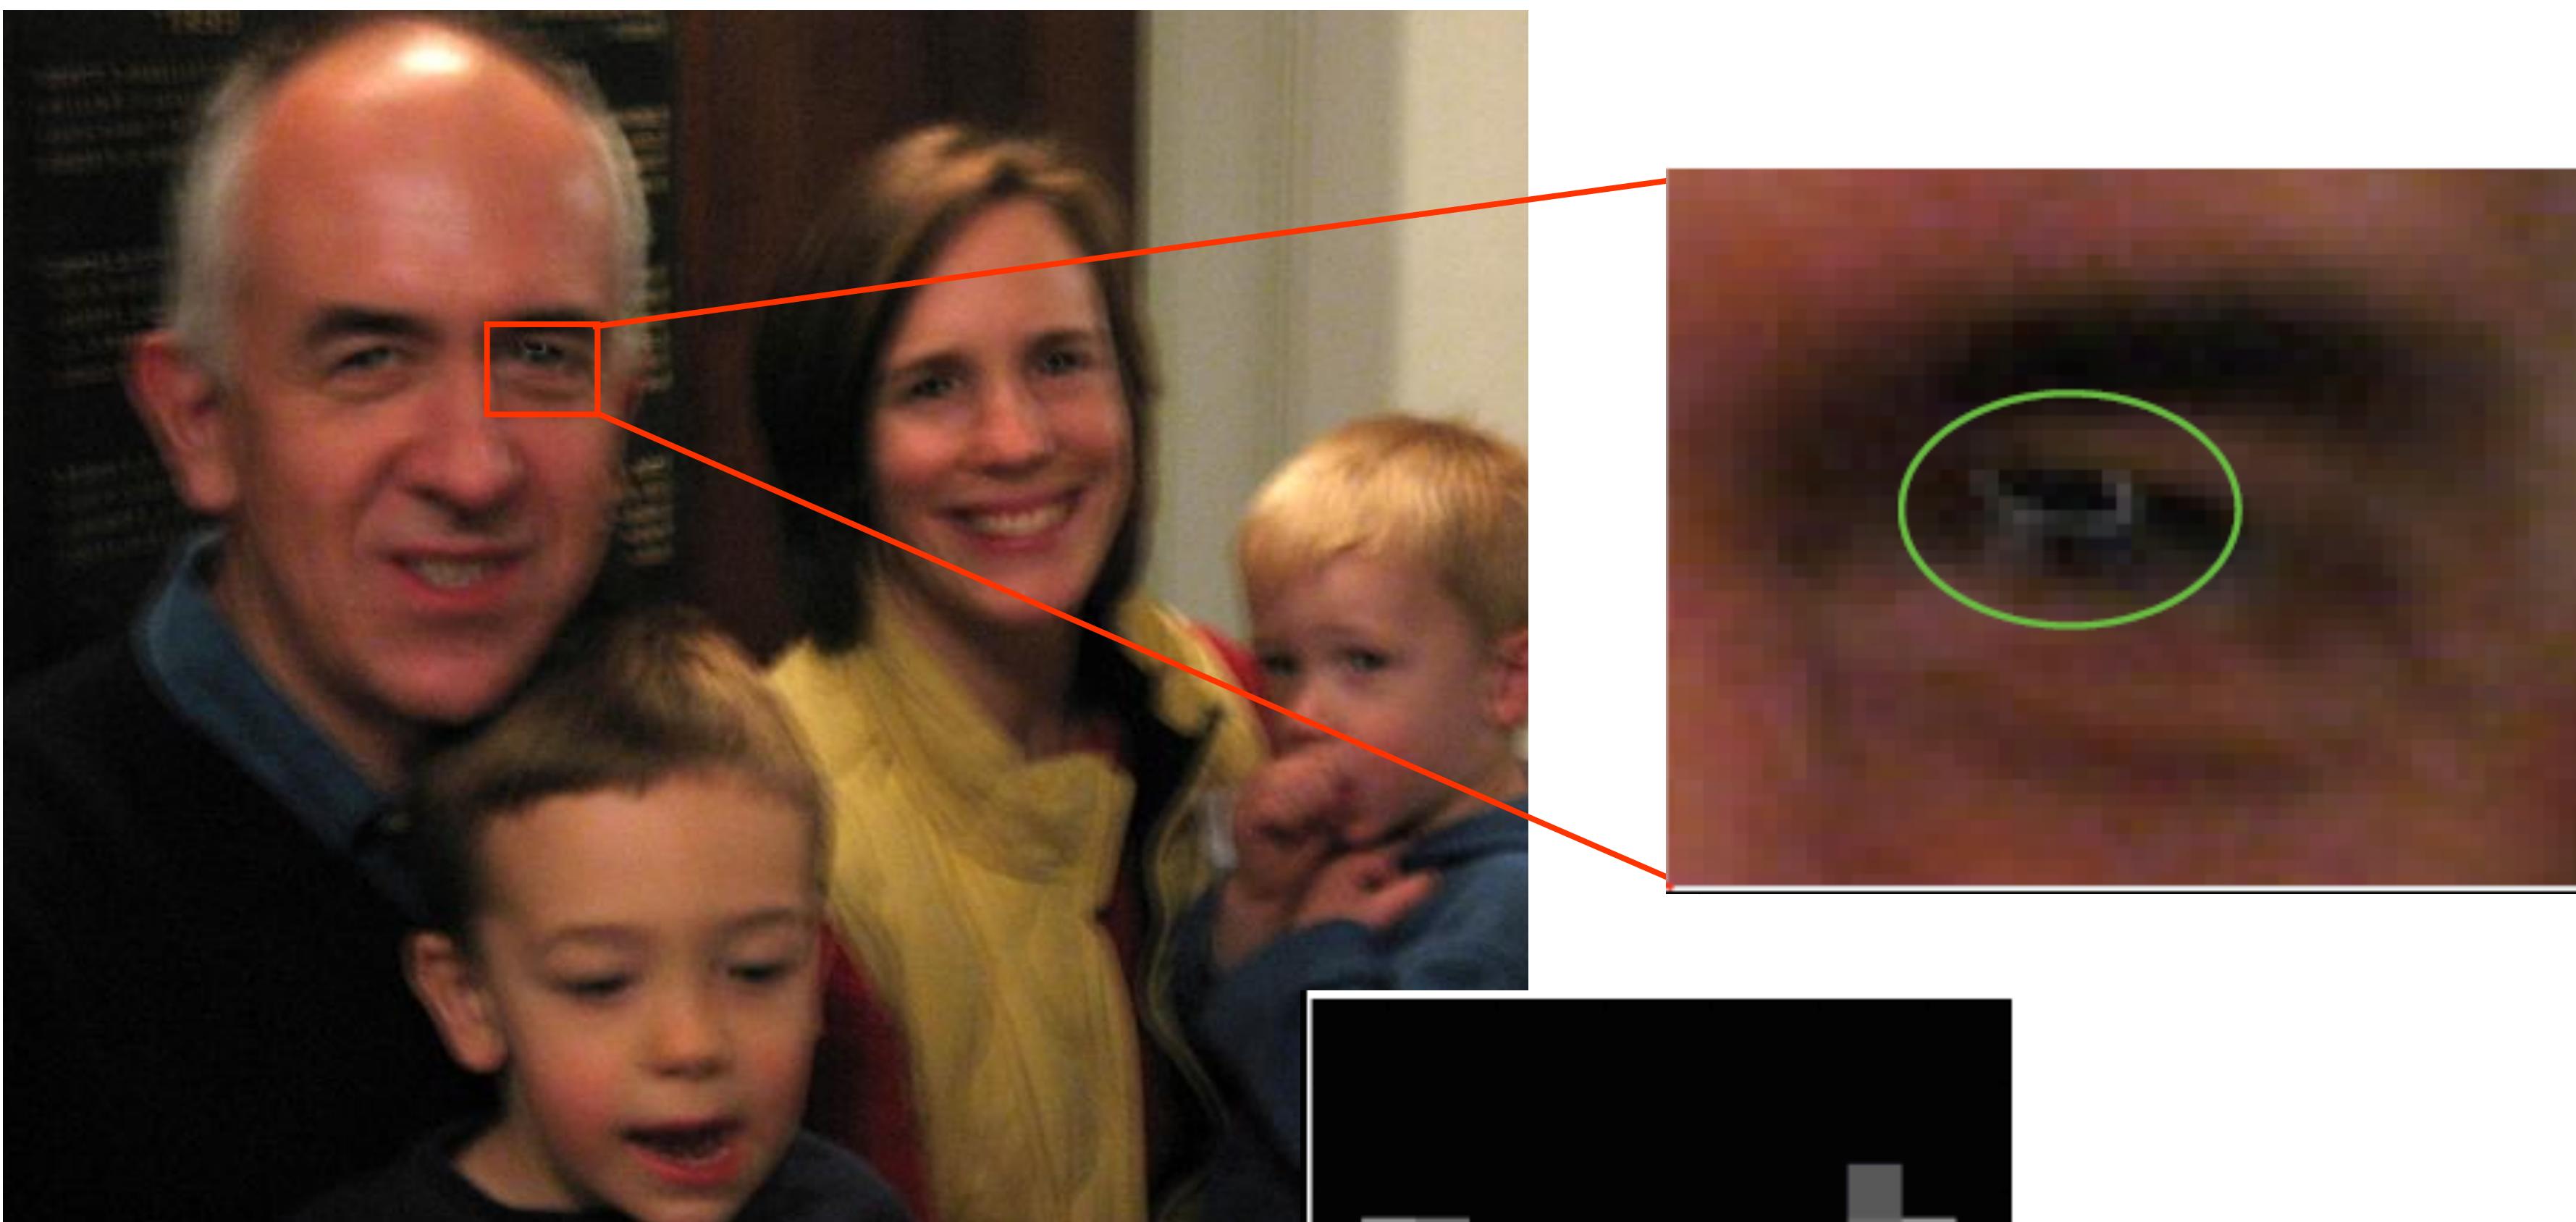


Convolution weights



Convolution output

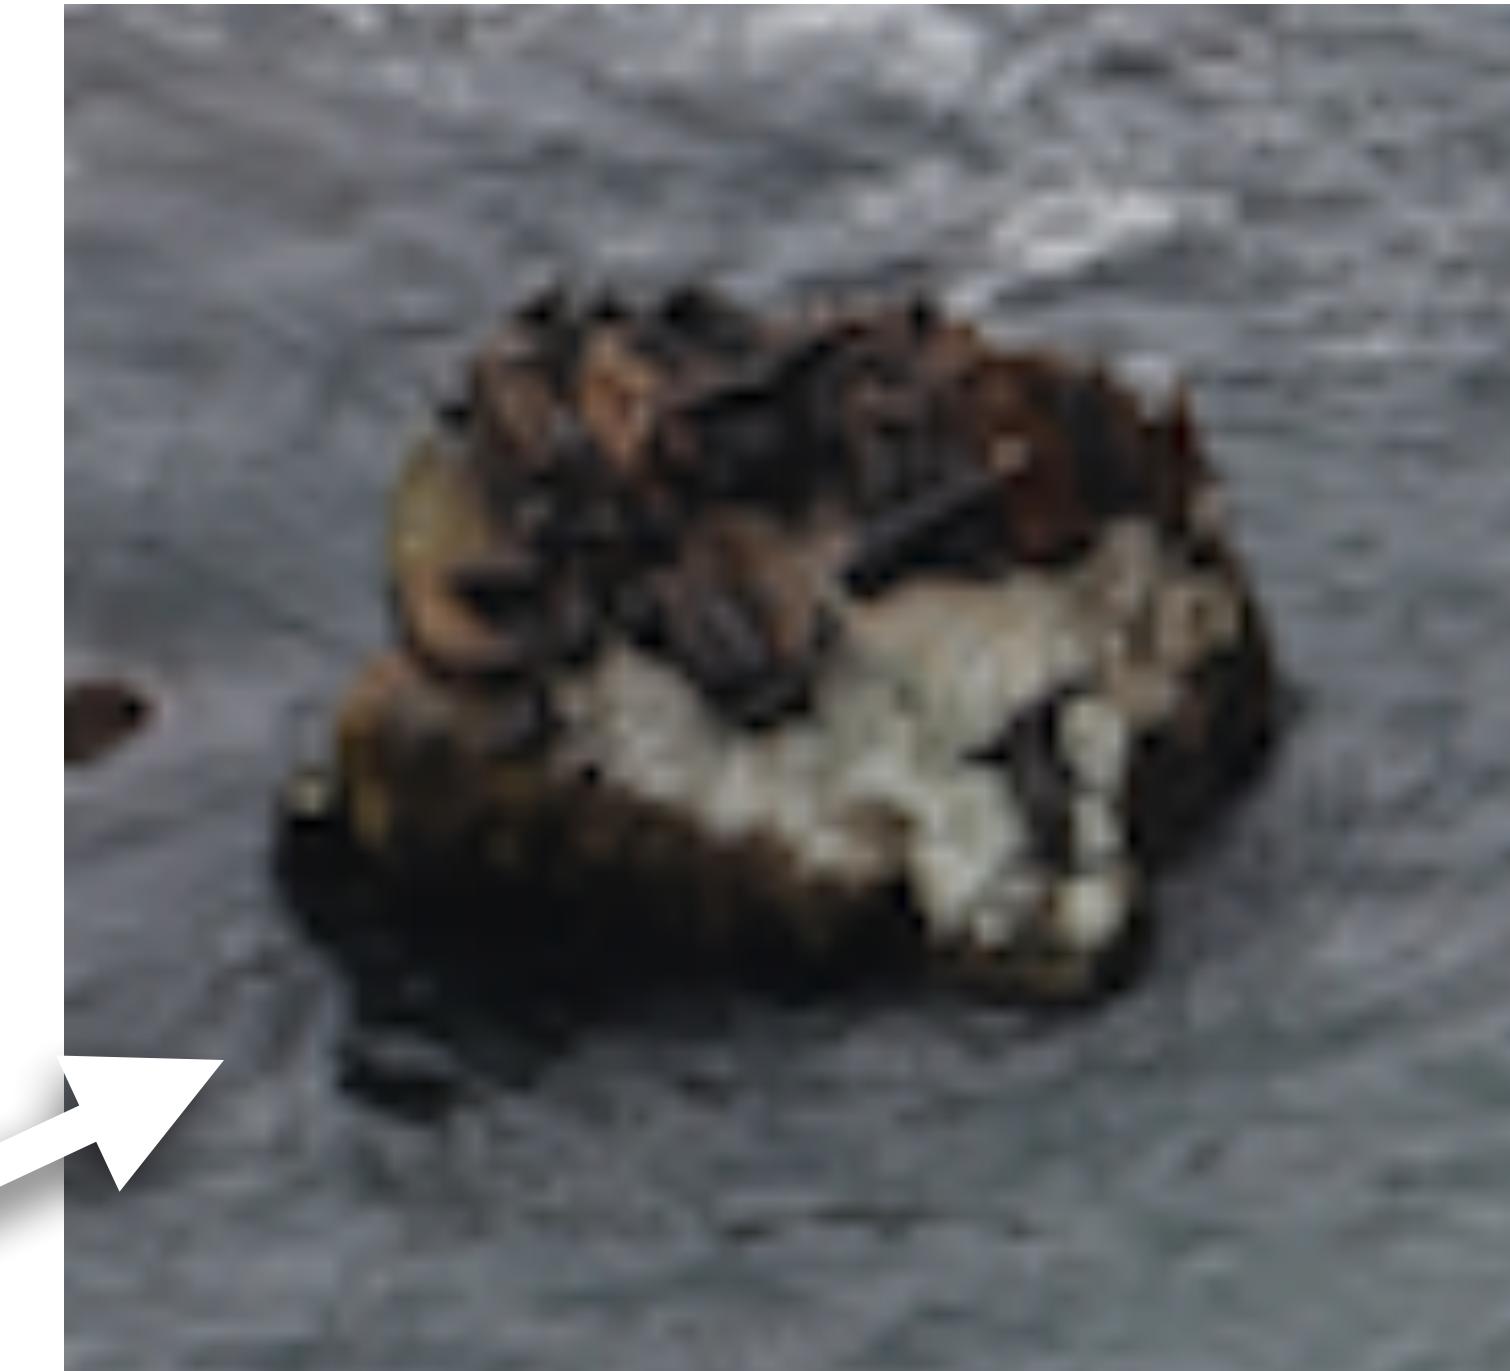
Camera shake



Blur kernel

Source: Torralba, Freeman, Isola

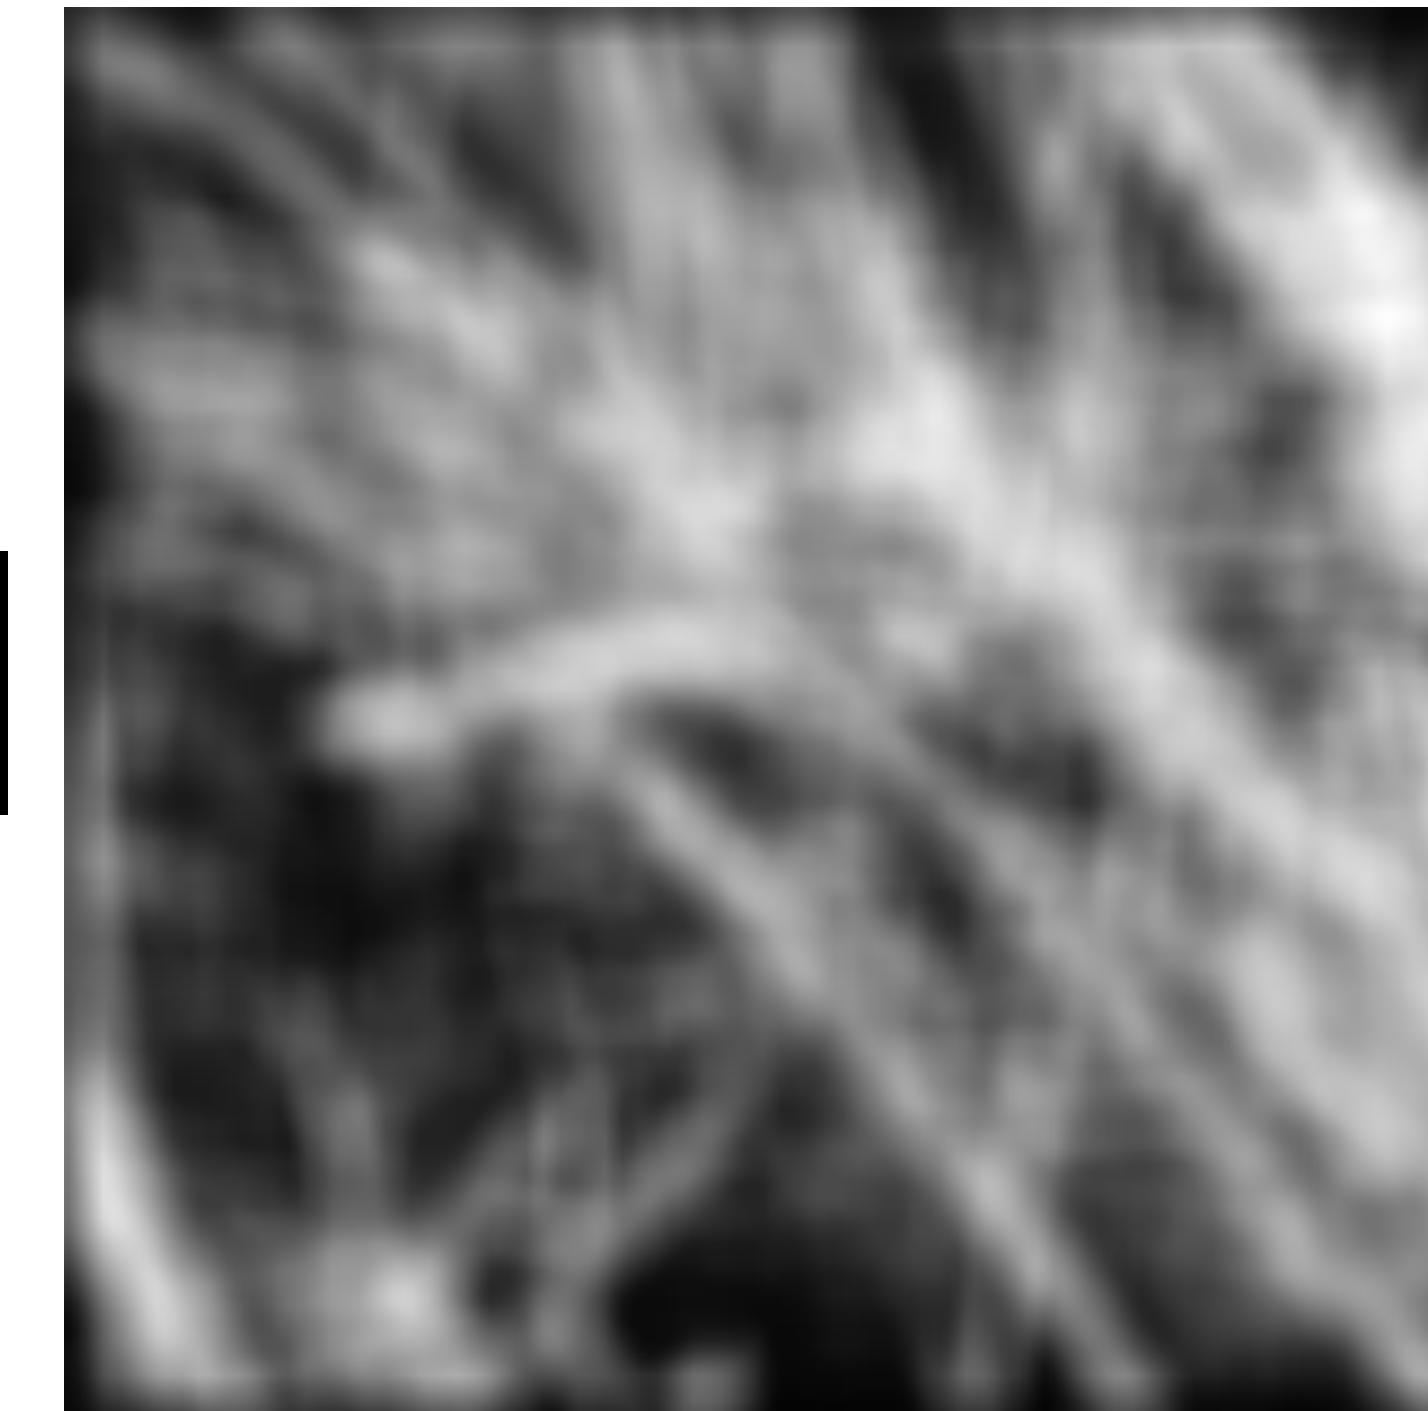
Blur occurs in many natural situations



Source: Torralba, Freeman, Isola

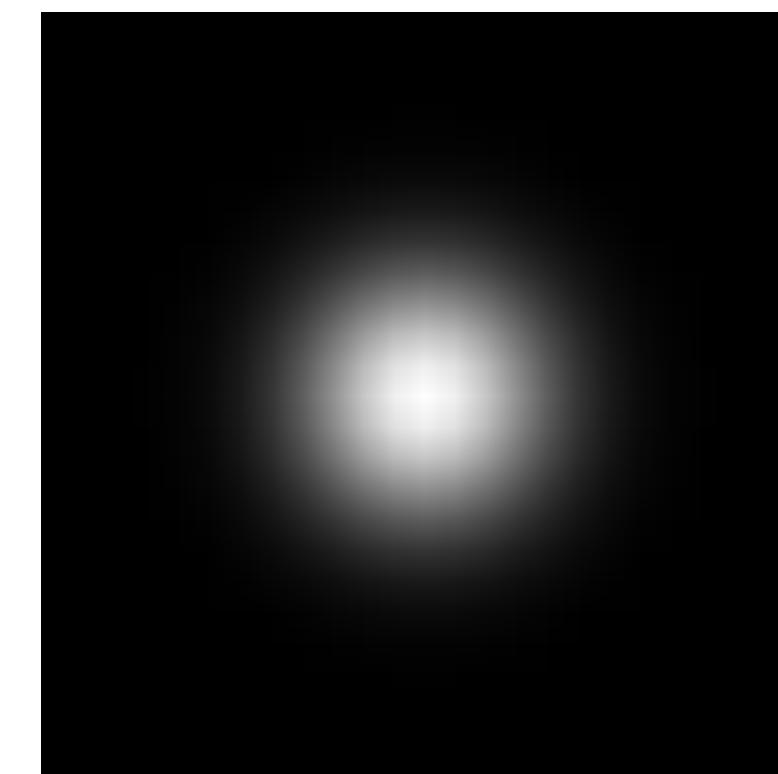
Smoothing with box filter revisited

- What's wrong with this picture?
- What's the solution?



Smoothing with box filter revisited

- What's wrong with this picture?
- What's the solution?
- To eliminate edge effects, weight contribution of neighborhood pixels according to their closeness to the center

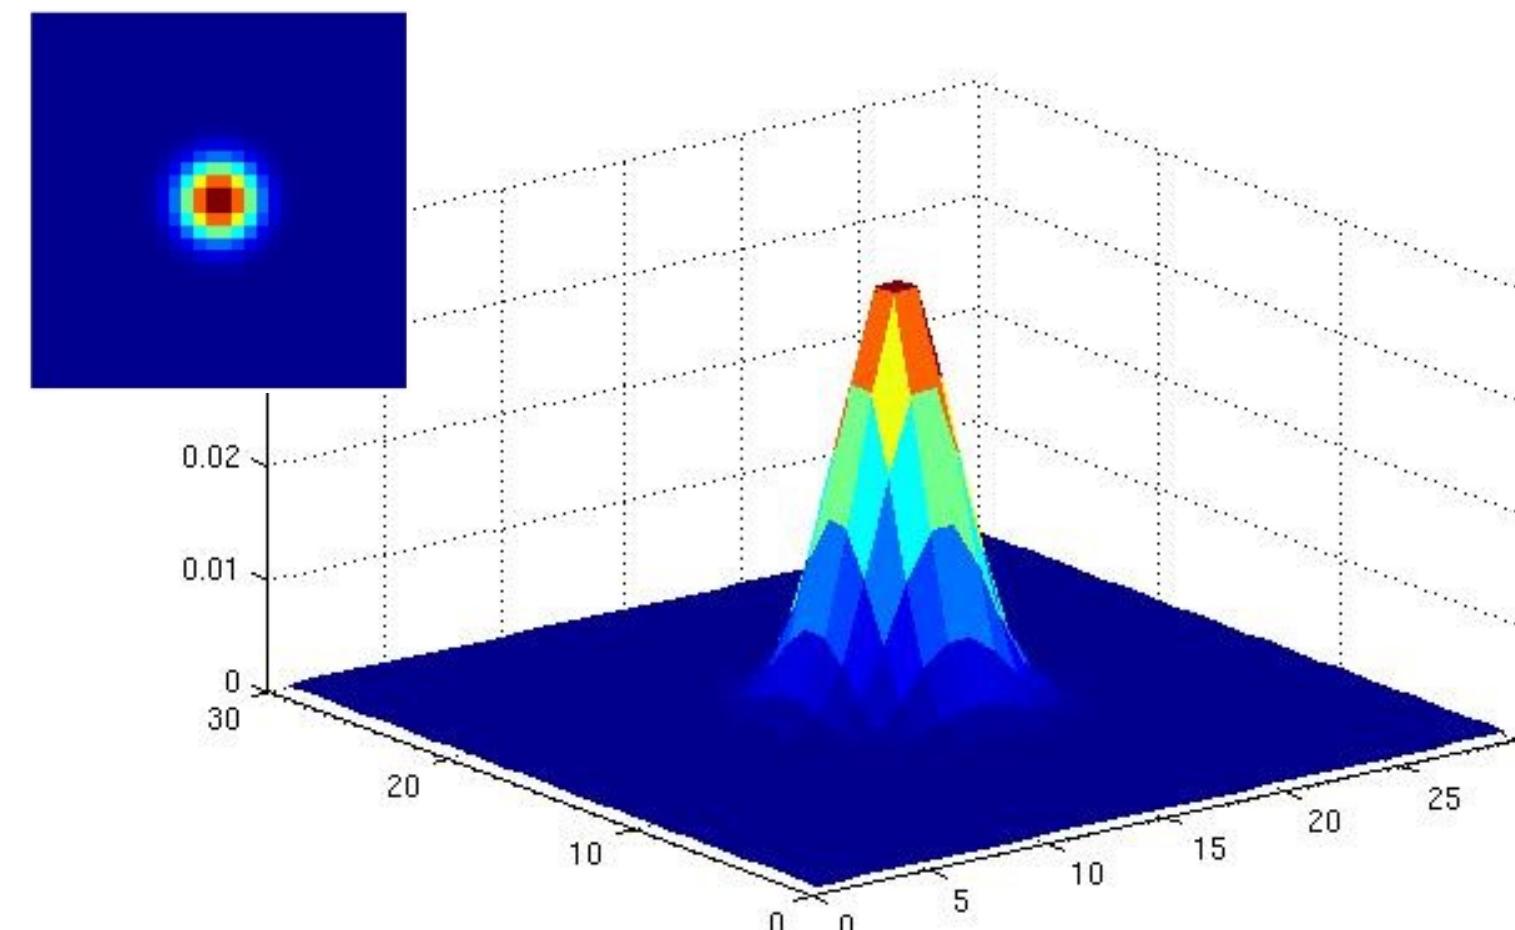


“fuzzy blob”

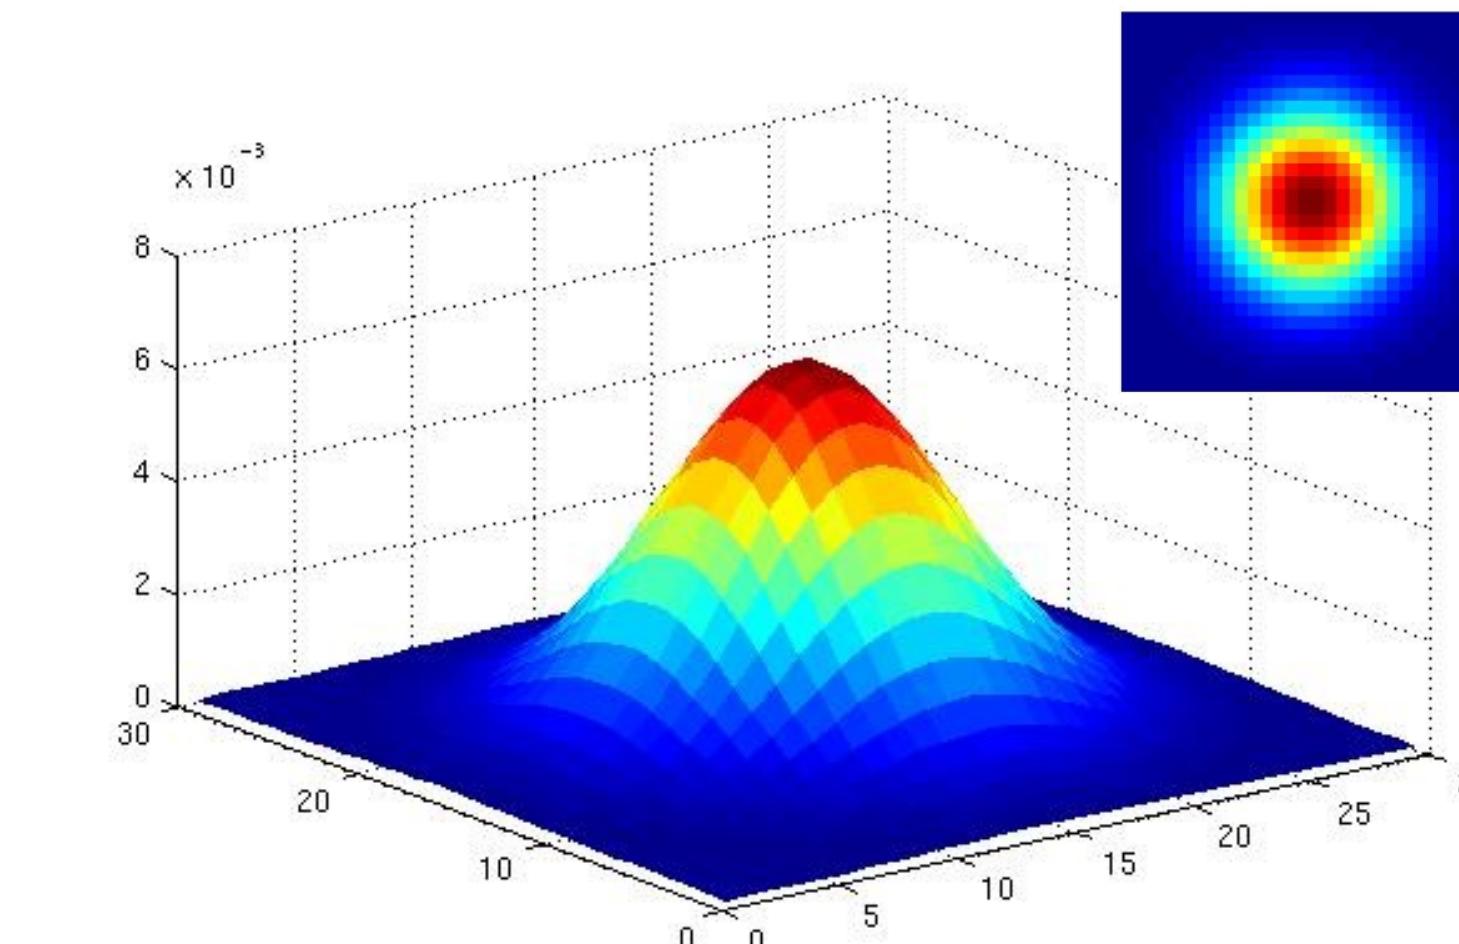
Source: S. Lazebnik

Gaussian kernel

$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$



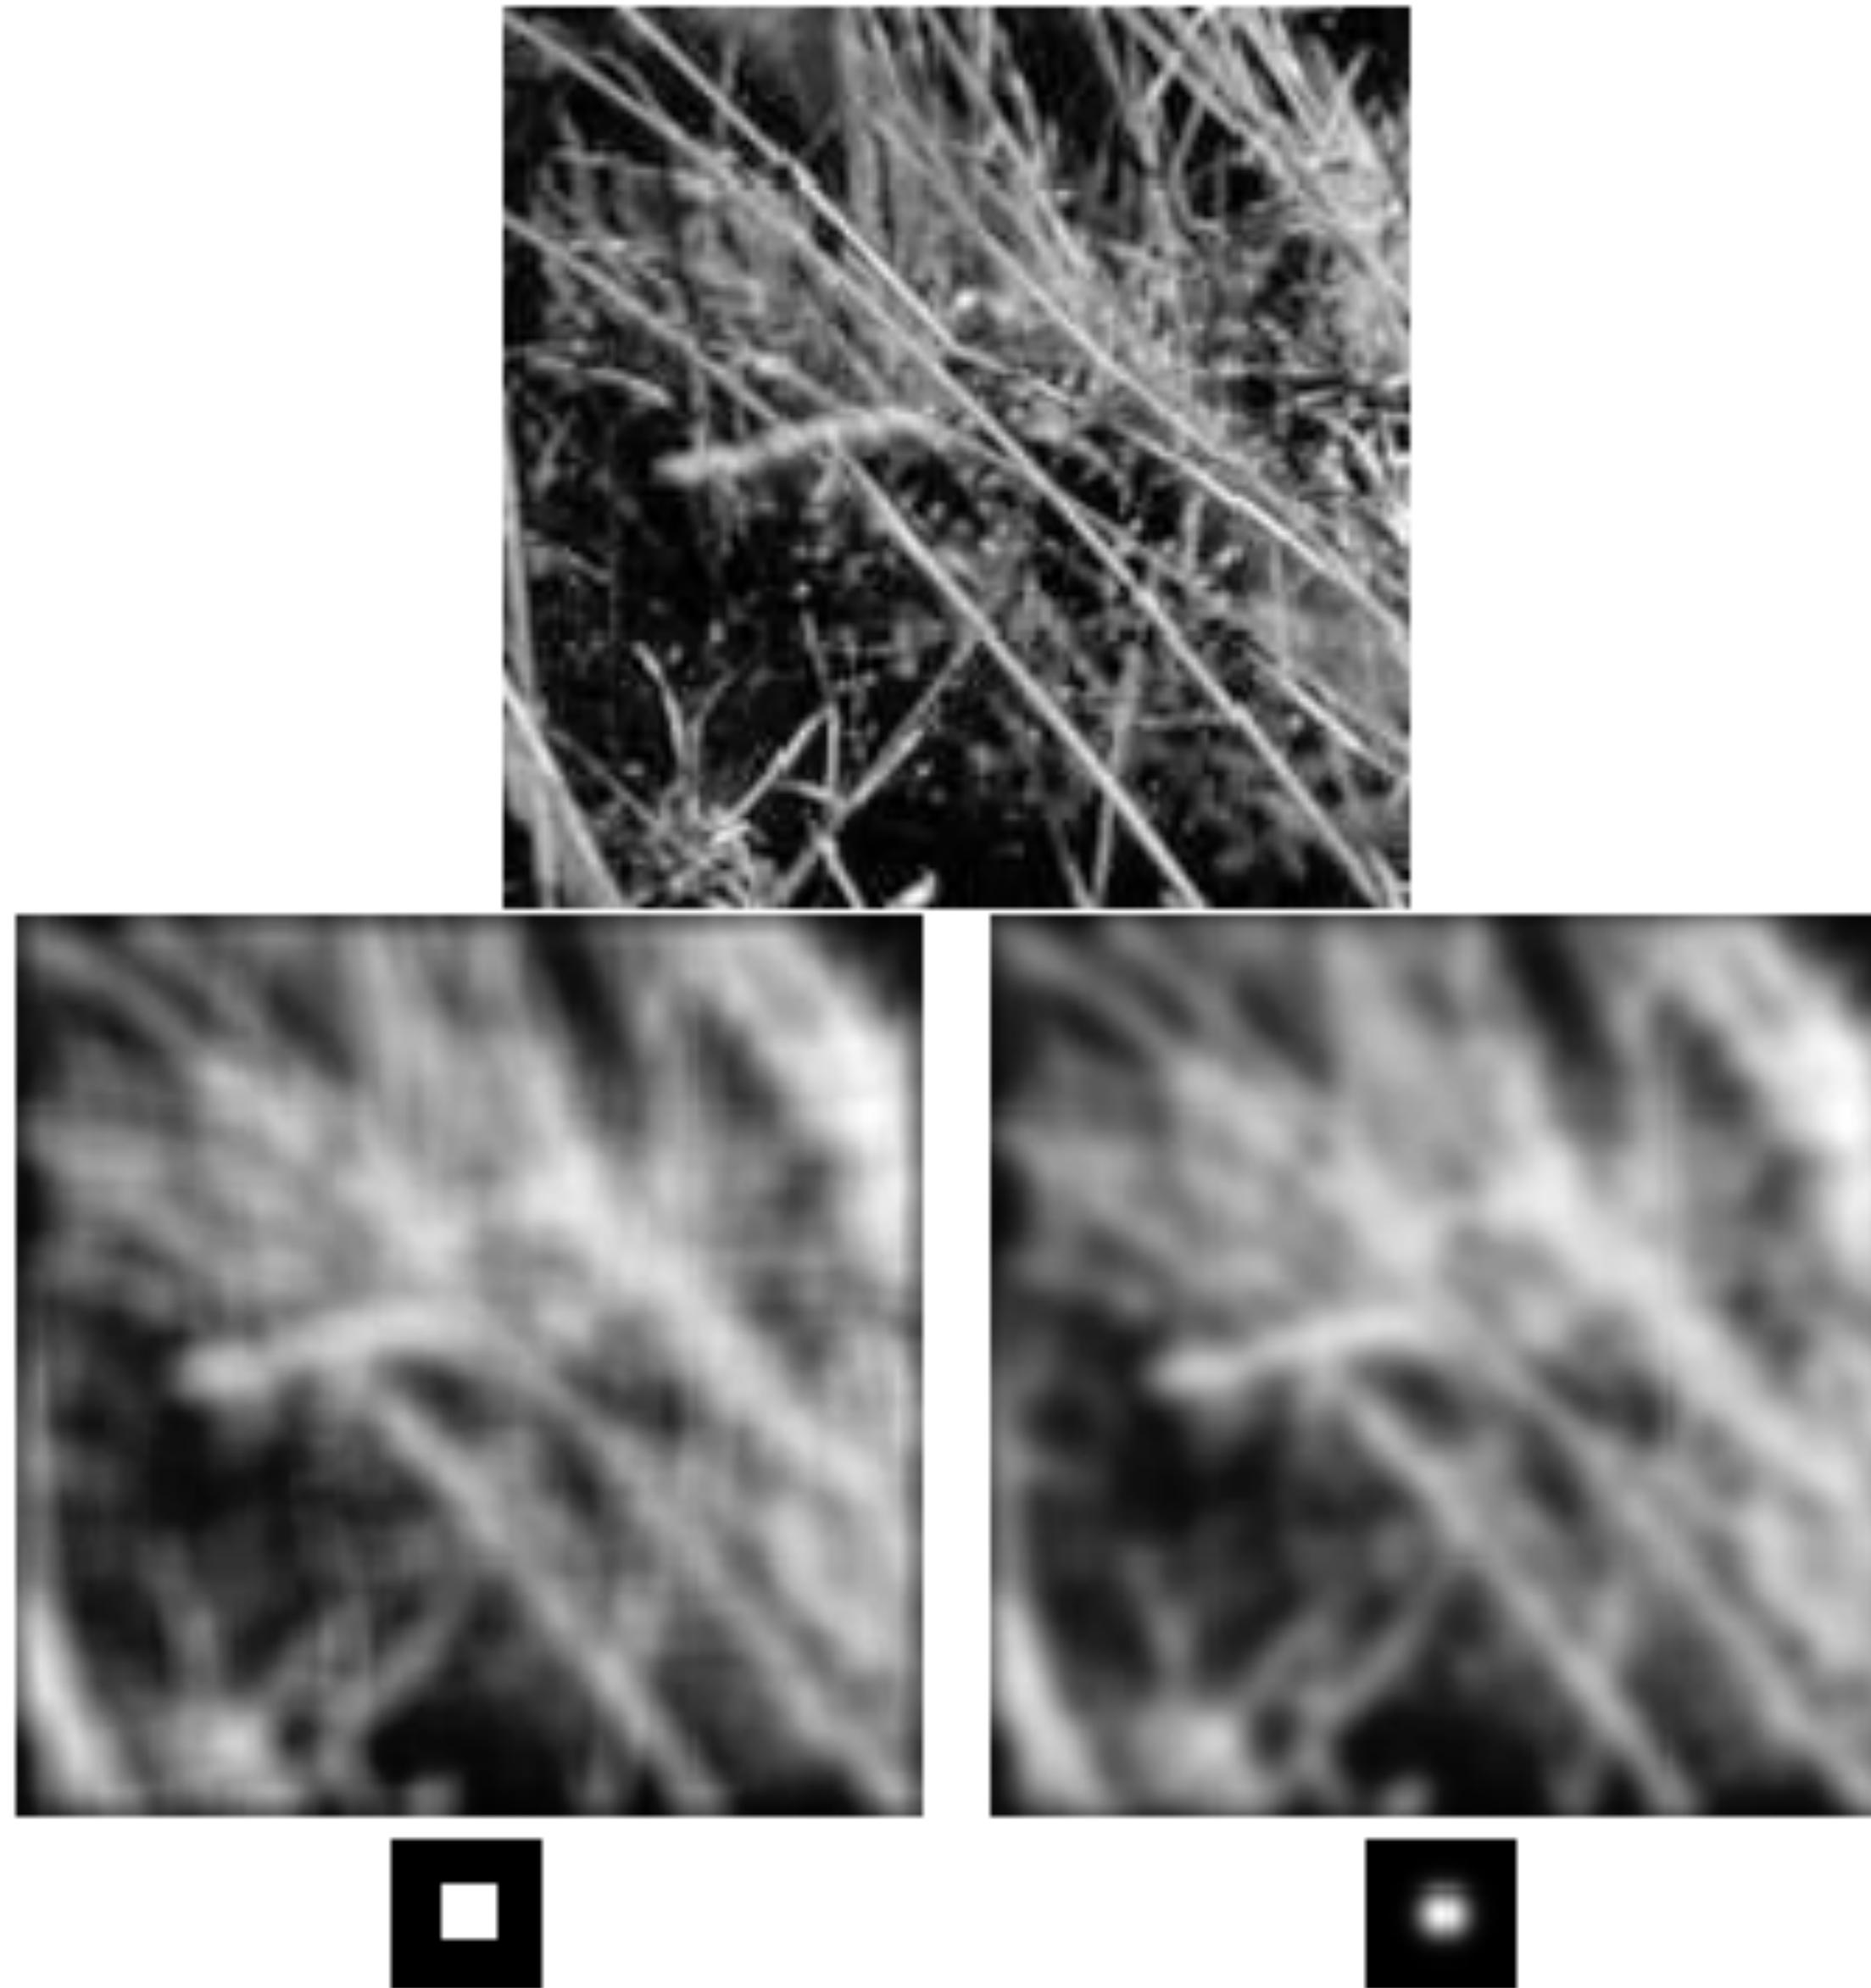
$\sigma = 2$ with 30×30 kernel



$\sigma = 5$ with 30×30 kernel

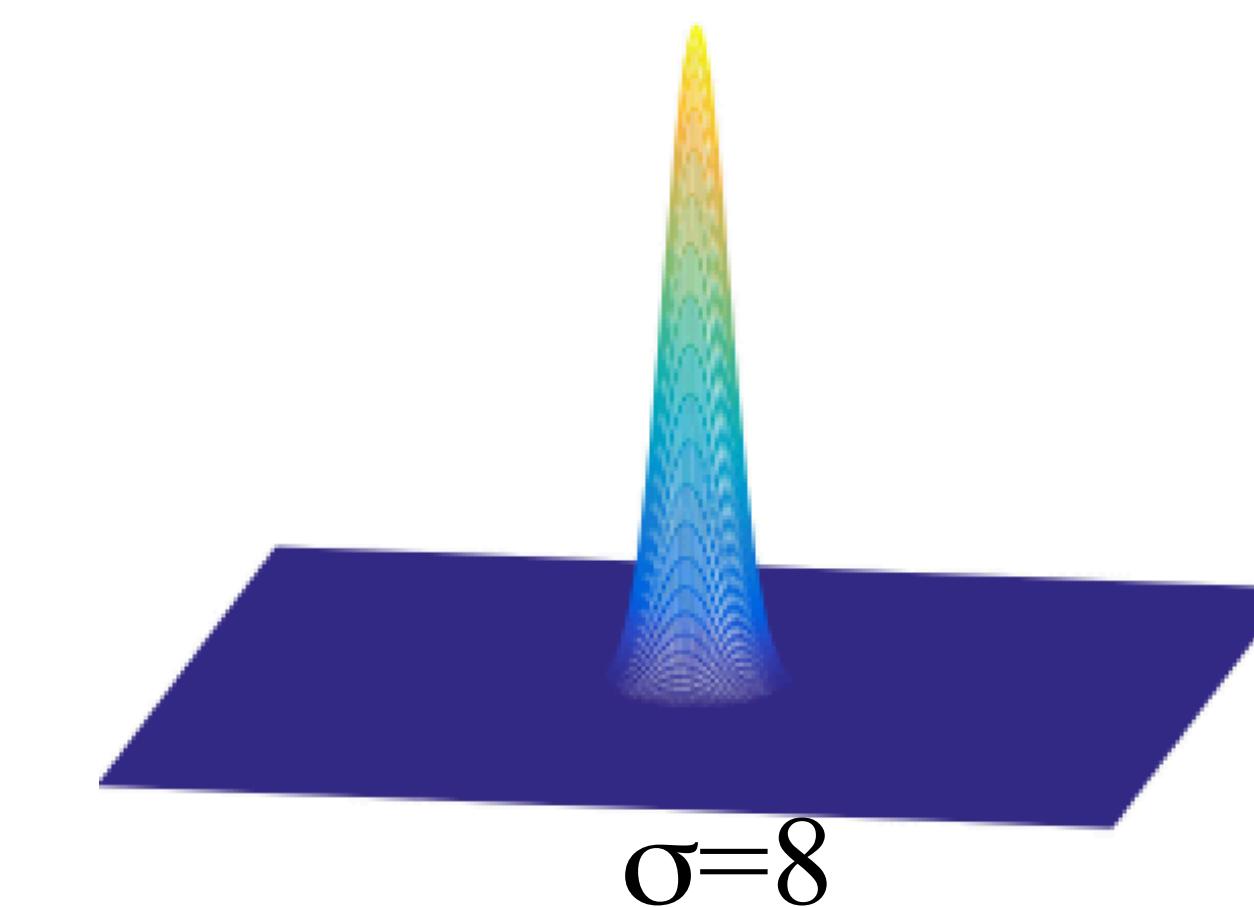
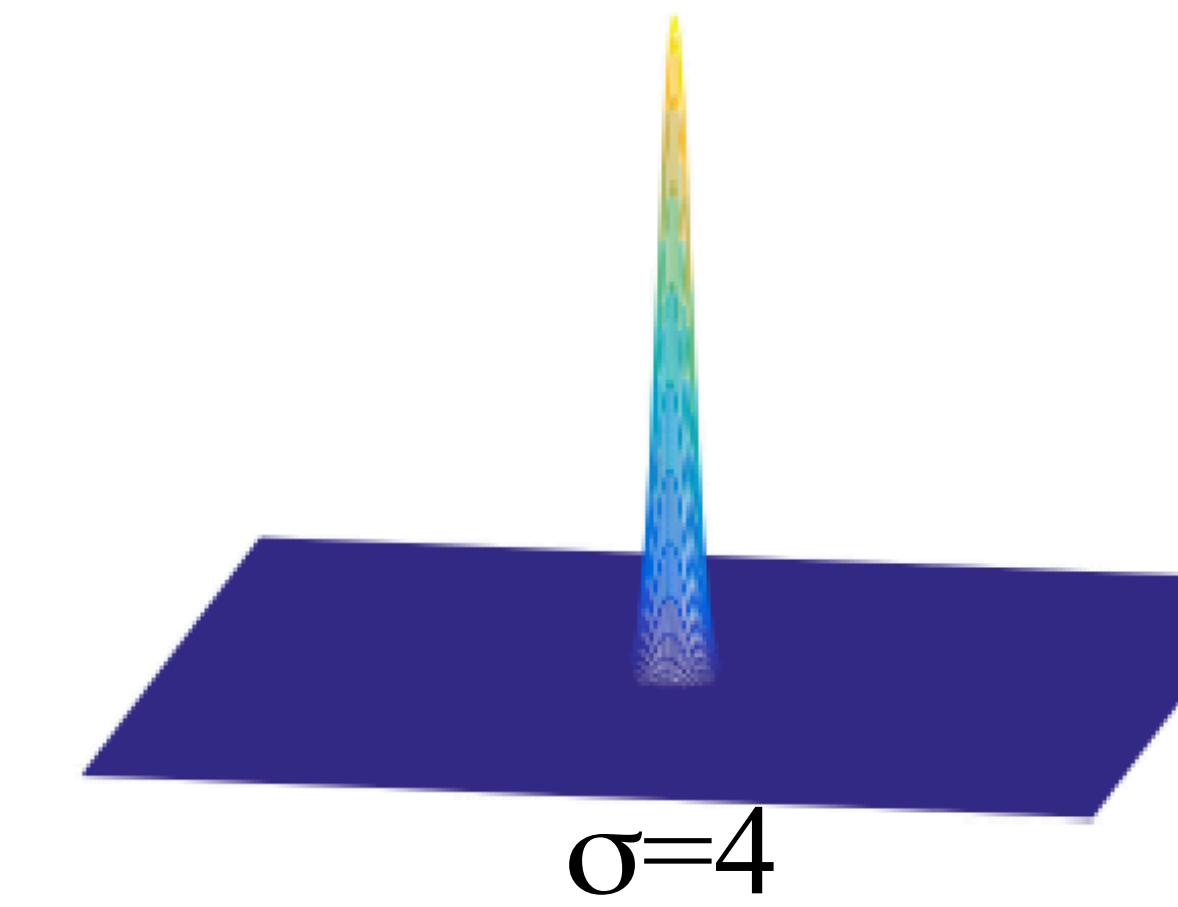
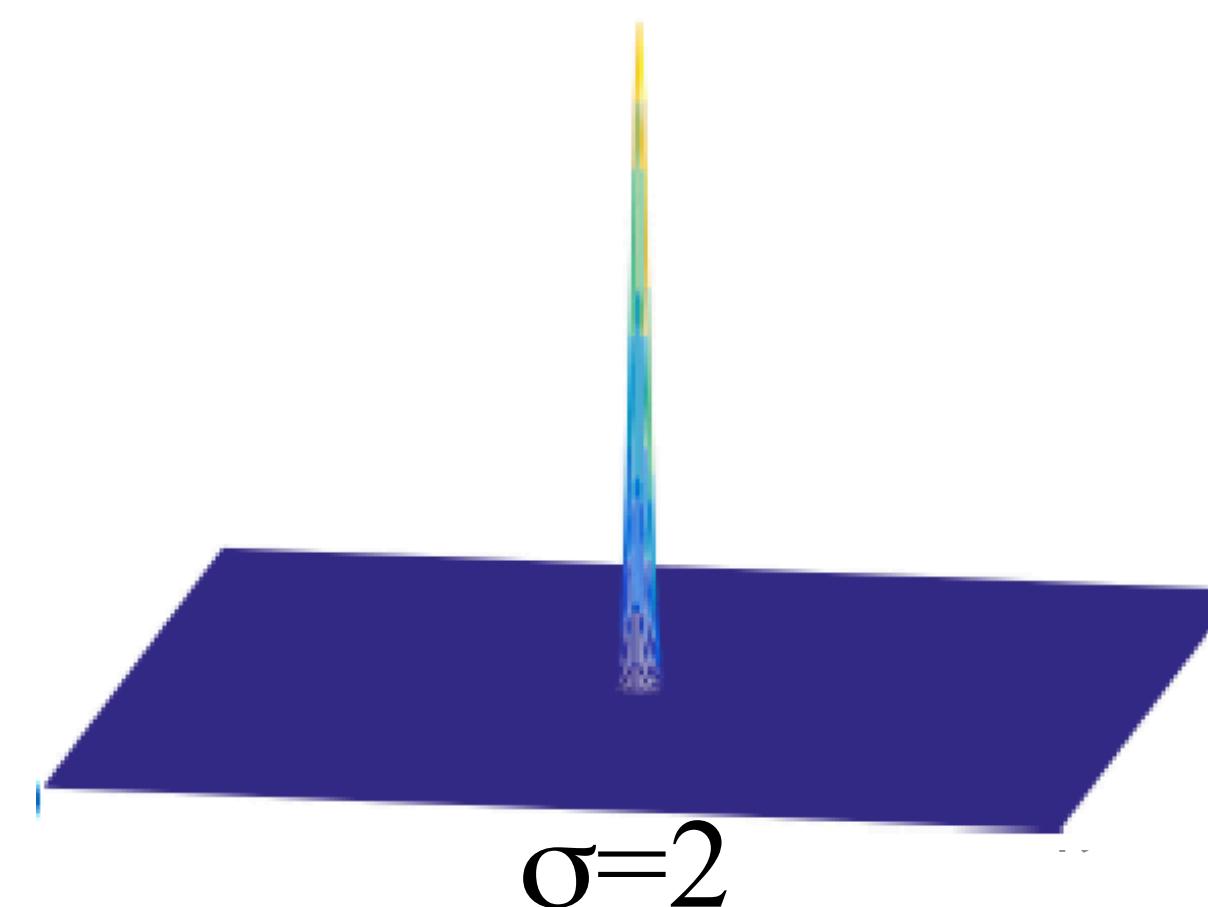
- Constant factor in front makes kernel sum to 1 (can also omit it and just divide by sum of filter weights).

Gaussian vs. box filtering



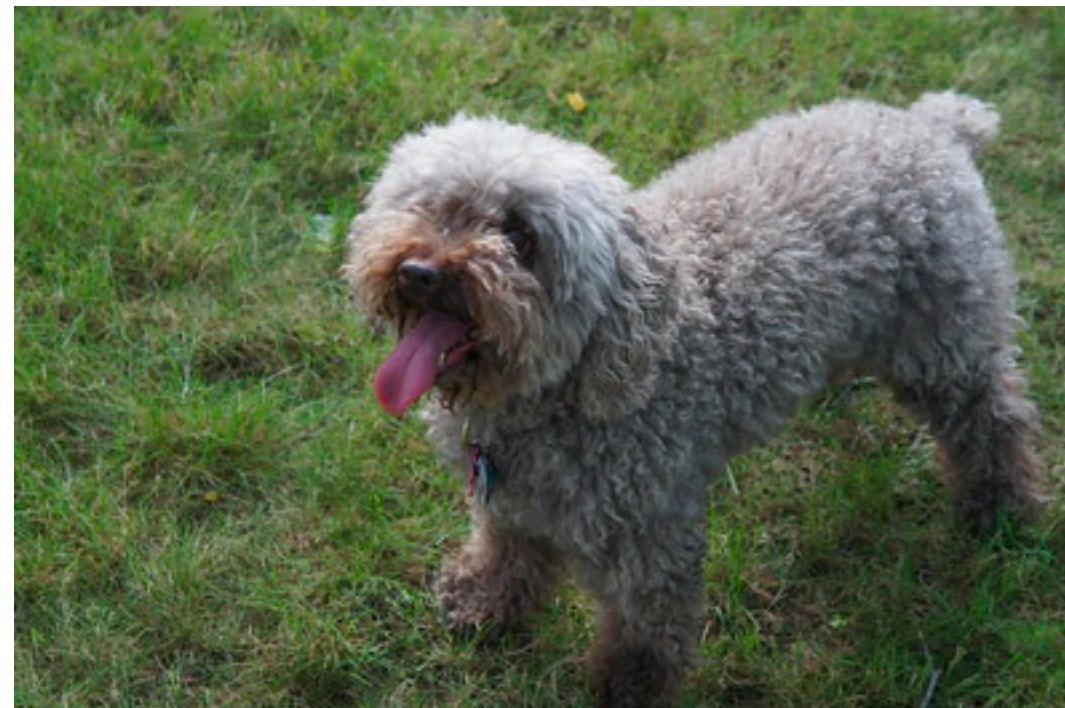
Source: S. Lazebnik

Gaussian standard deviation

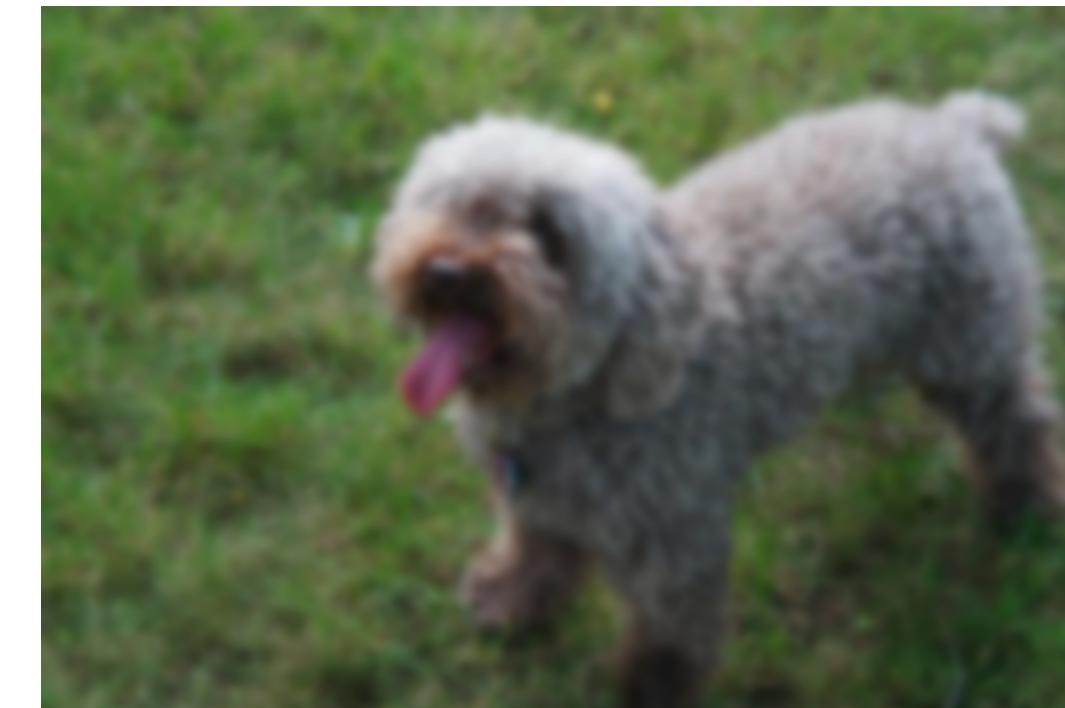


Gaussian filters

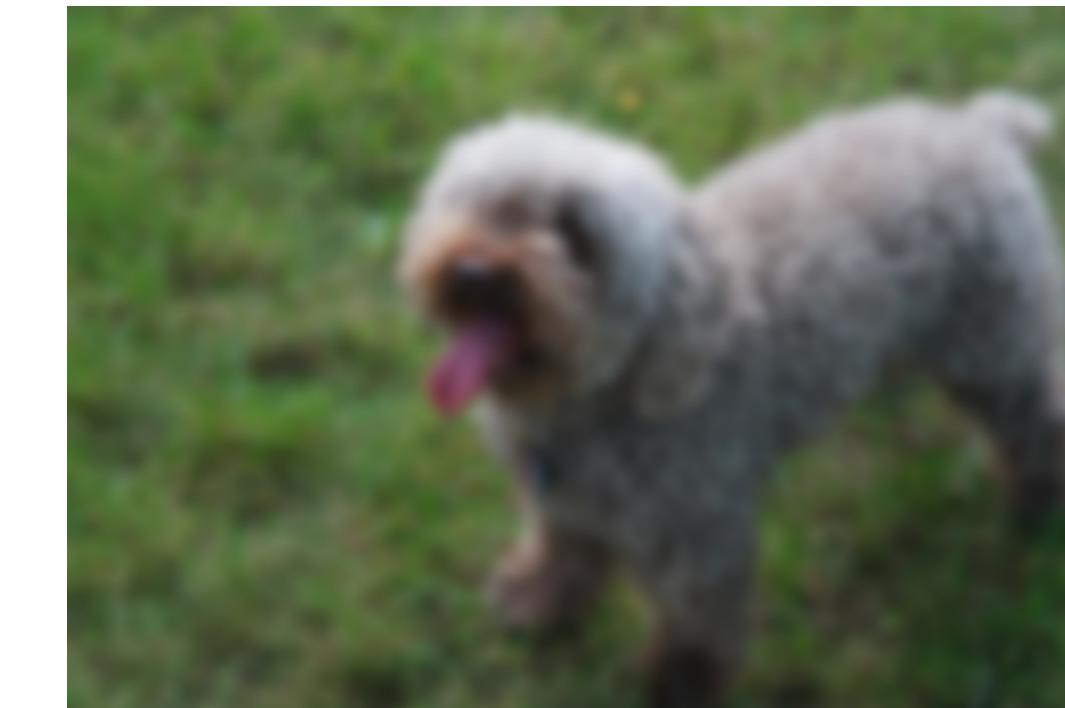
- Convolution with self is another Gaussian
 - Can smooth with small- σ kernel, repeat, get same result as larger- σ
 - Convolving two times with Gaussian kernel with std. dev. σ is same as convolving once with kernel with std. dev. $\sigma\sqrt{2}$



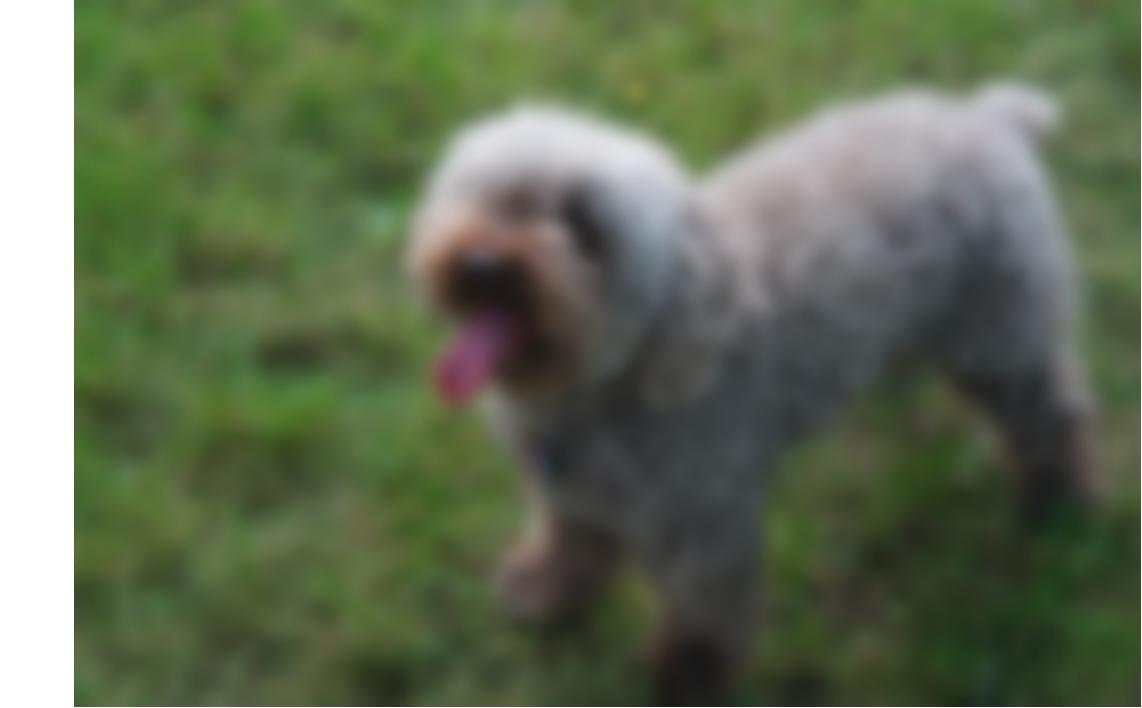
I



blur(blur(I)))



blur(blur(blur(blur(I)))))



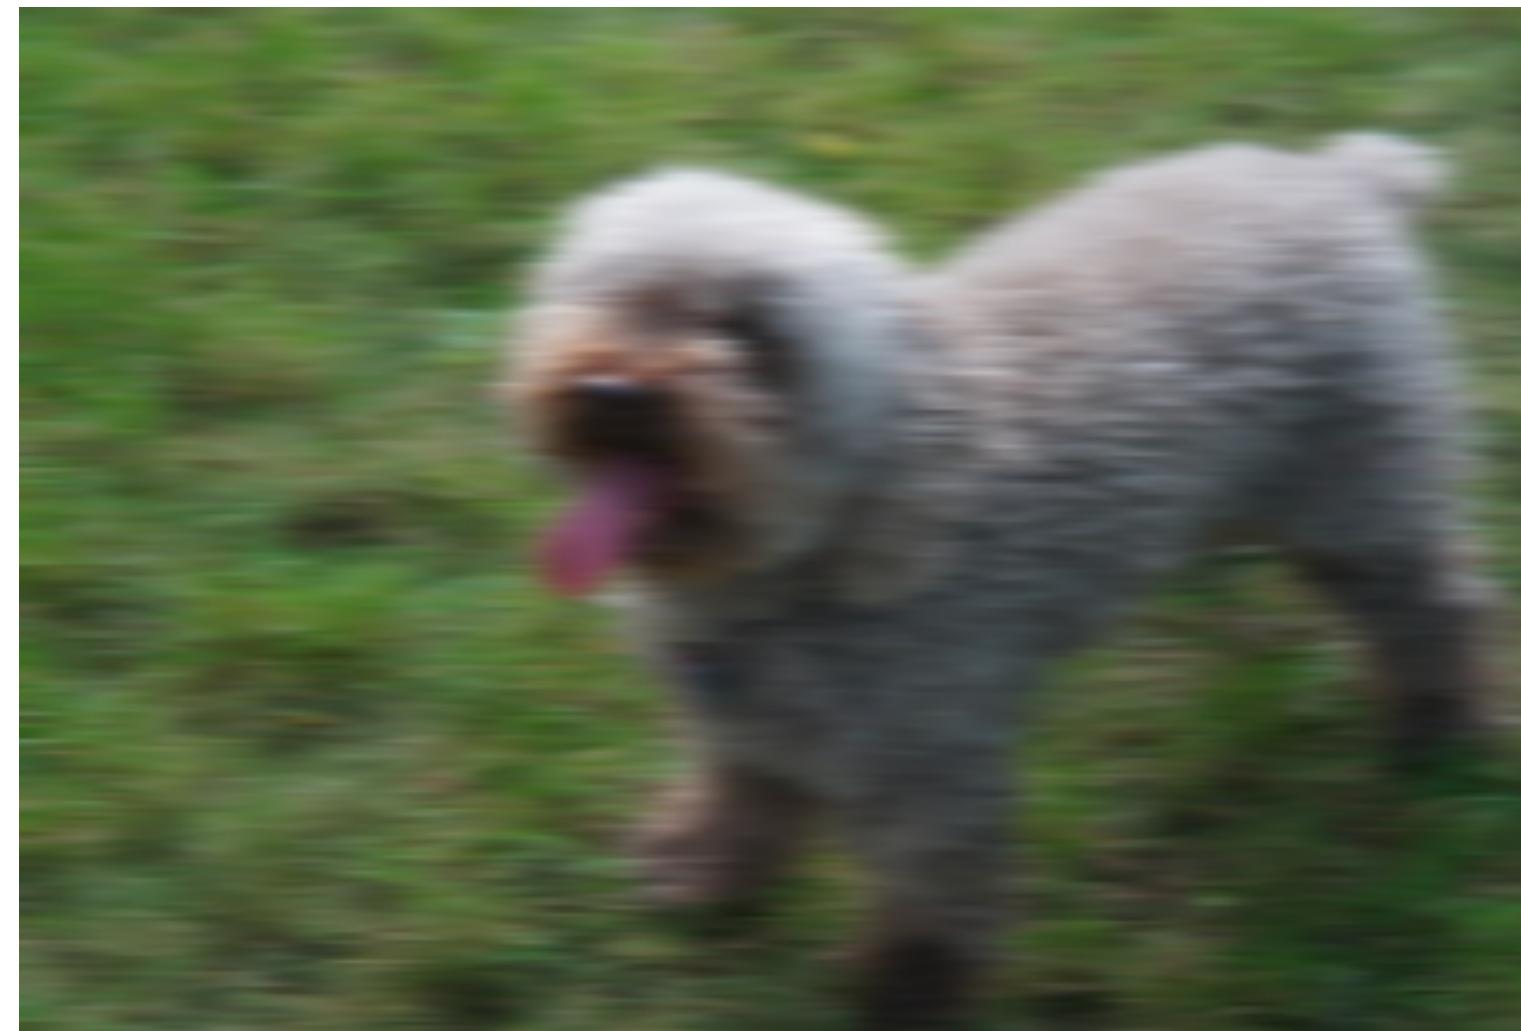
blur(blur(blur(blur(blur(blur(I)))))))

Gaussian filters

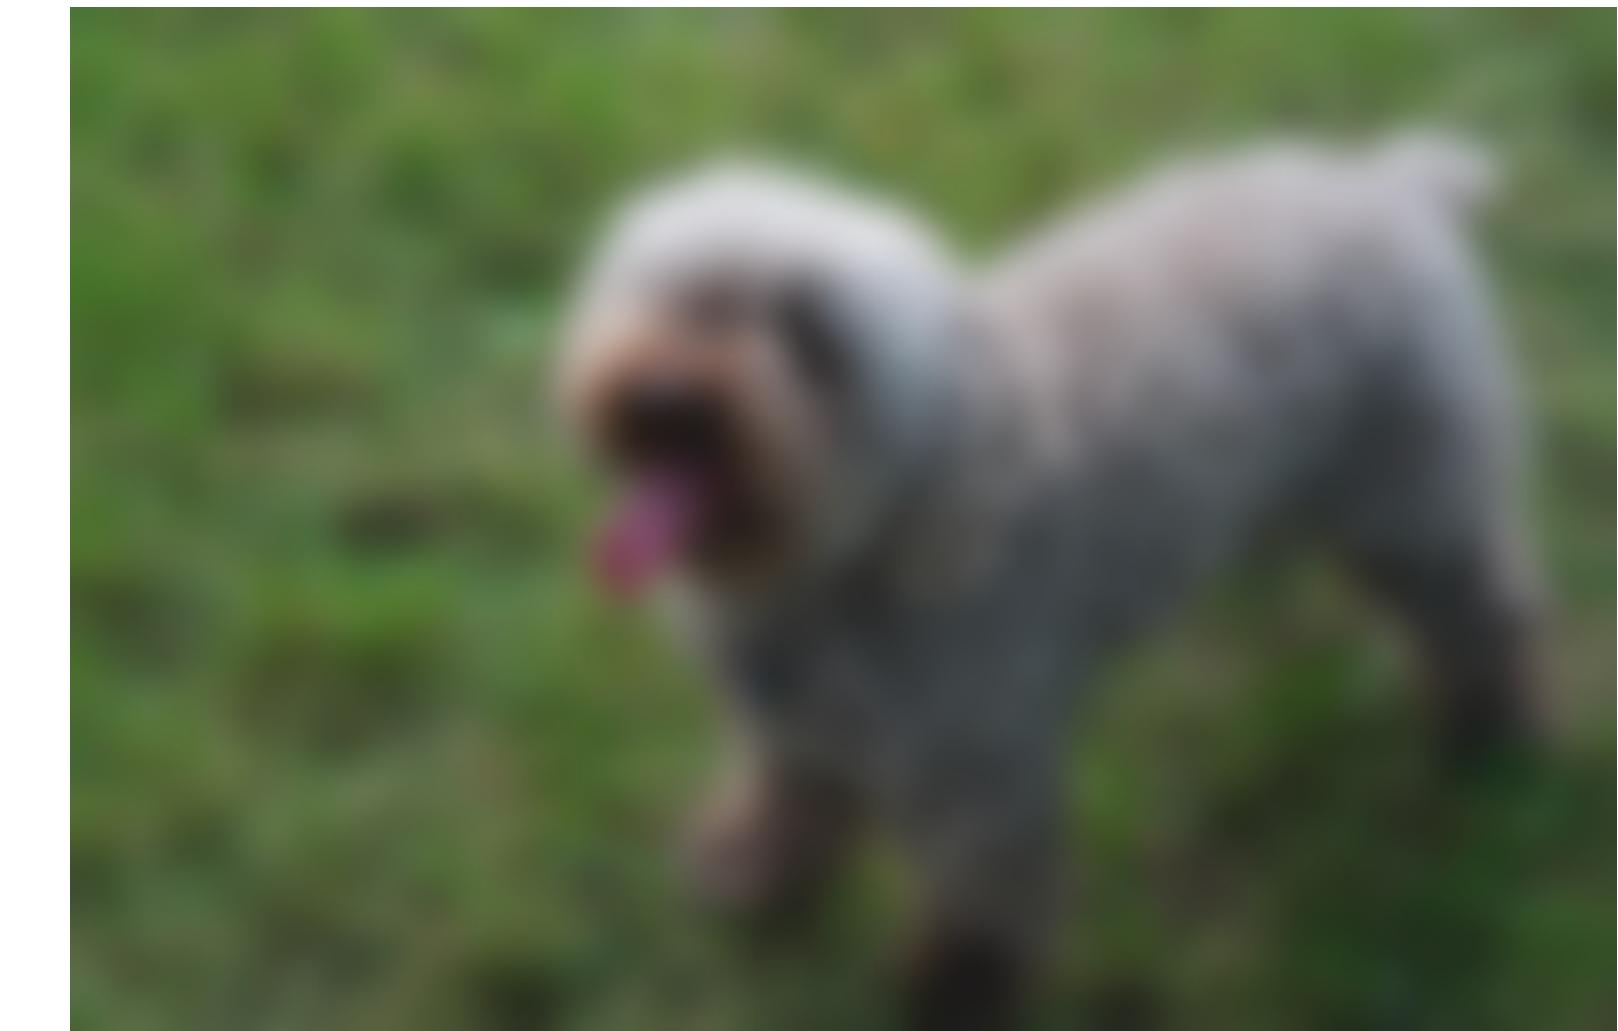
- It's a **separable kernel**
 - Blur with 1D Gaussian in one direction, then the other.
 - Faster to compute. $O(n)$ time for an n^*n kernel instead of $O(n^2)$
 - Learn more about this in Problem Set 1!



I



$\text{blur}_x(I)$



$\text{blur}_y(\text{blur}_x(I))$

Edges: recall last lecture...

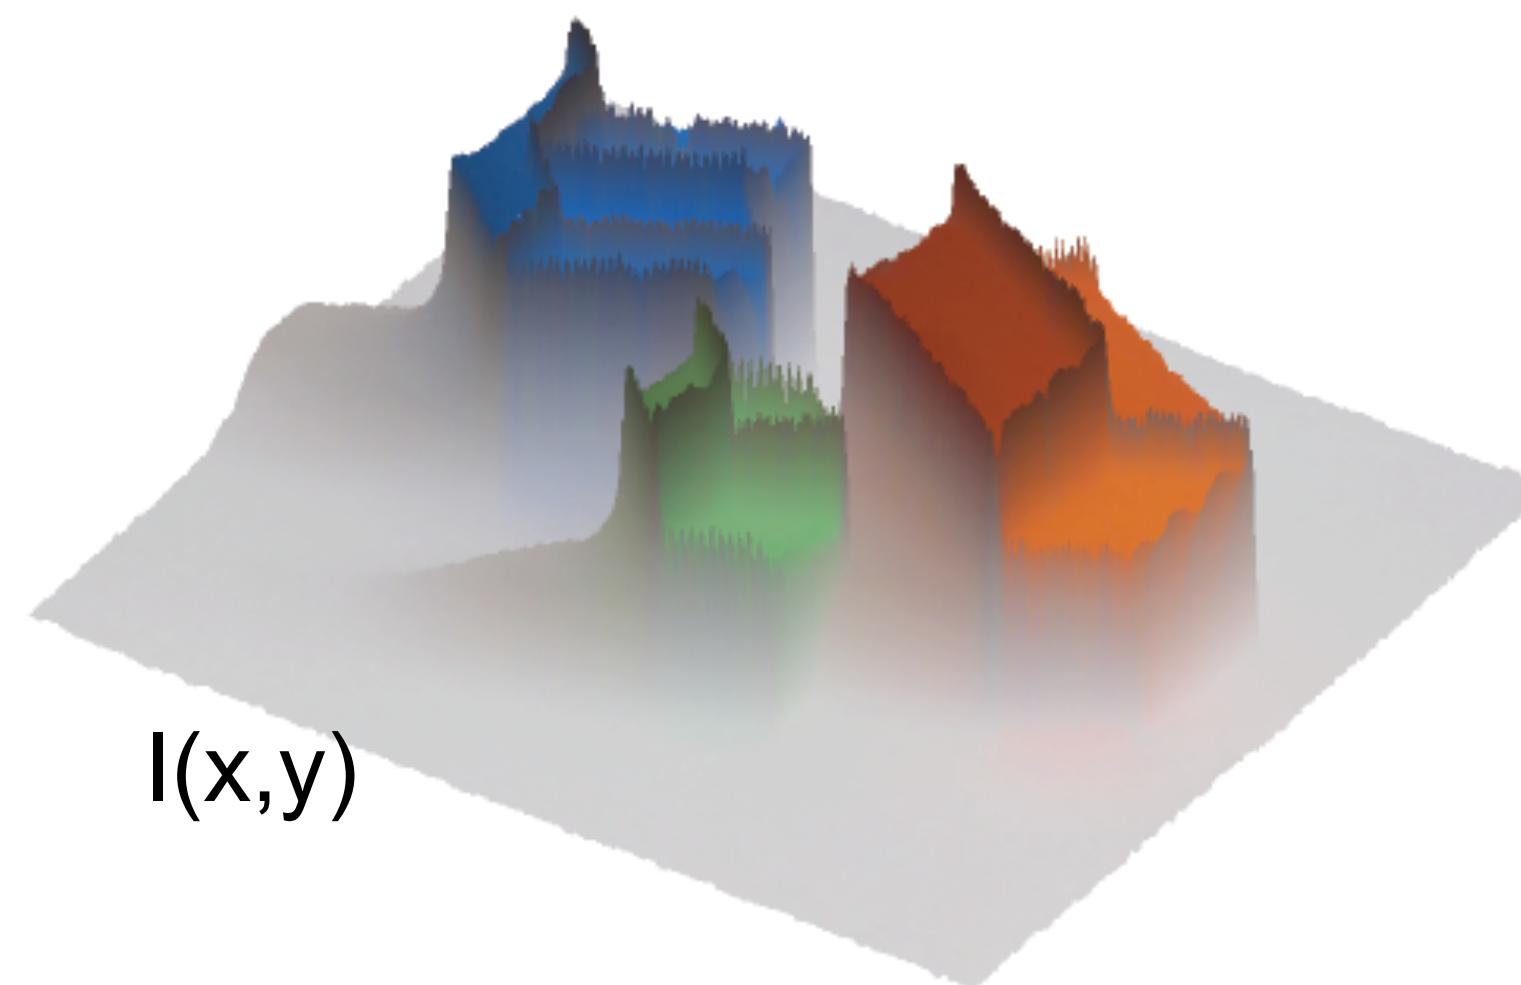


Image gradient:

$$\nabla \mathbf{I} = \left(\frac{\partial \mathbf{I}}{\partial x}, \frac{\partial \mathbf{I}}{\partial y} \right)$$

Approximation image derivative:

$$\frac{\partial \mathbf{I}}{\partial x} \simeq \mathbf{I}(x, y) - \mathbf{I}(x - 1, y)$$

Edge strength

$$E(x, y) = |\nabla \mathbf{I}(x, y)|$$

Edge orientation:

$$\theta(x, y) = \angle \nabla \mathbf{I} = \arctan \frac{\frac{\partial \mathbf{I}}{\partial y}}{\frac{\partial \mathbf{I}}{\partial x}}$$

Edge normal:

$$\mathbf{n} = \frac{\nabla \mathbf{I}}{|\nabla \mathbf{I}|}$$

Discrete derivatives

$$d_0 = [1, -1]$$

$$f \circ d_0 = f[n] - f[n - 1]$$

$$d_1 = [1, 0, -1]/2$$

$$f \circ d_1 = \frac{f[n + 1] - f[n - 1]}{2}$$

$$[-1 \ 1]$$



$g[m,n]$

$\circ [-1, 1] =$

$h[m,n]$



$f[m,n]$

Source: Torralba, Freeman, Isola

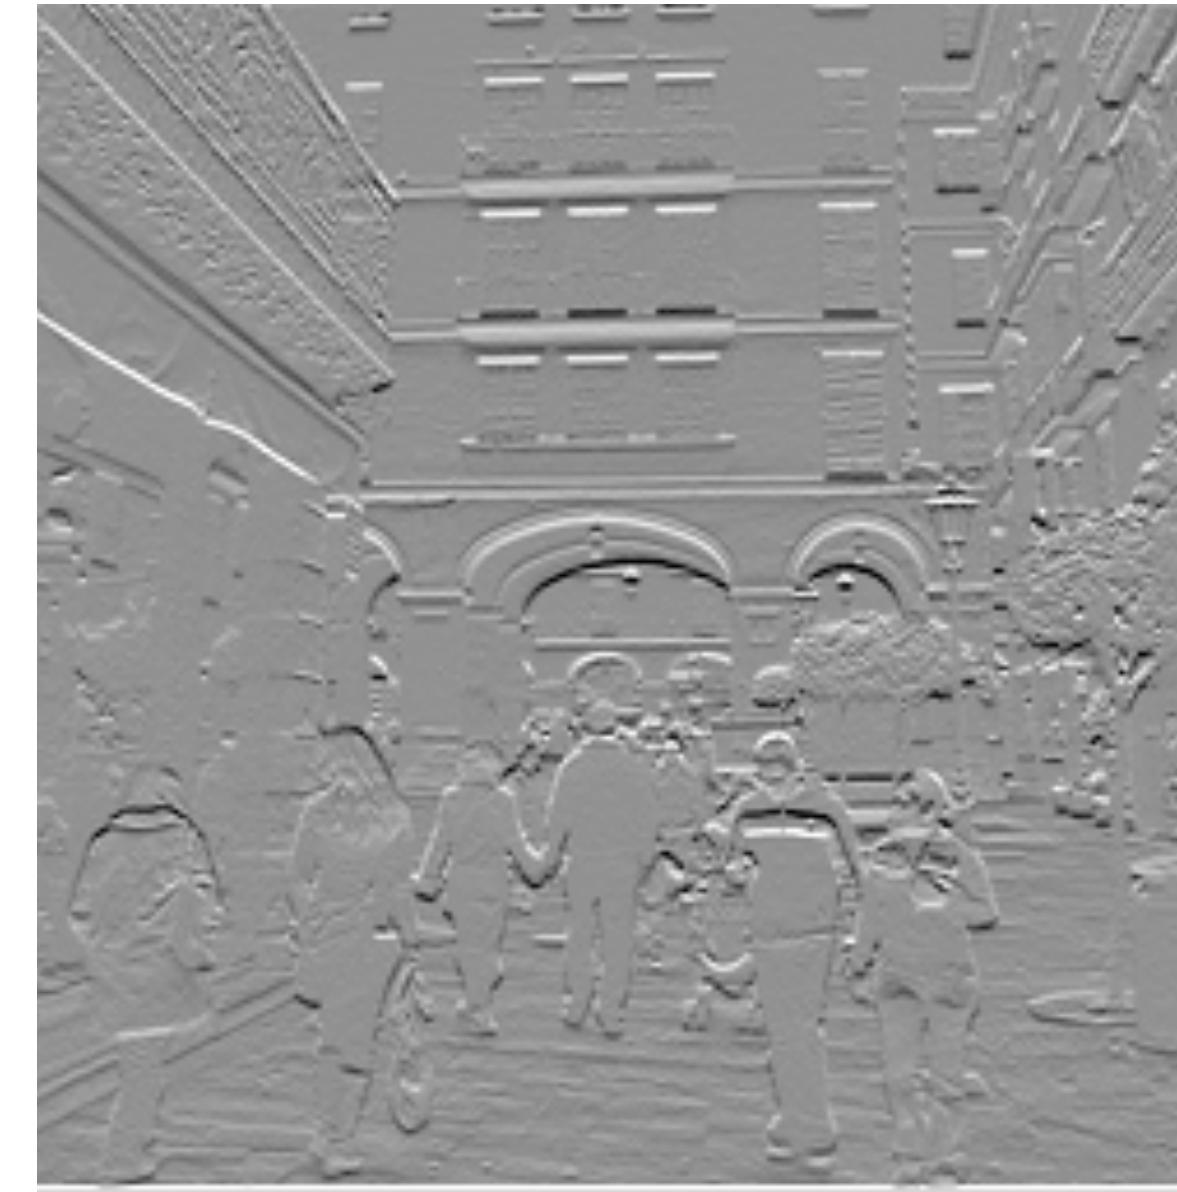
$$[-1 \ 1]^\top$$



$g[m,n]$

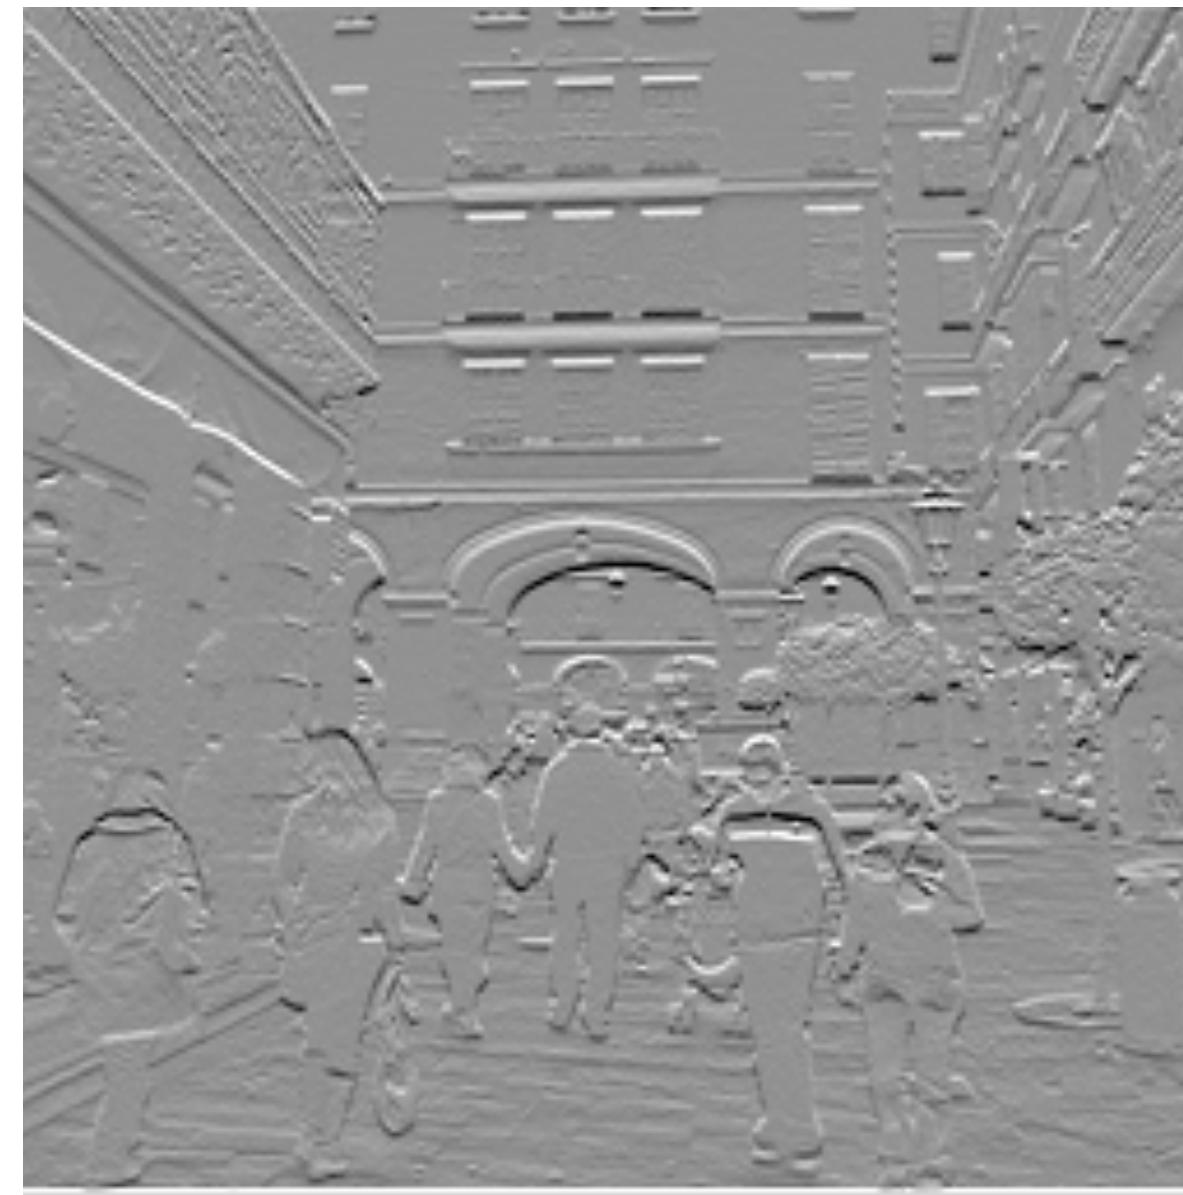
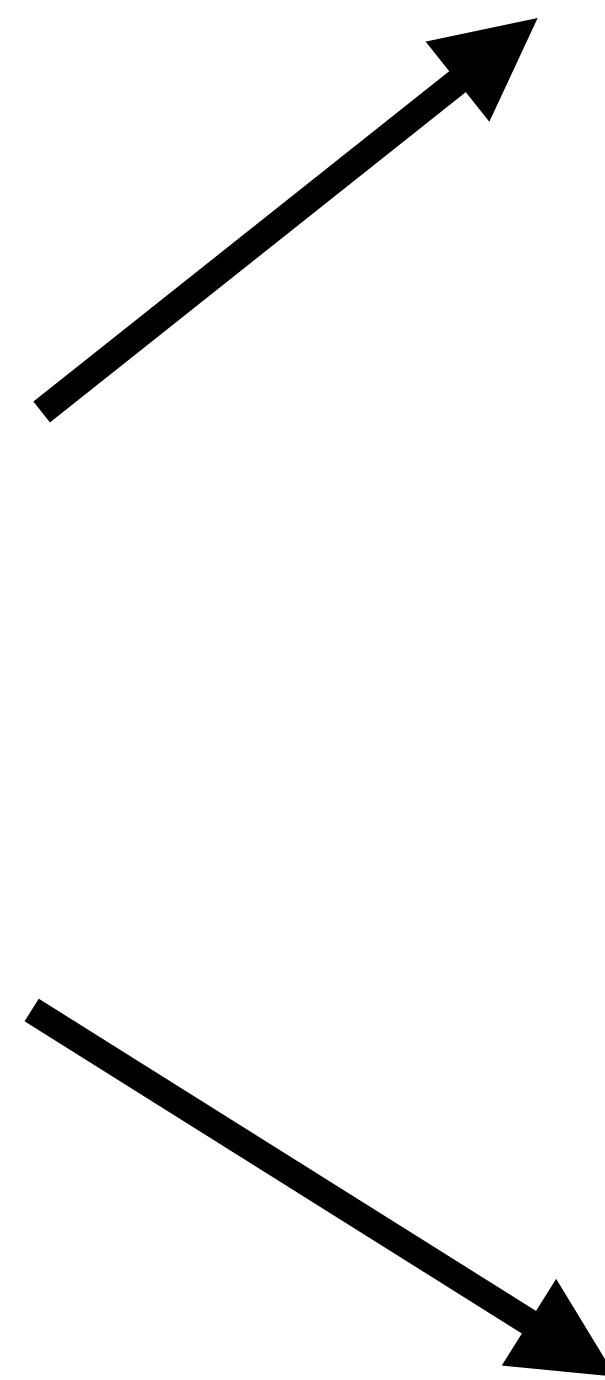
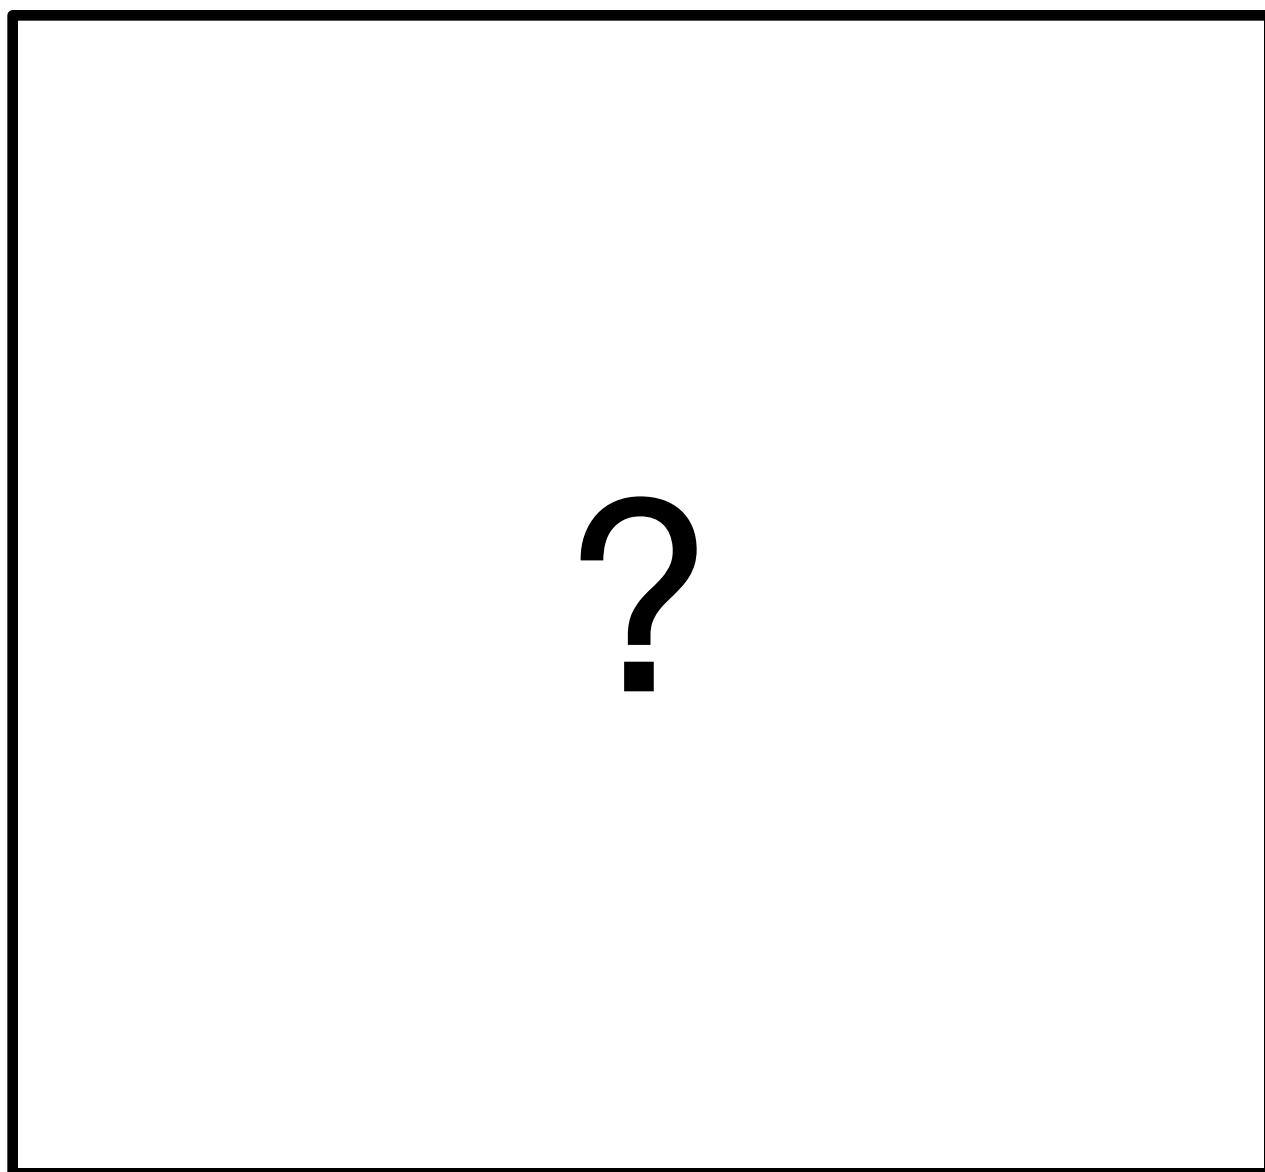
$$\circ [-1, 1]^\top =$$

$h[m,n]$



$f[m,n]$

Can we recover the image?



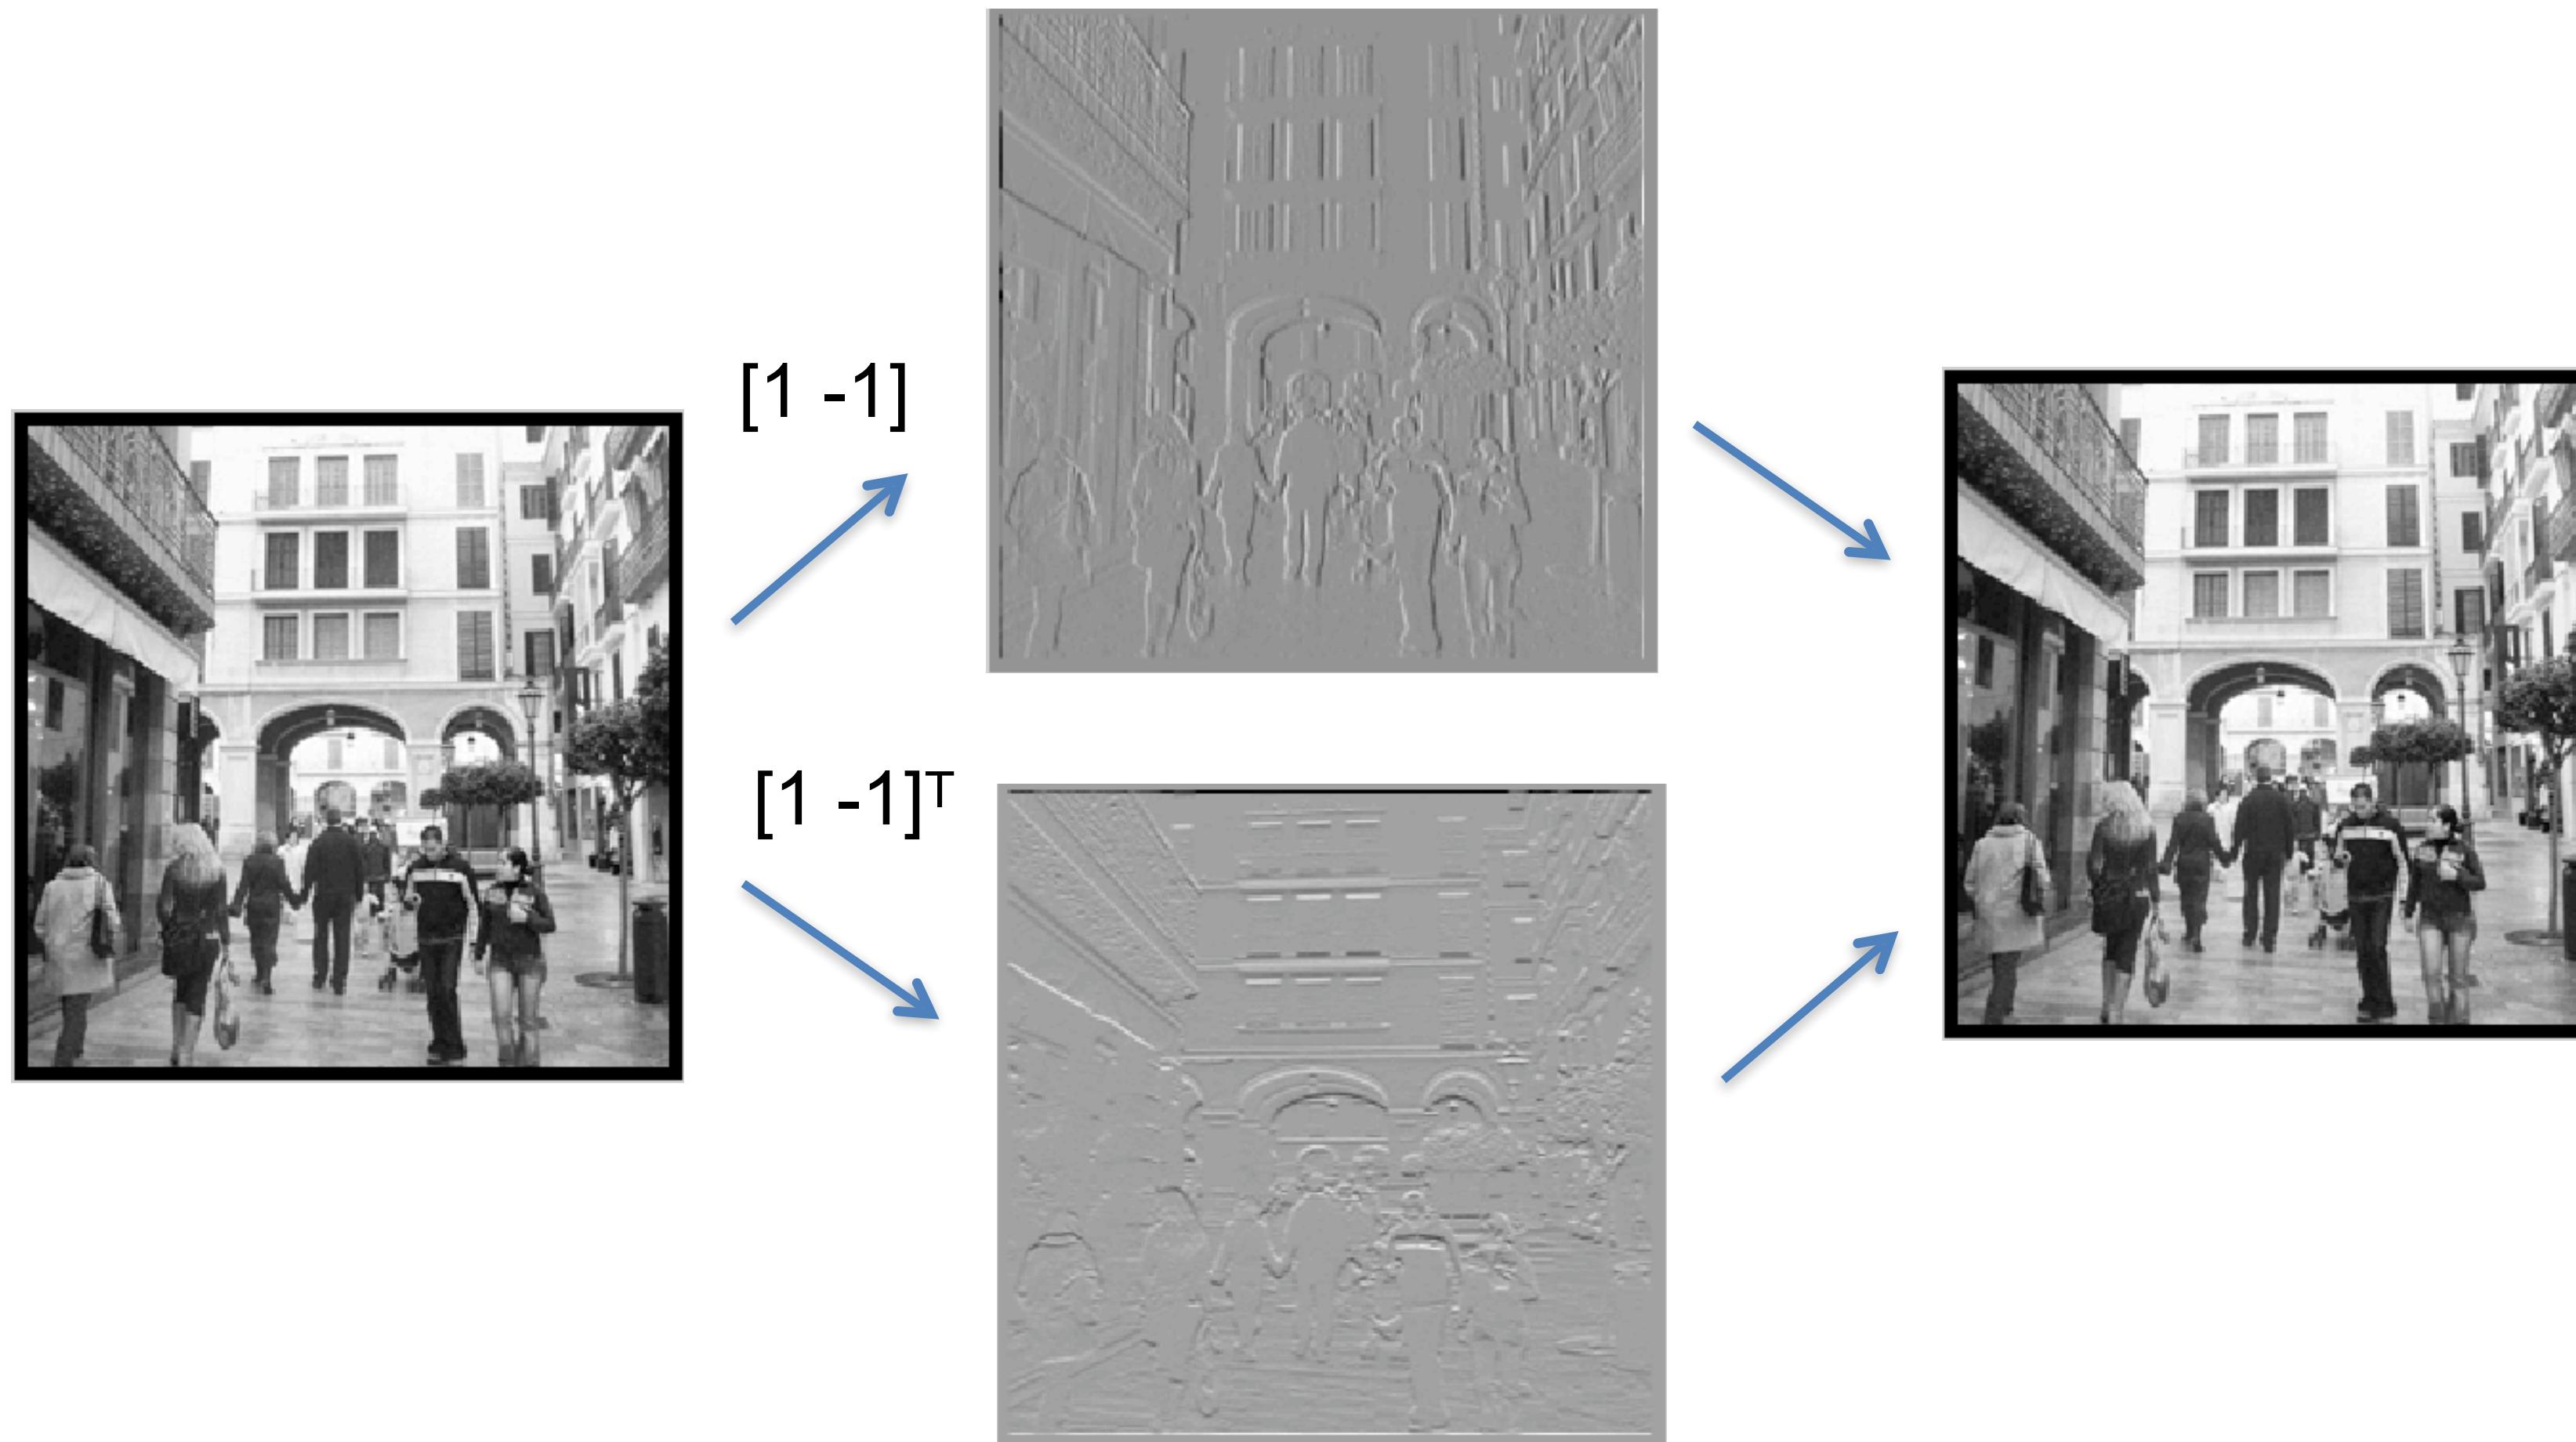
Reconstruction from 2D derivatives

In 2D, we have multiple derivatives (along n and m)

$$\begin{matrix} c \\ c \end{matrix} = \begin{matrix} [-1 & 1] \\ [-1 & 1]^T \end{matrix} c$$

and we compute the pseudo-inverse of the full matrix.

Reconstruction from 2D derivatives

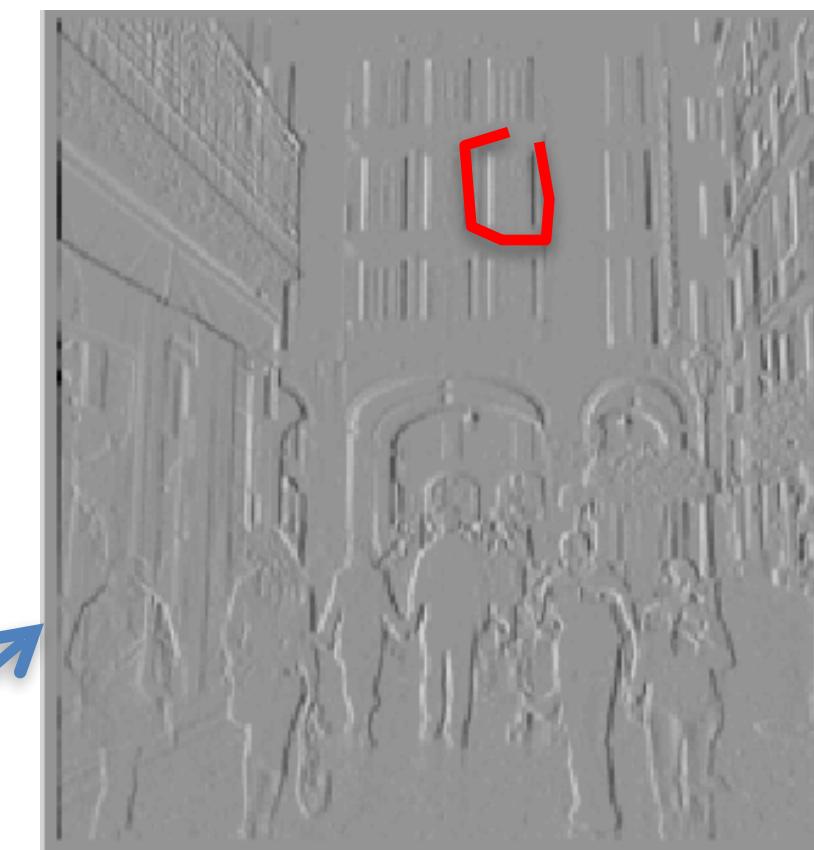


Source: Torralba, Freeman, Isola

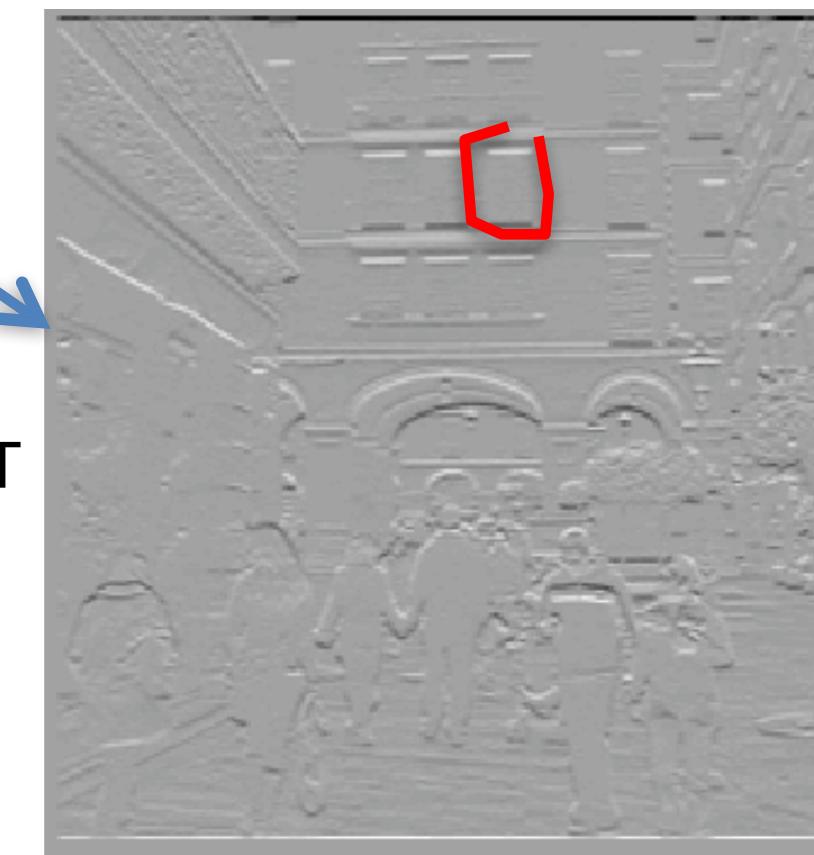
Editing the edge image



$$[1 \ -1]$$



$$[1 \ -1]^T$$



Thresholding edges

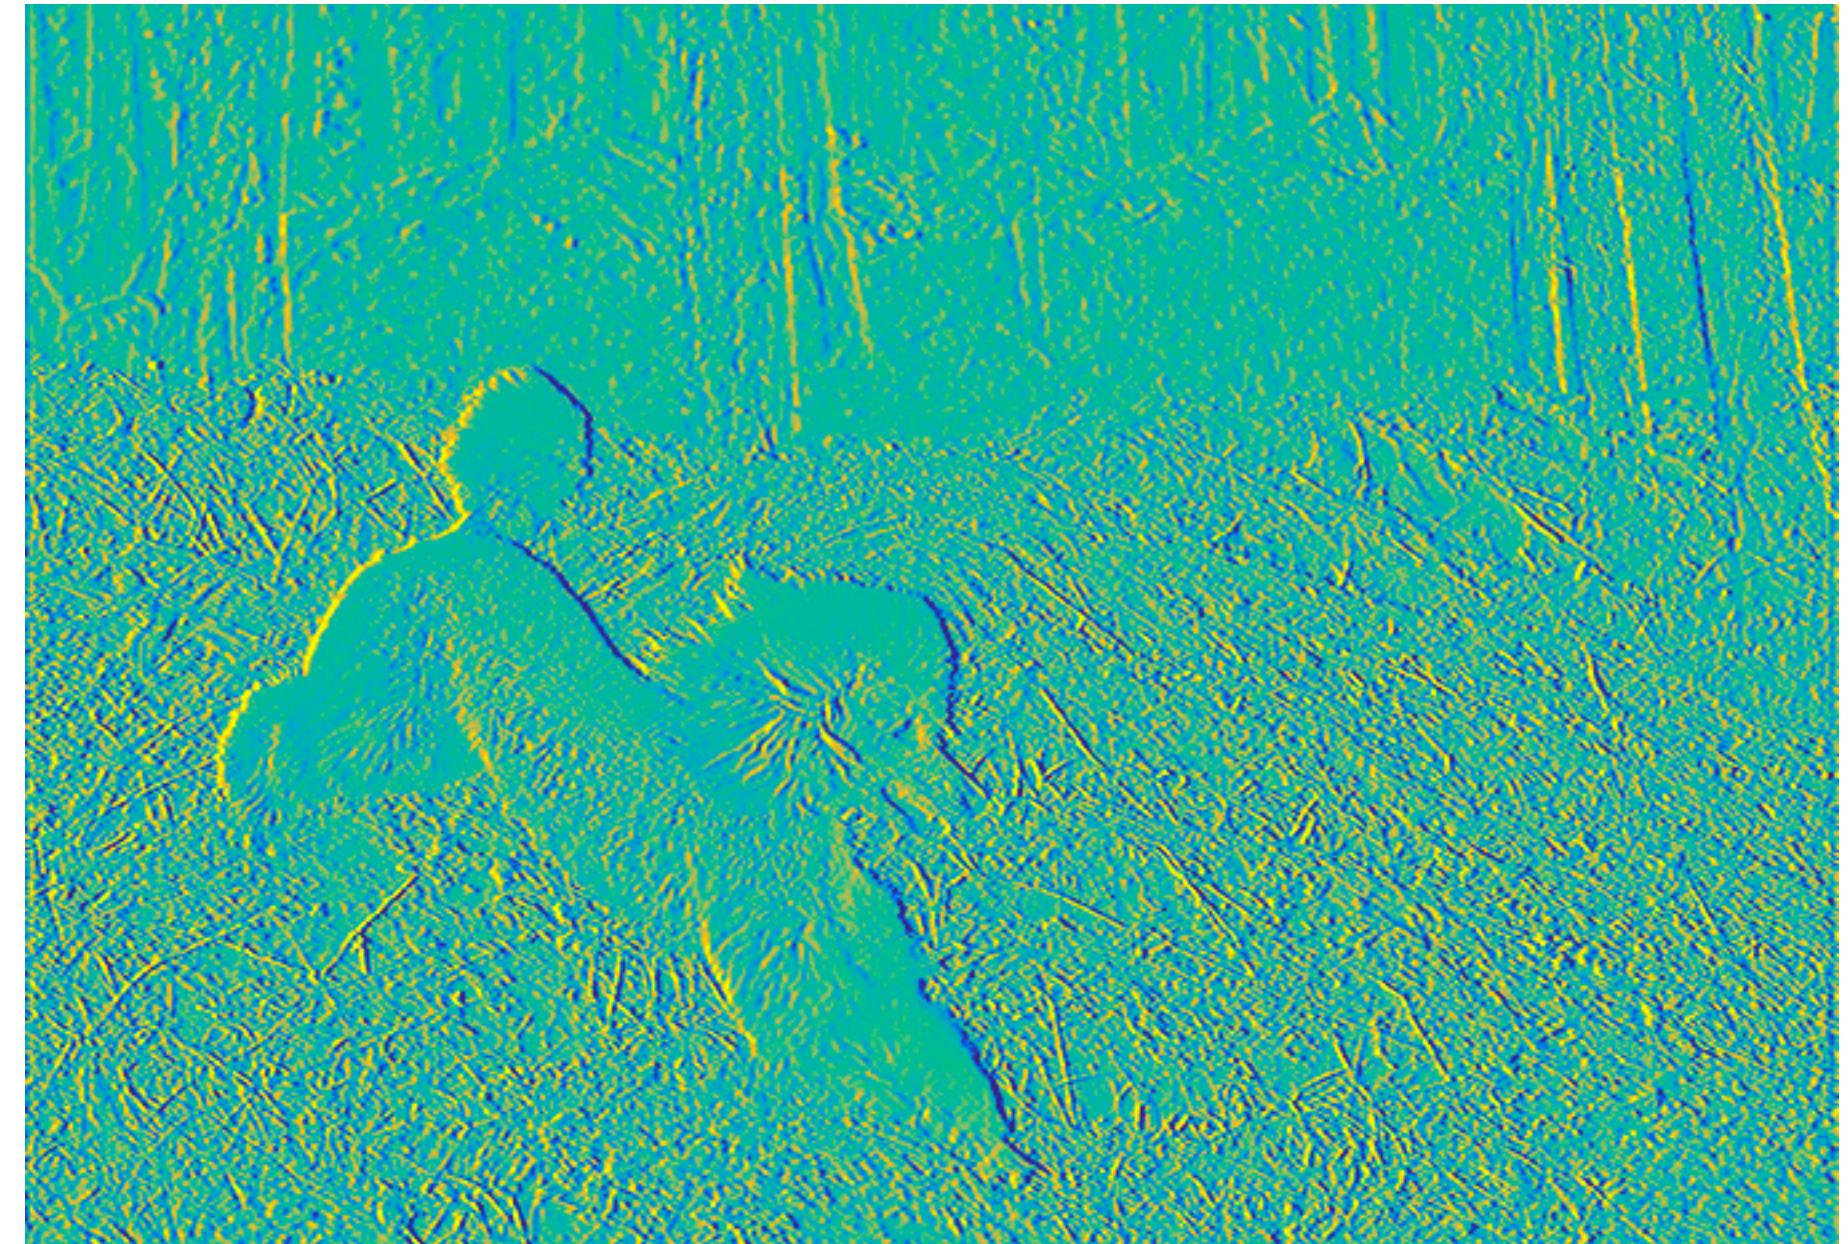


Source: Torralba, Freeman, Isola

Issues with derivative filters

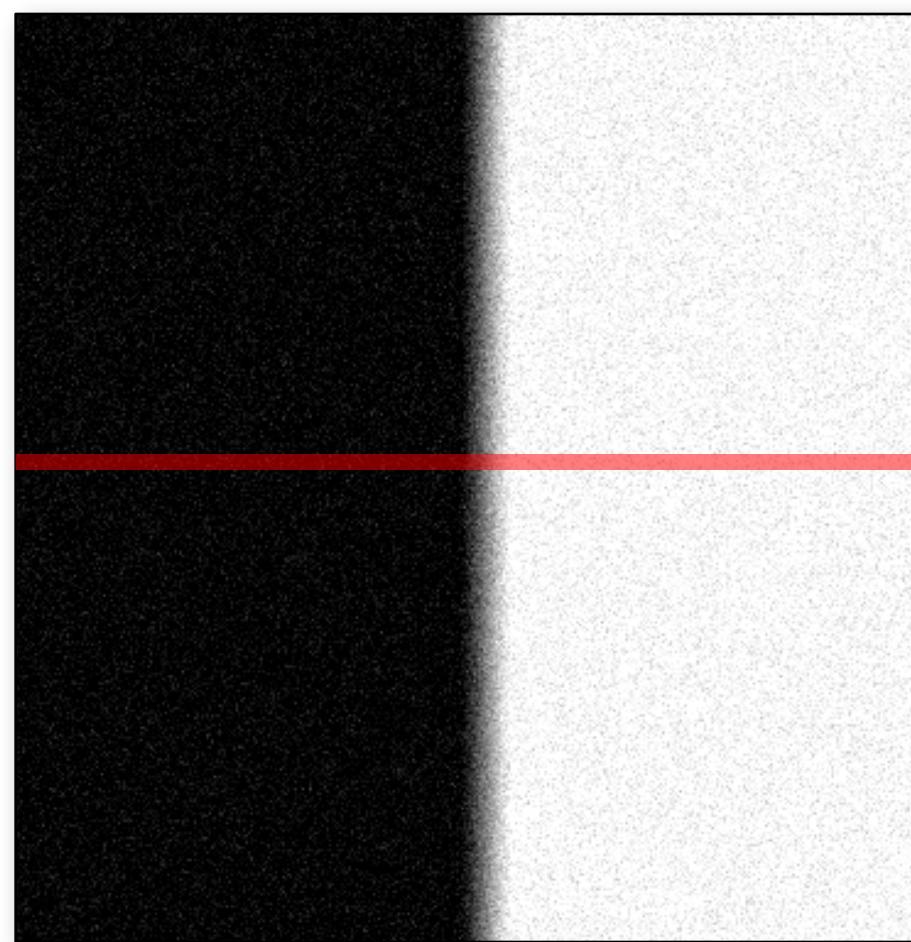


[1 -1]

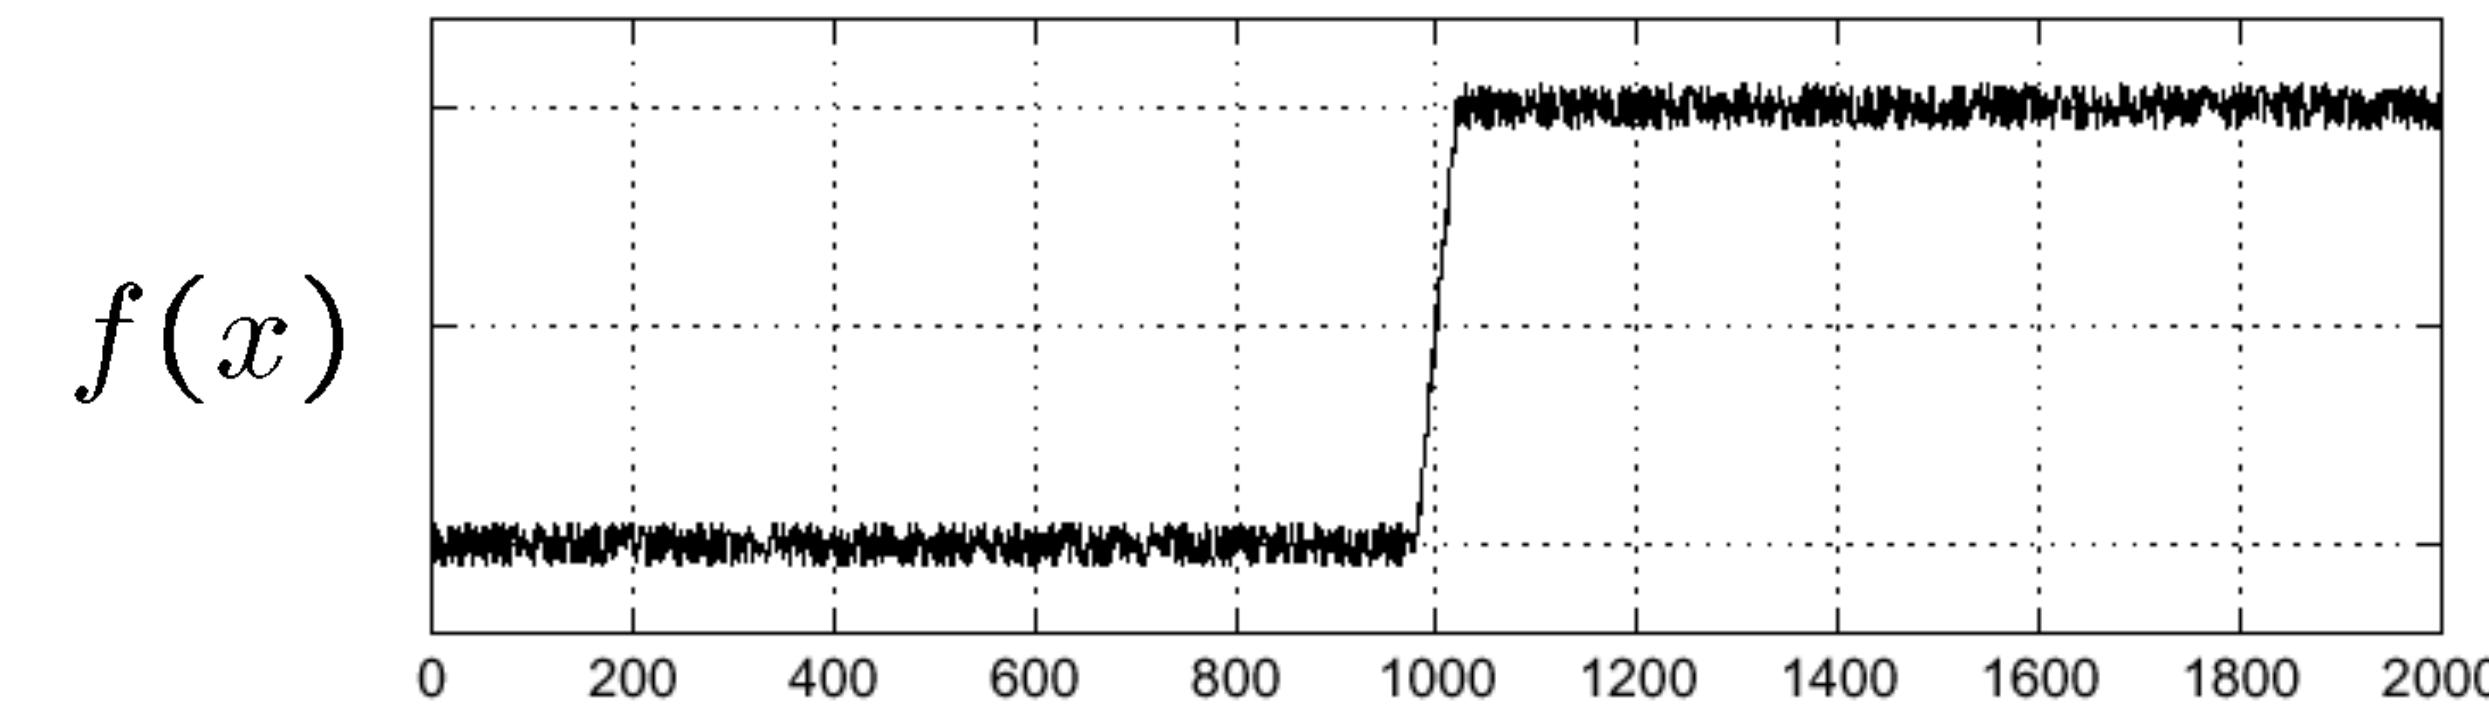


- Sensitive to edges at small spatial scales
- Also sensitive to noise
- You'll see this in Problem Set 1

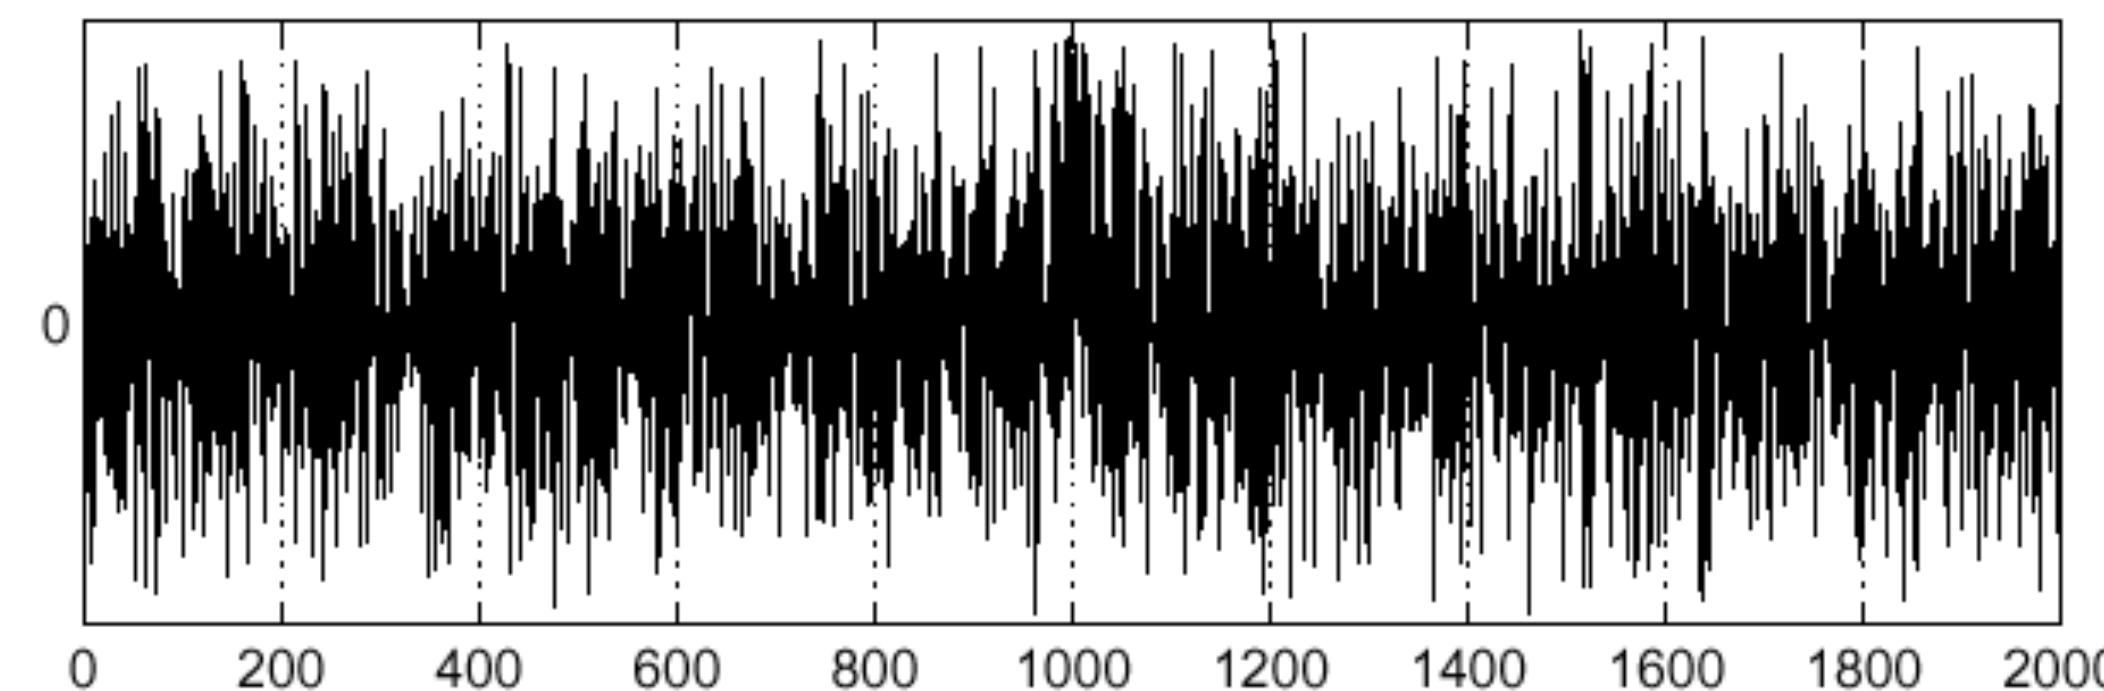
Why is this happening?



Noisy input image

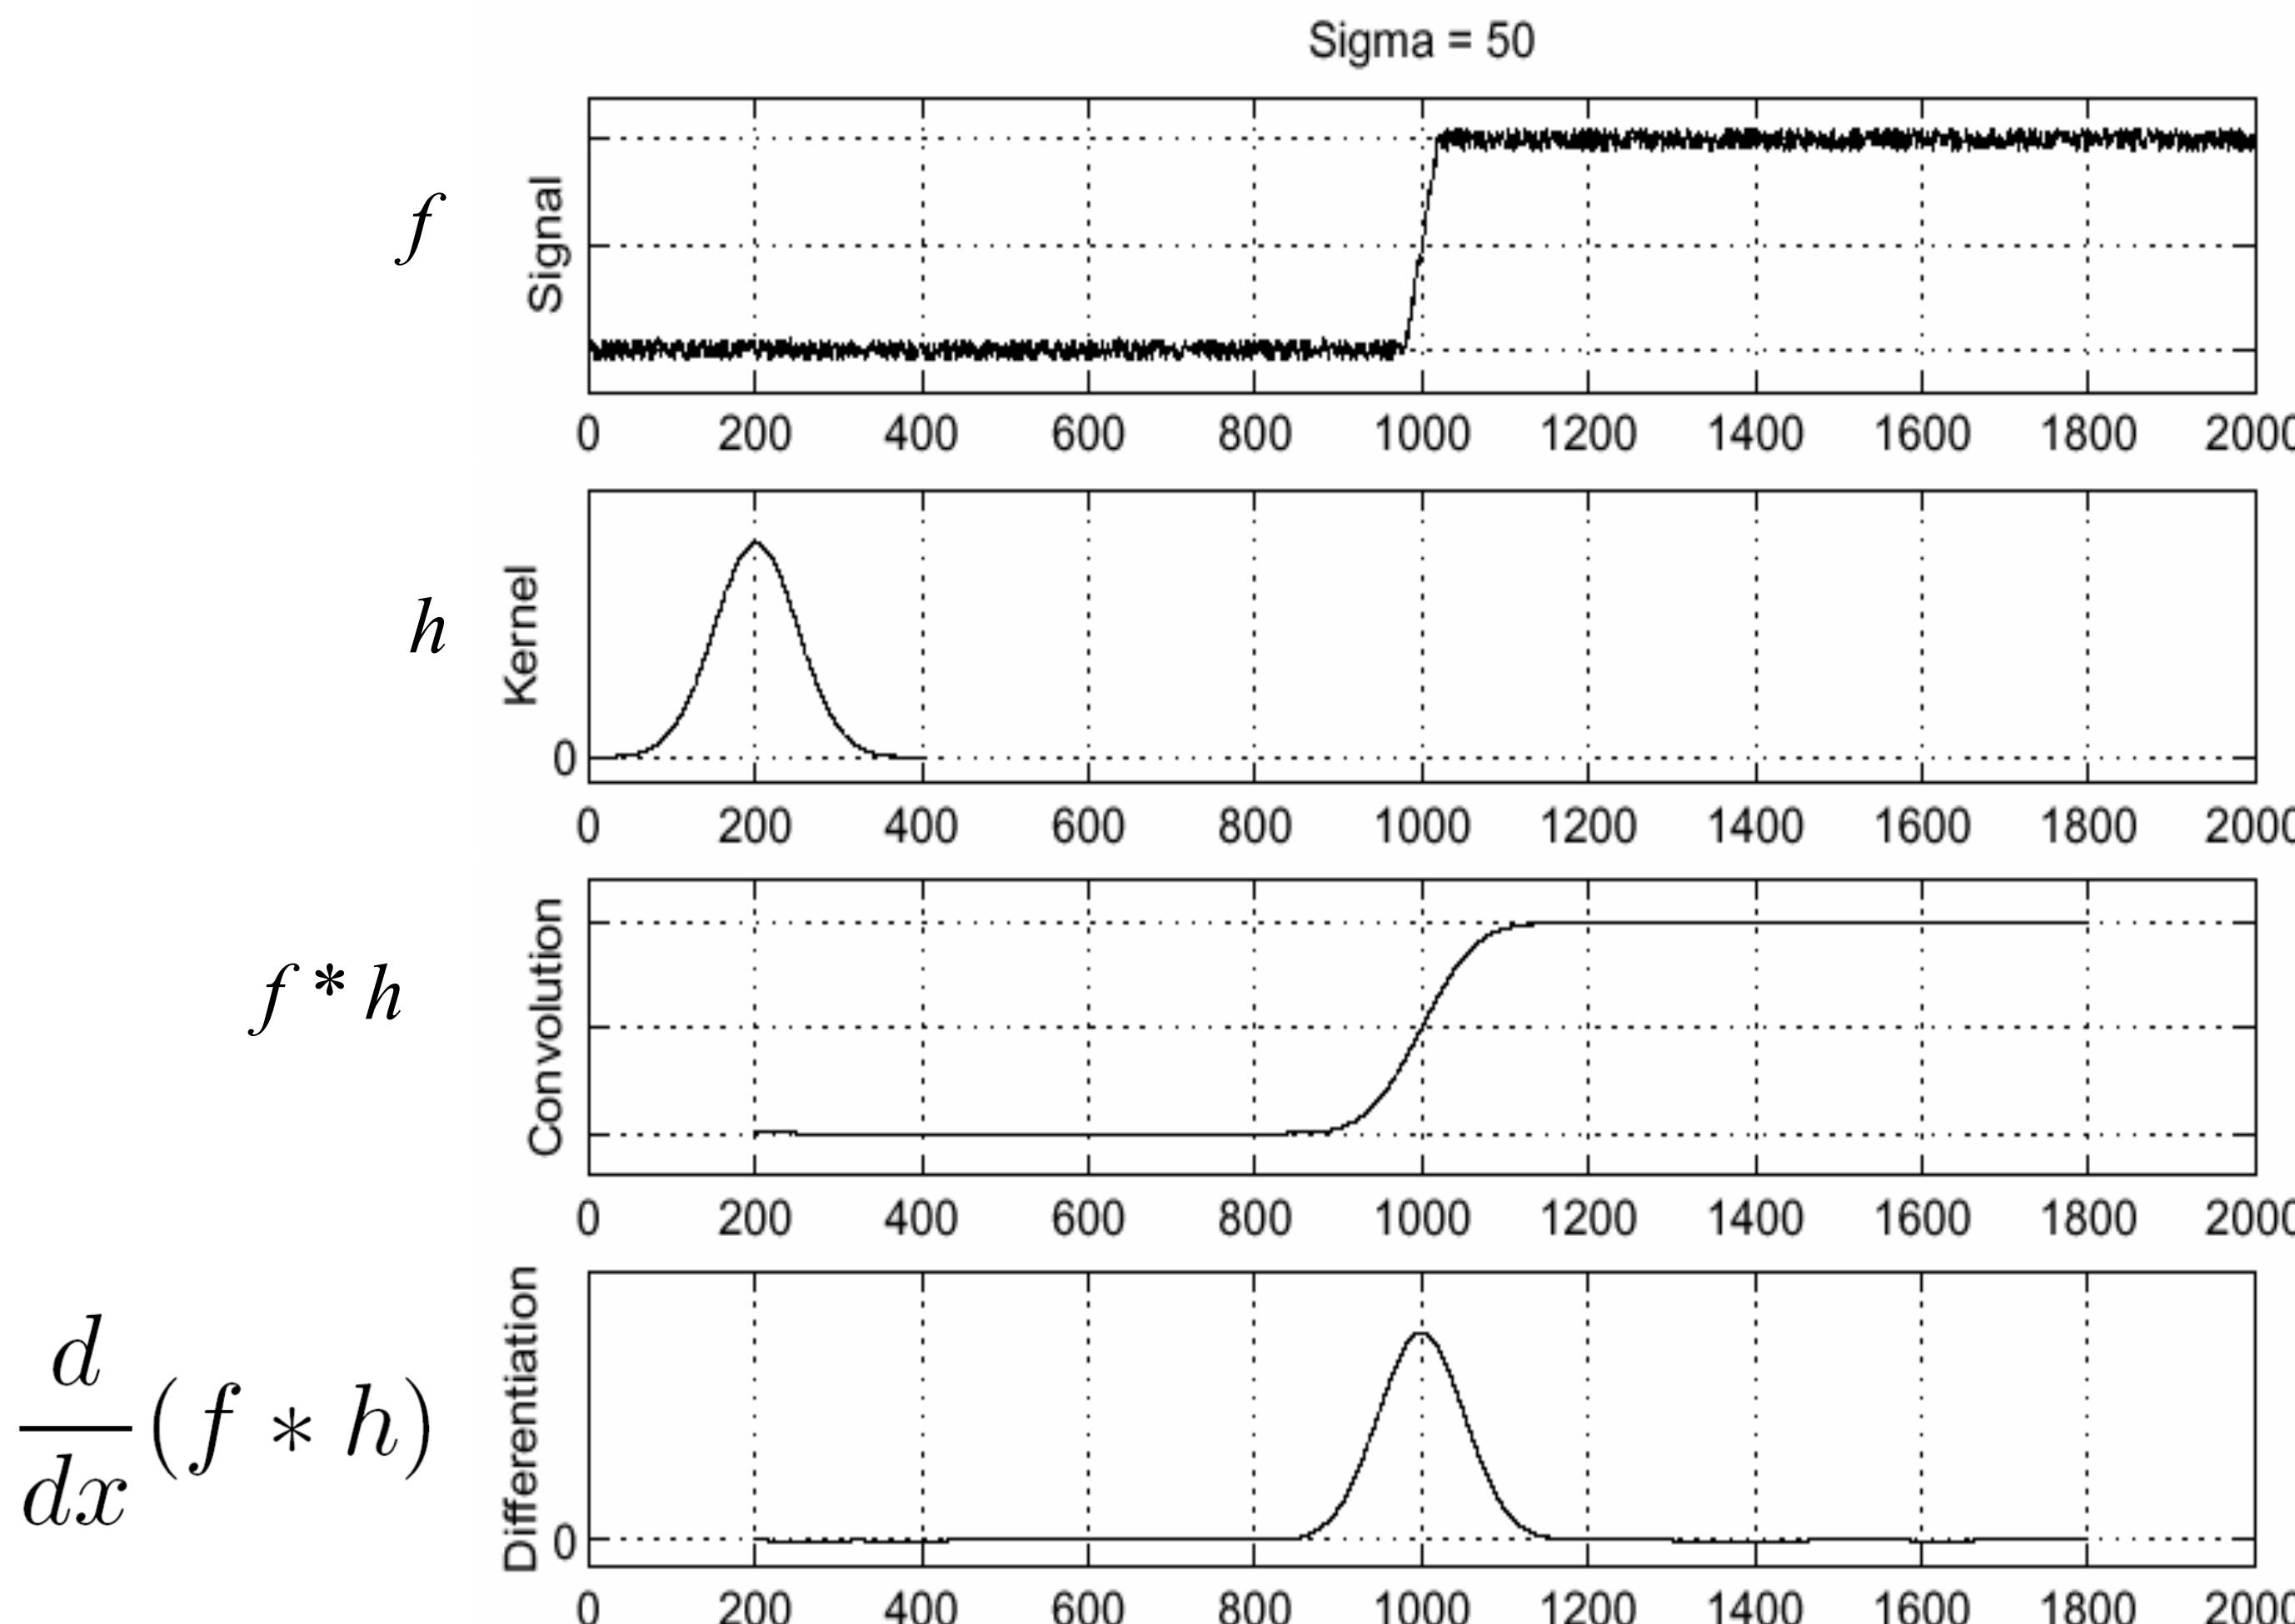


$$\frac{d}{dx}f(x)$$



Where is the edge?

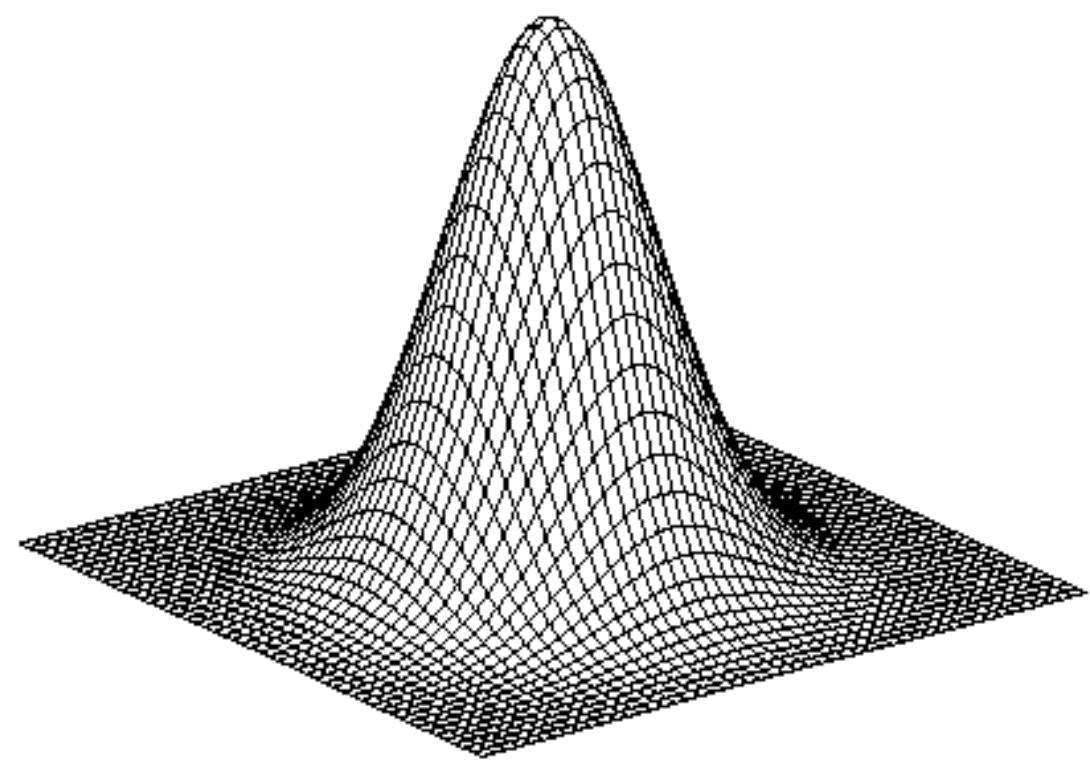
Solution: smooth first



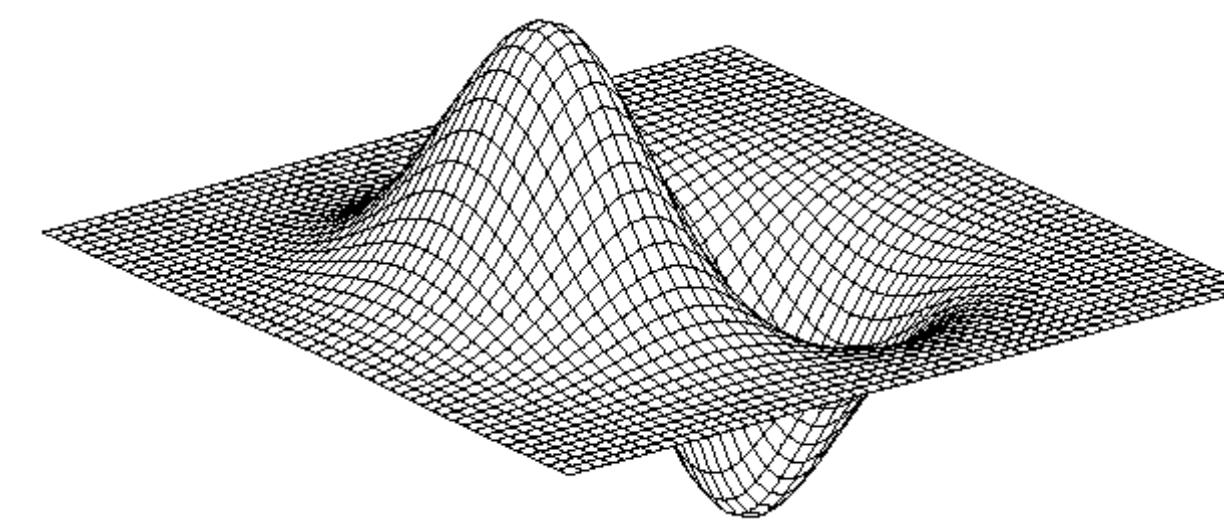
To find edges, look for peaks in $\frac{d}{dx}(f * h)$

Source: S. Seitz

Derivative of Gaussian filter



Gaussian

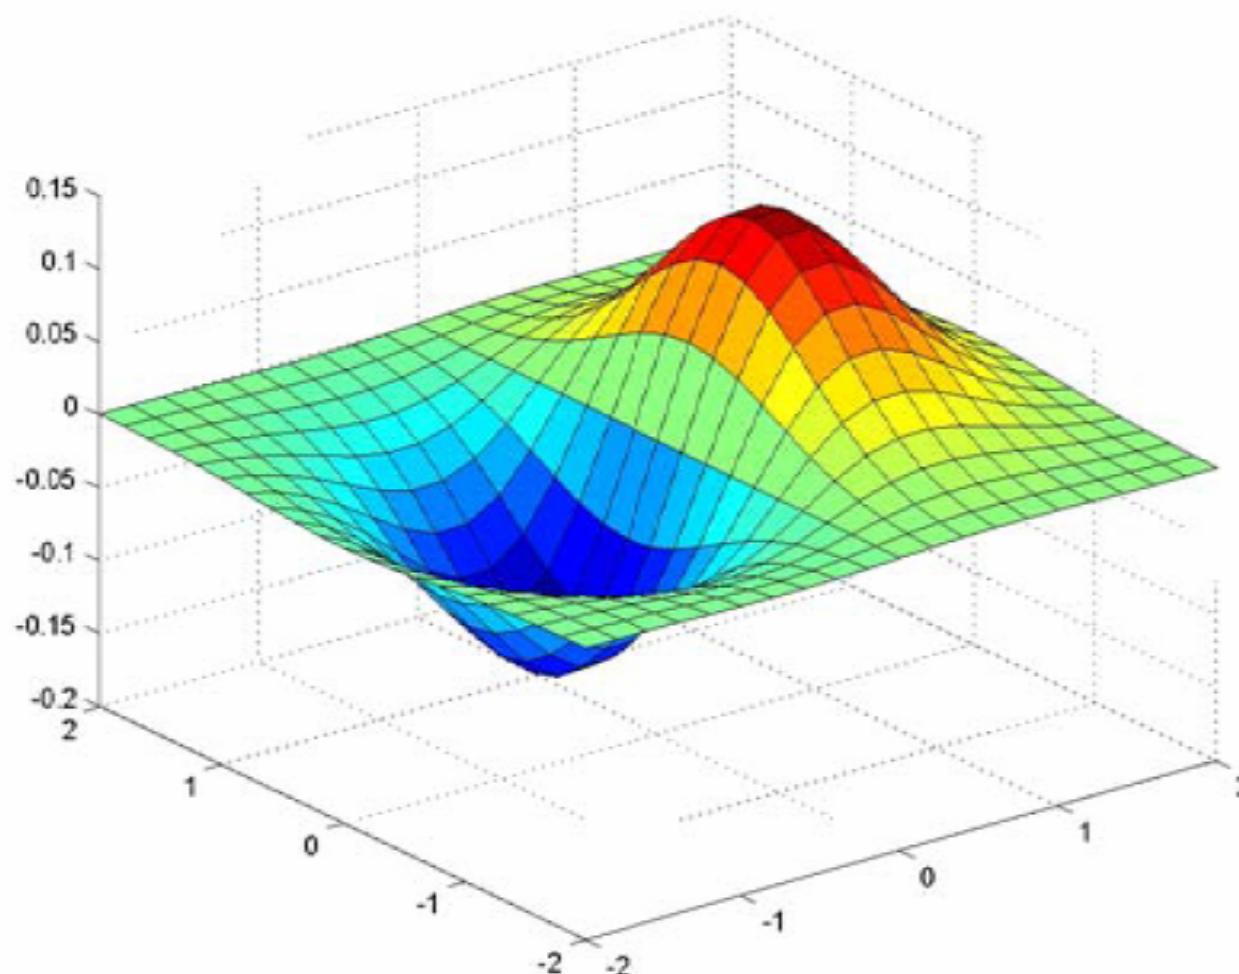


Derivative of Gaussian (x)

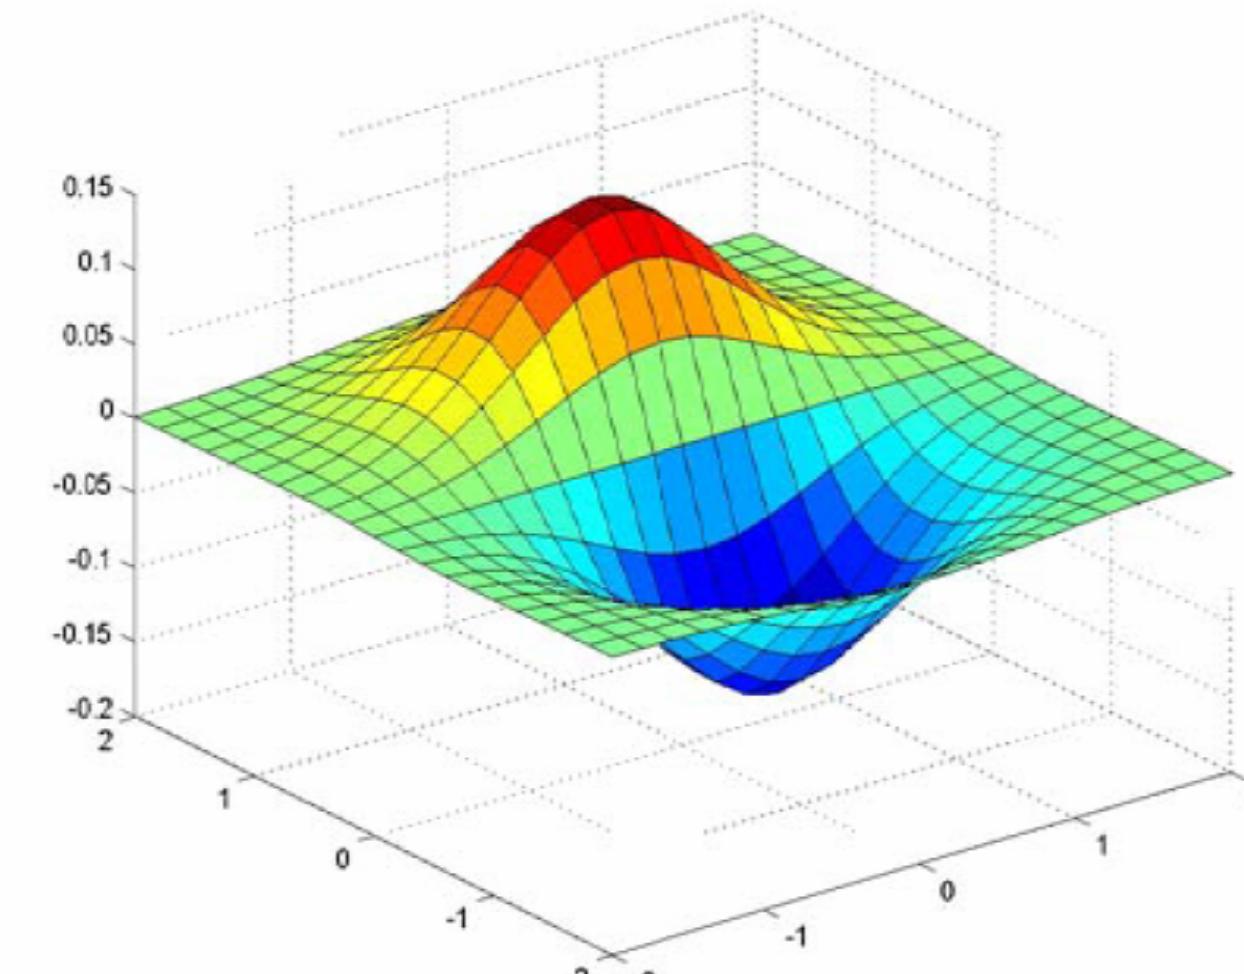
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

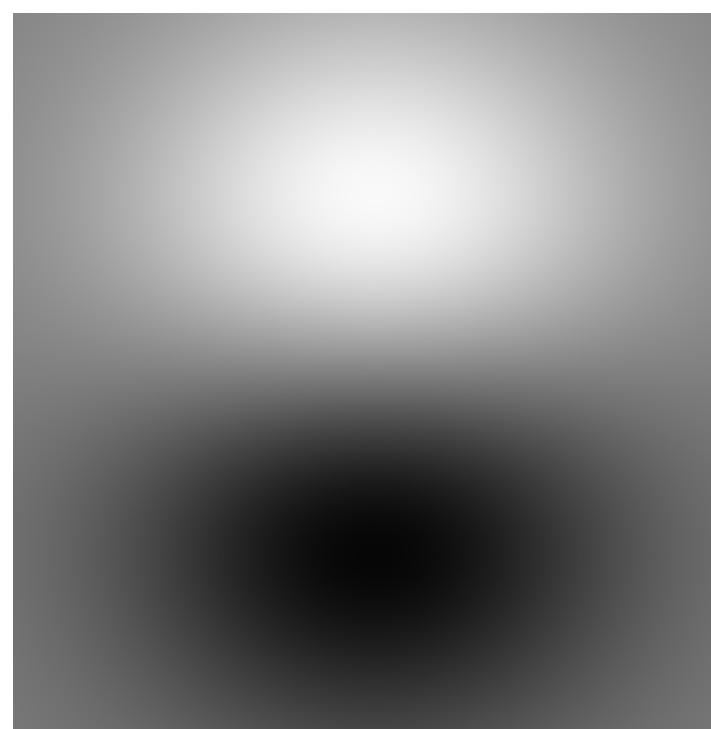
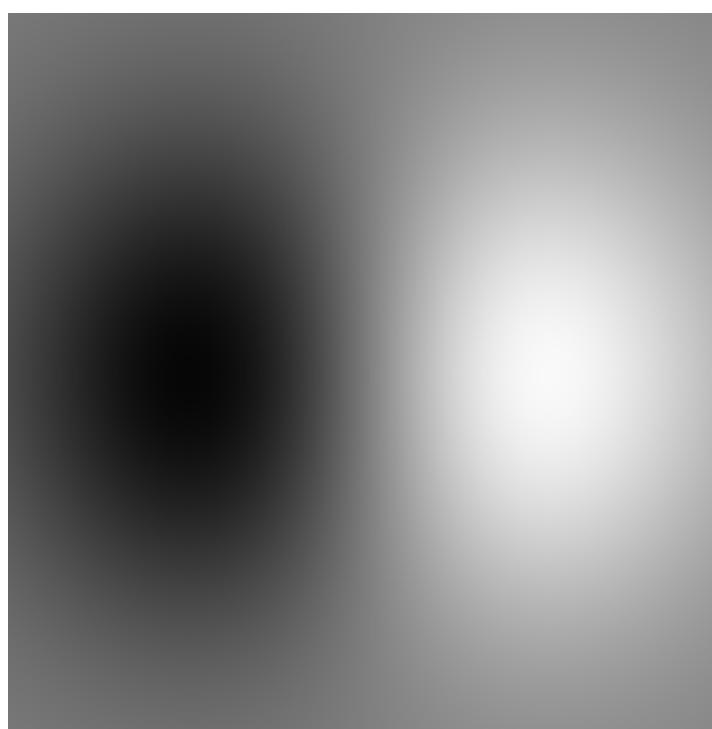
Derivative of Gaussian filter



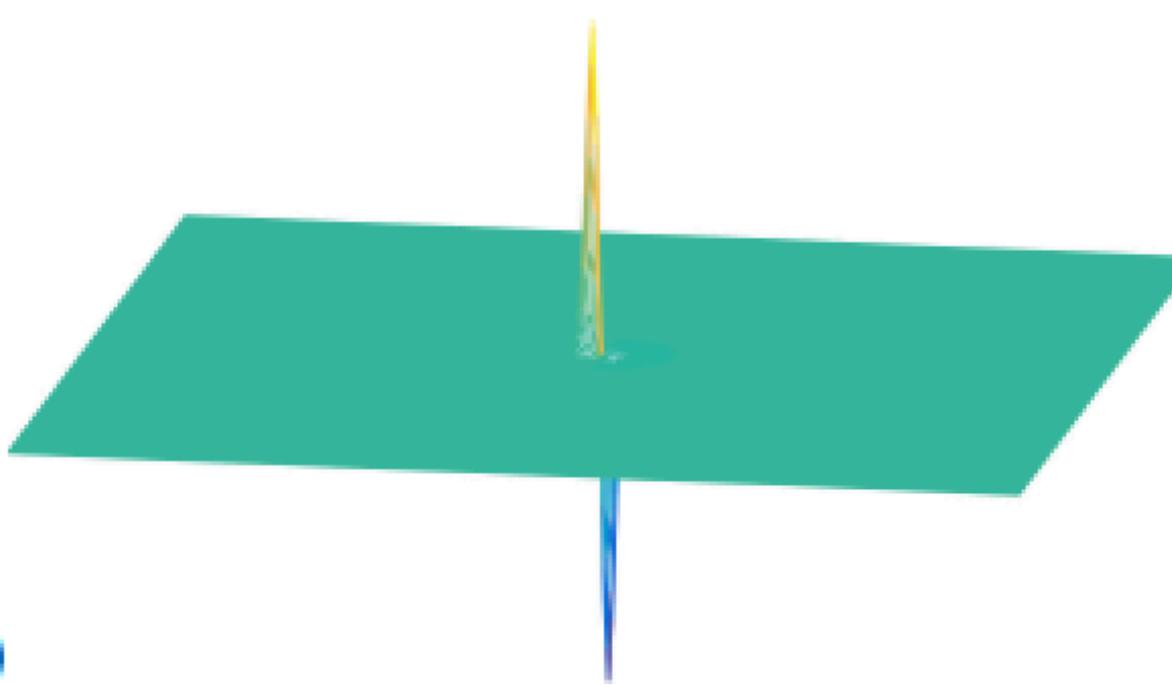
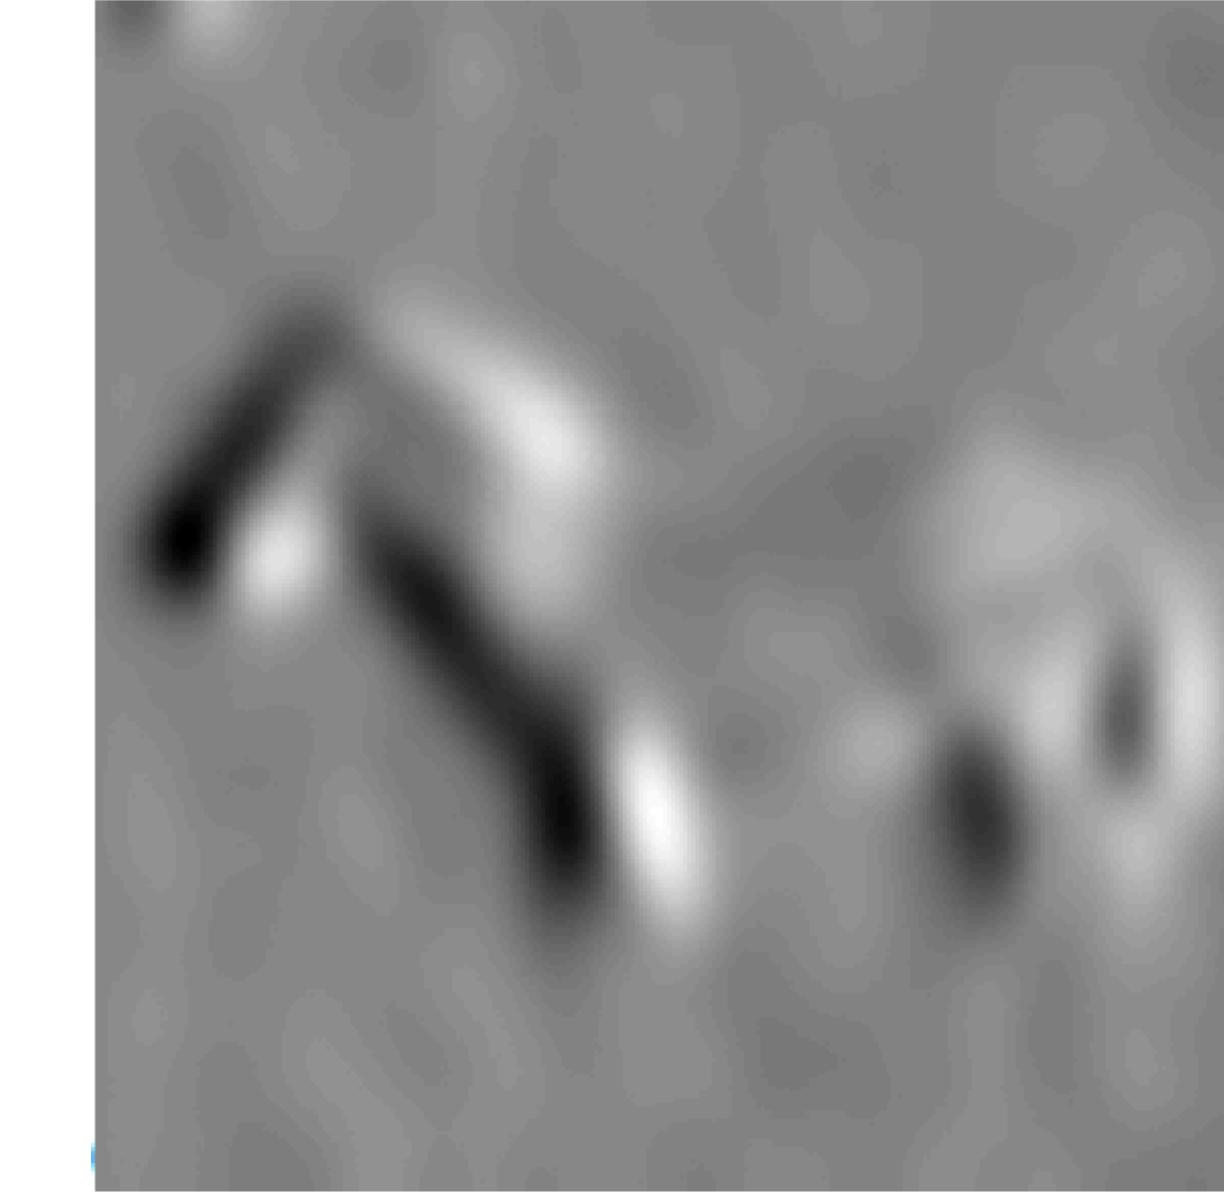
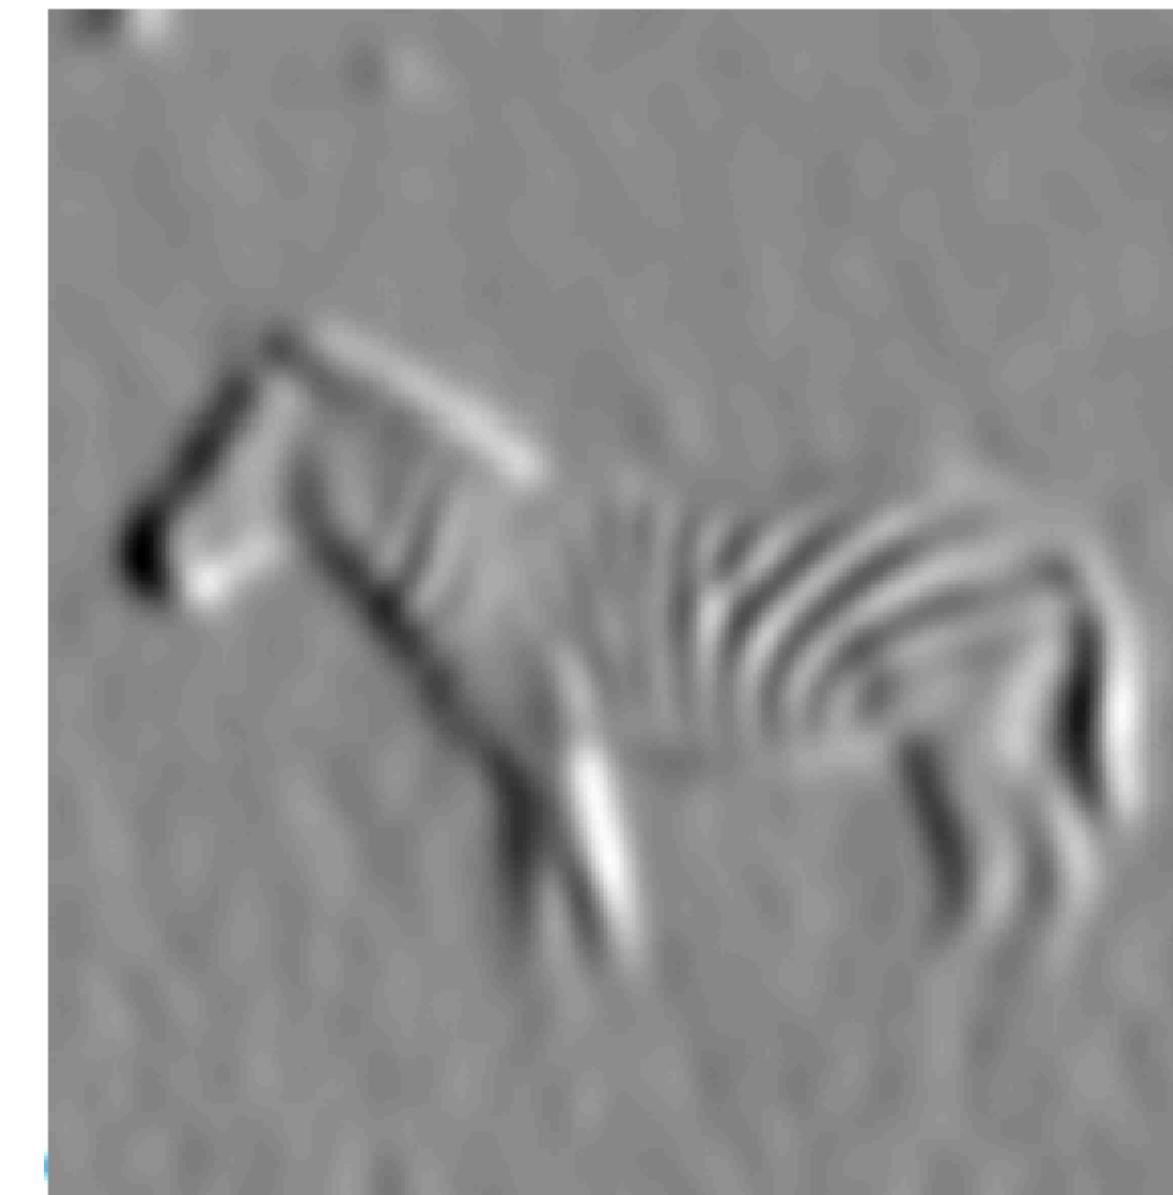
x-direction



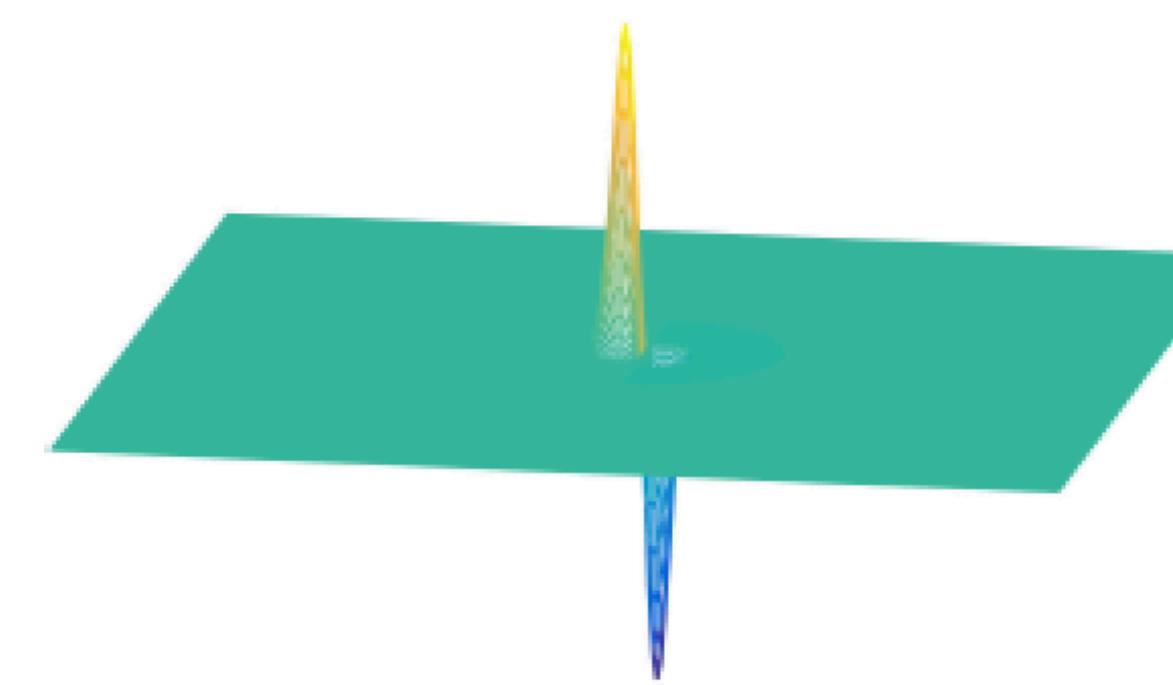
y-direction



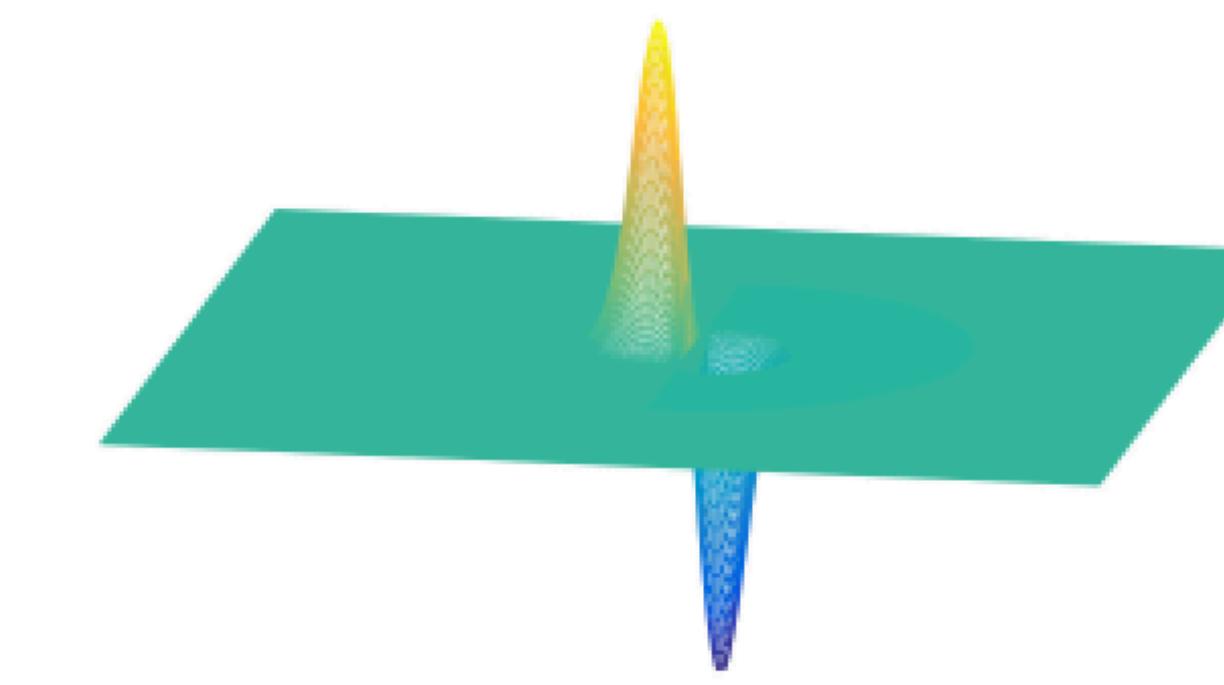
Derivatives of Gaussians: Scale



$\sigma=2$



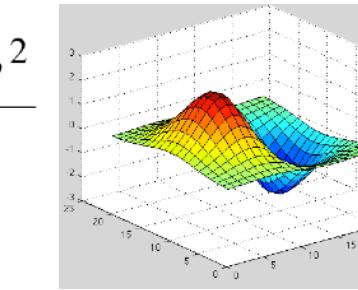
$\sigma=4$



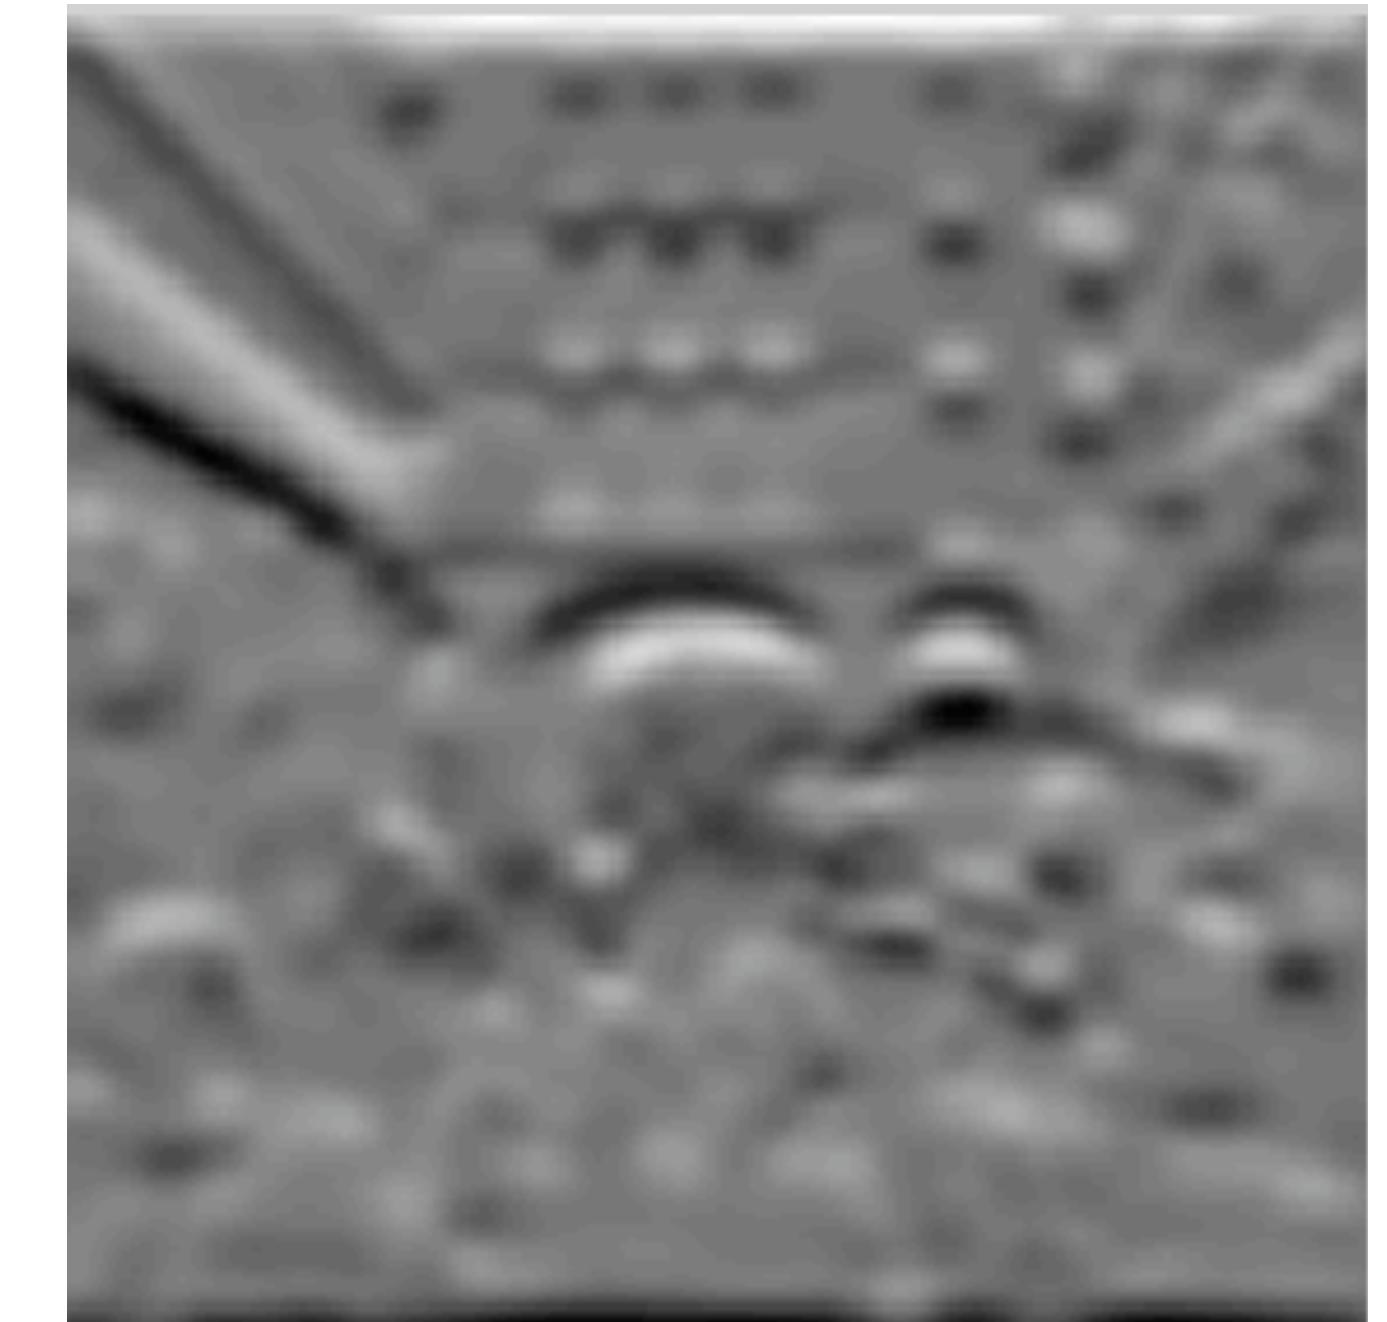
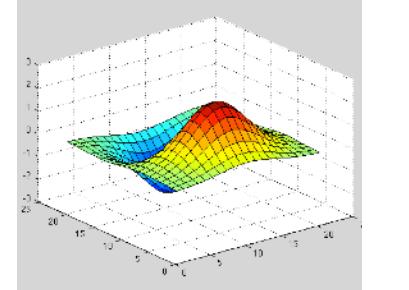
$\sigma=8$

Picks up larger-scale edges

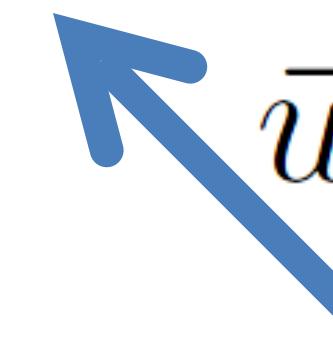
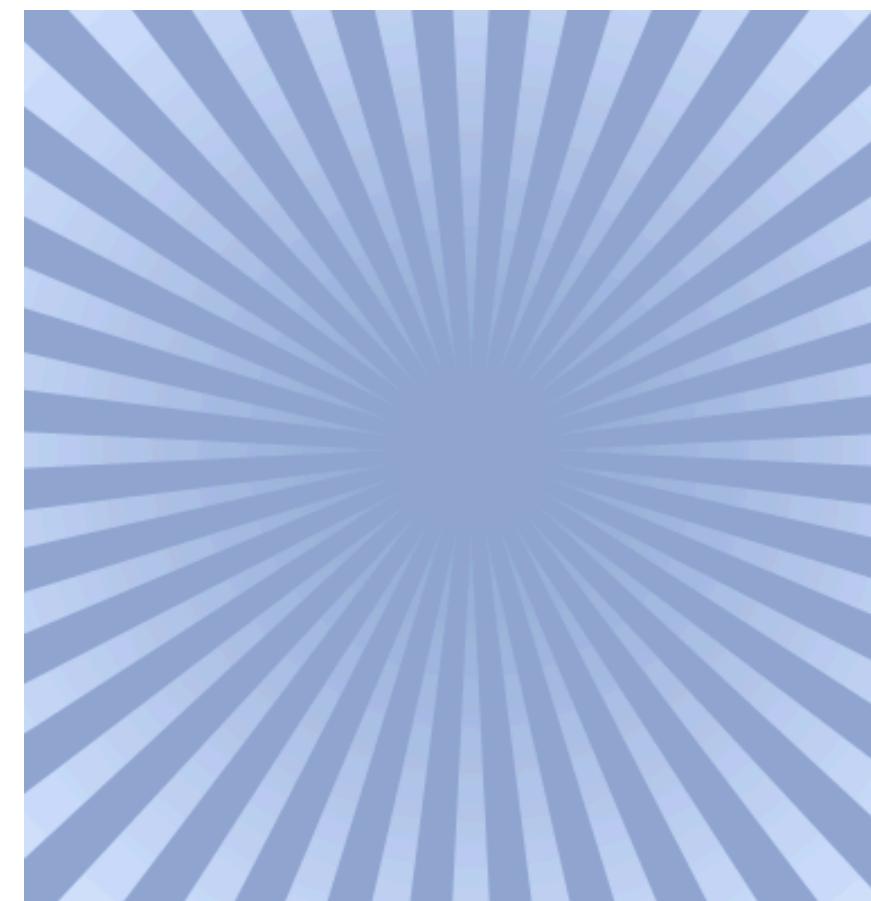
$$g_x(x,y) = \frac{\partial g(x,y)}{\partial x} = \frac{-x}{2\pi\sigma^4} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



$$g_y(x,y) = \frac{\partial g(x,y)}{\partial y} = \frac{-y}{2\pi\sigma^4} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



Computing a directional derivative

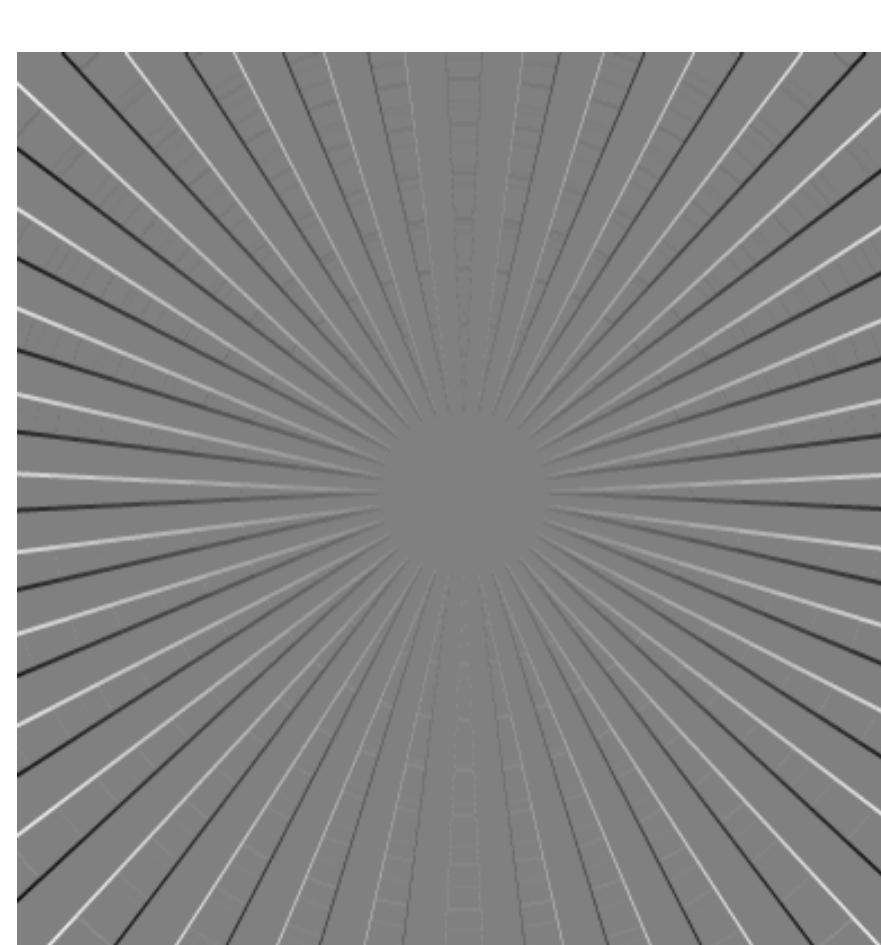


$$\nabla_{\vec{u}} f = ?$$

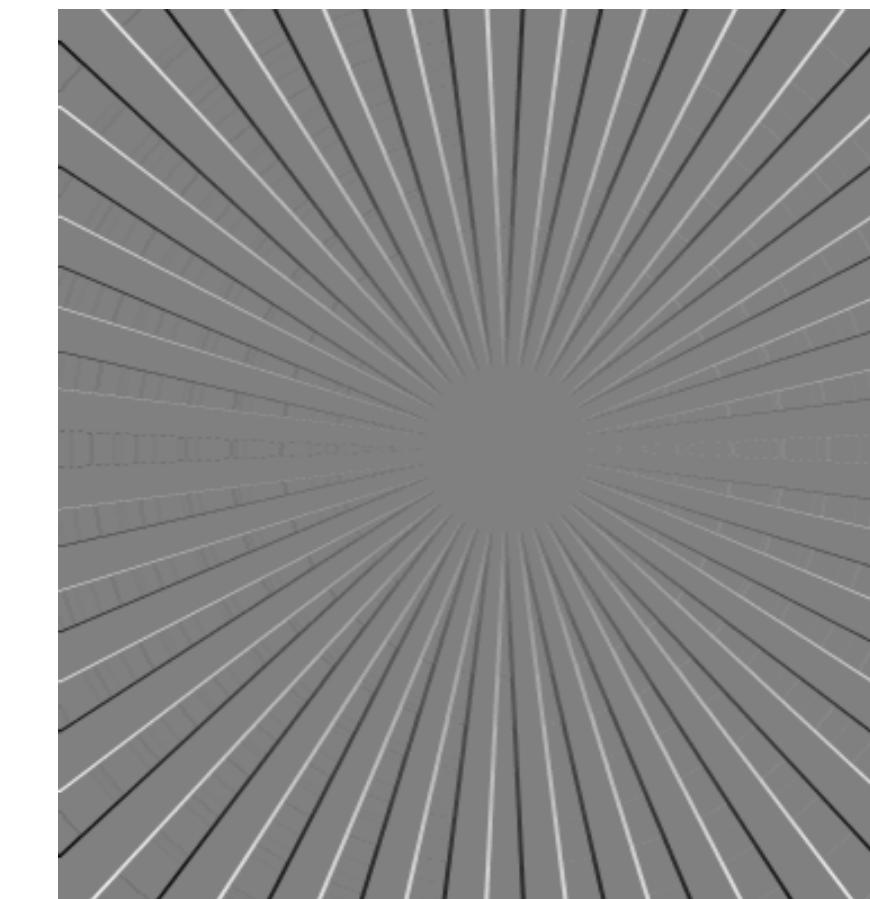
(From multivariable calculus)

$$\nabla_{\vec{u}} f(\vec{x}) = \nabla f(\vec{x}) \cdot \vec{u}$$

Directional derivative is a linear combination of partial derivatives



f



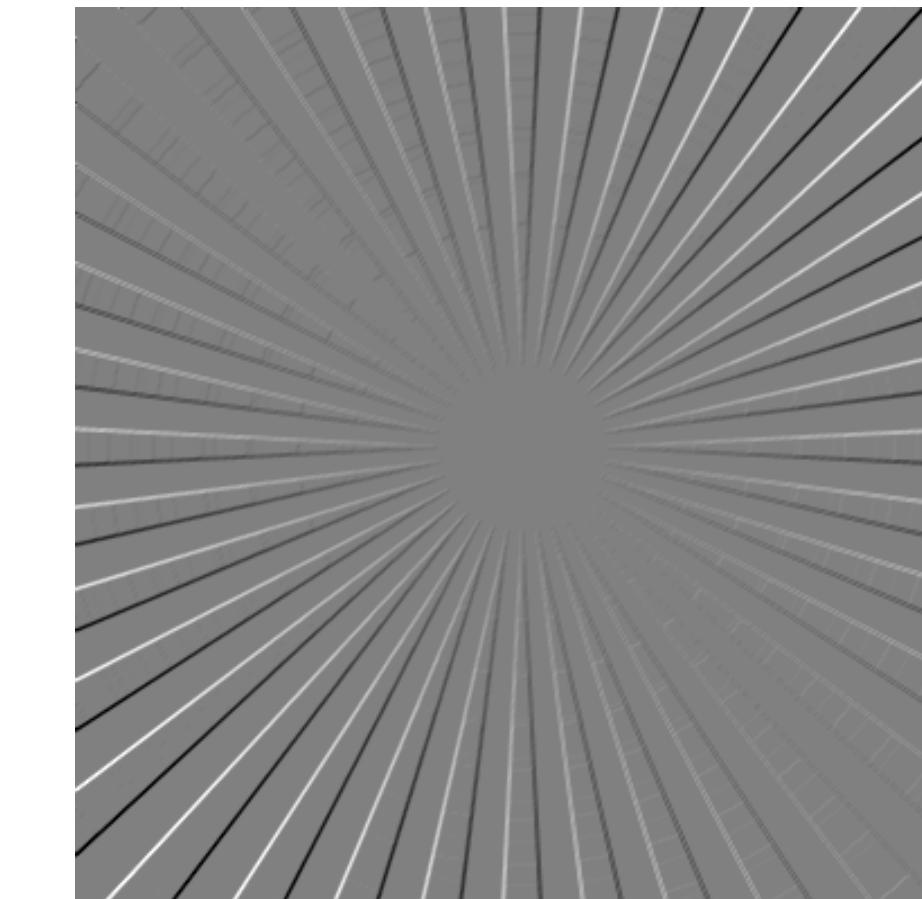
$$\frac{\partial f}{\partial x} \cdot u_x$$

+

$$\frac{\partial f}{\partial y} \cdot u_y$$

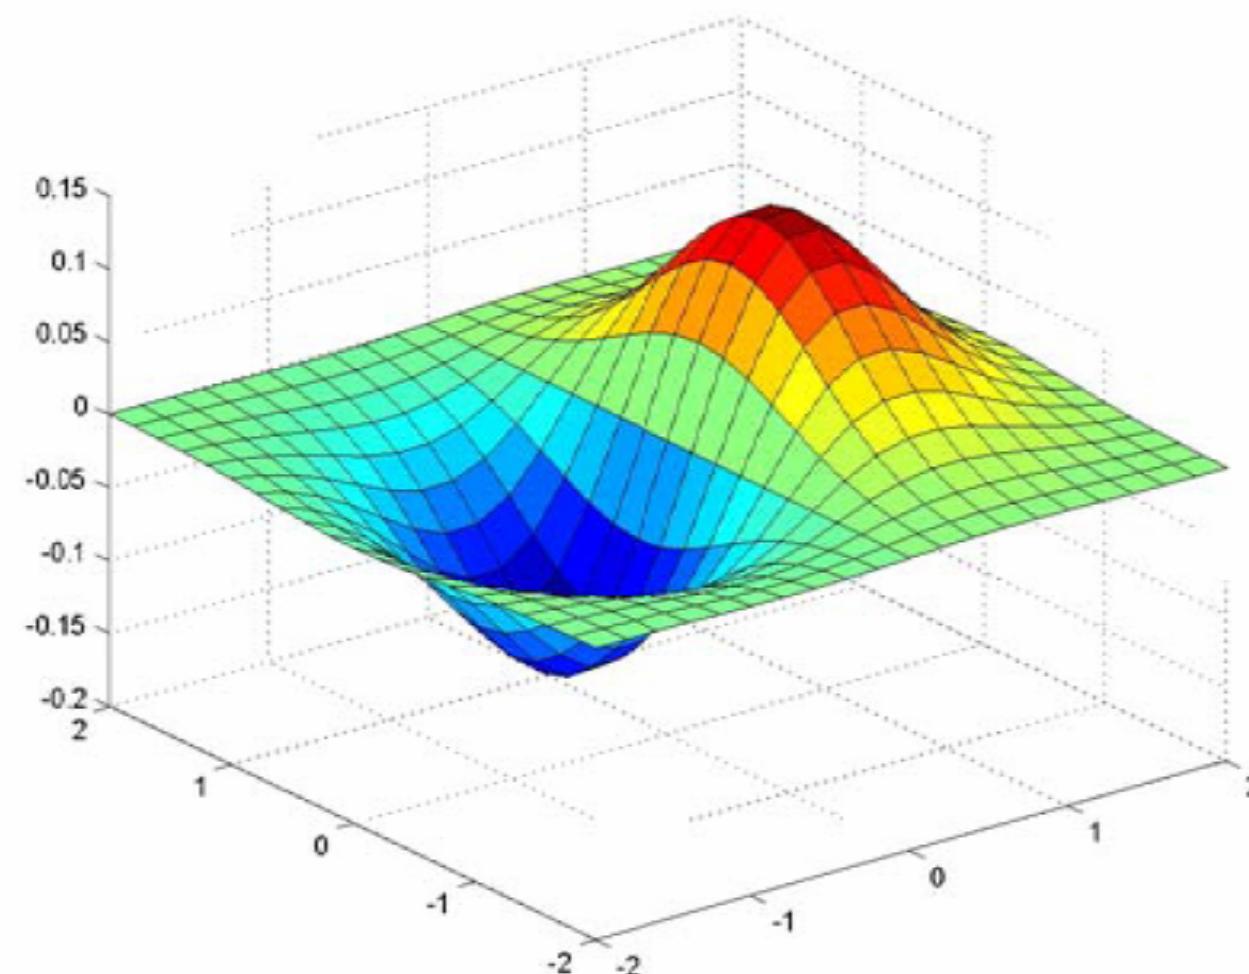
=

$$\nabla_{\vec{u}} f$$

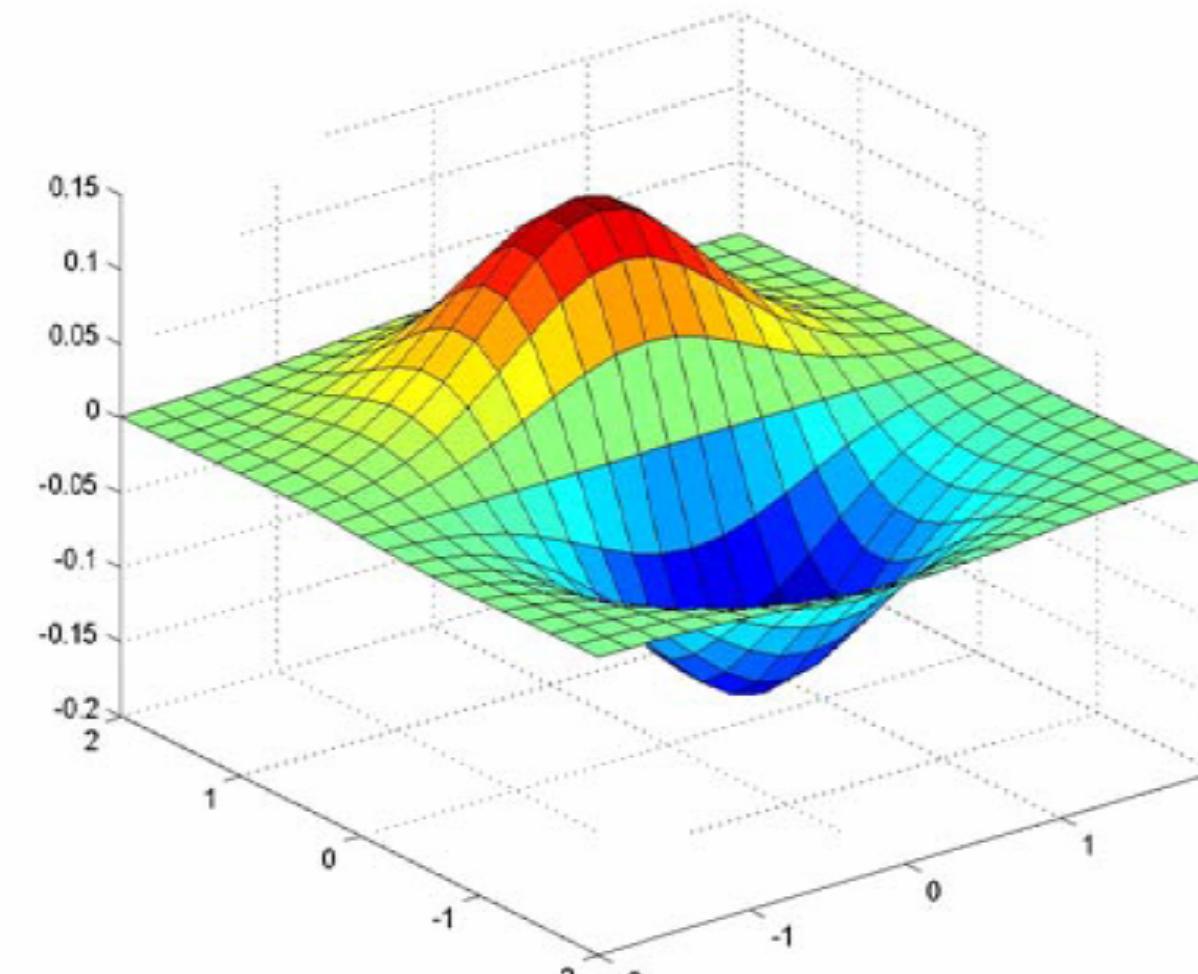


Source: N. Snavely

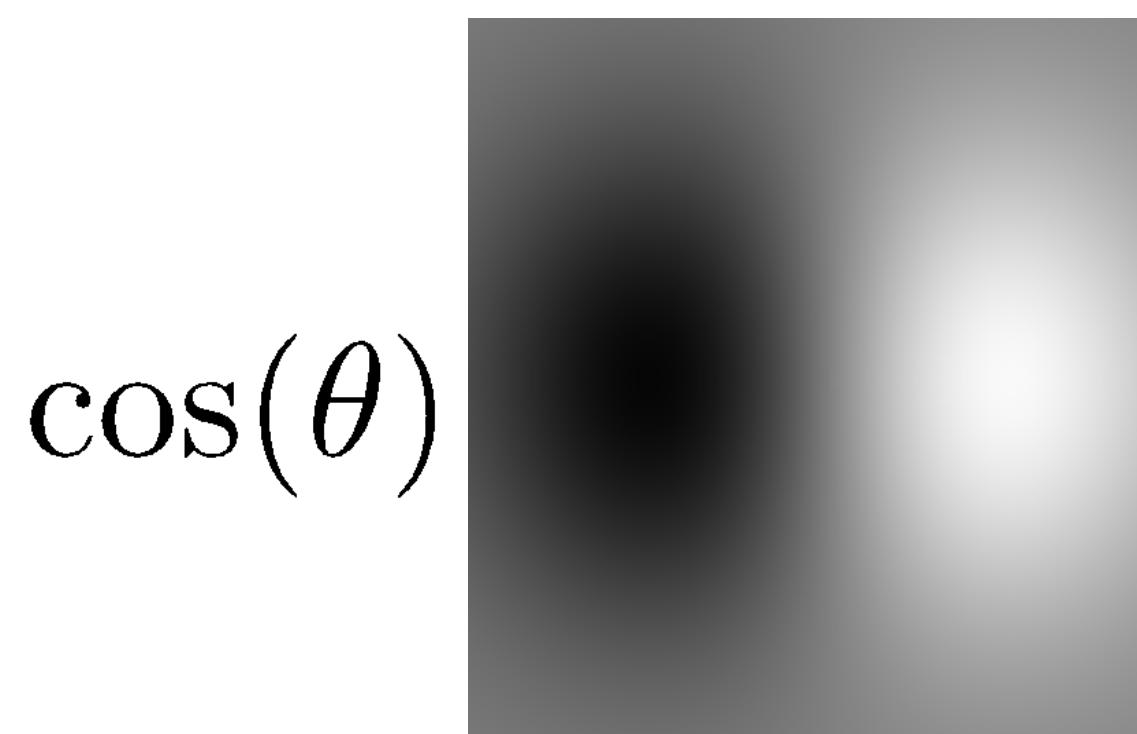
Derivative of Gaussian filter



x-direction

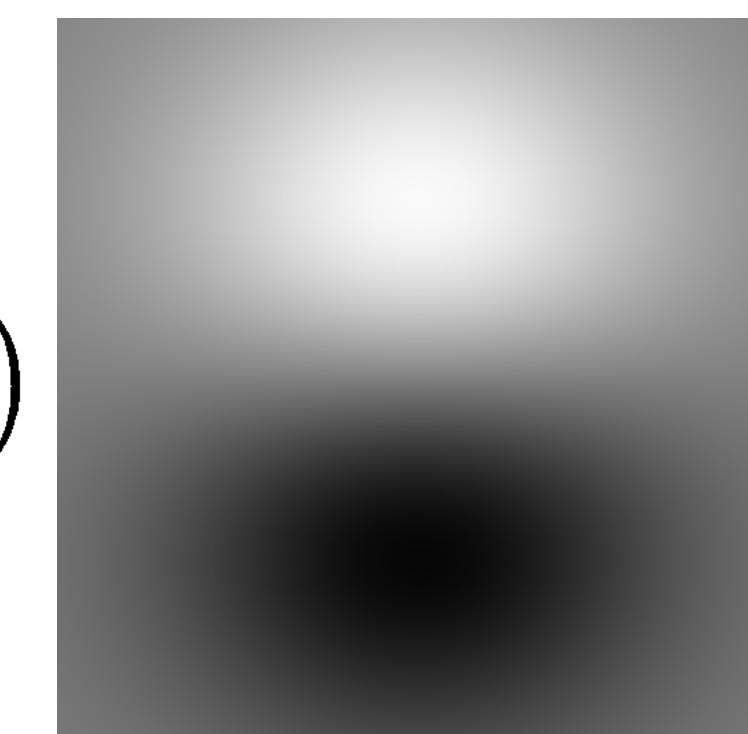


y-direction

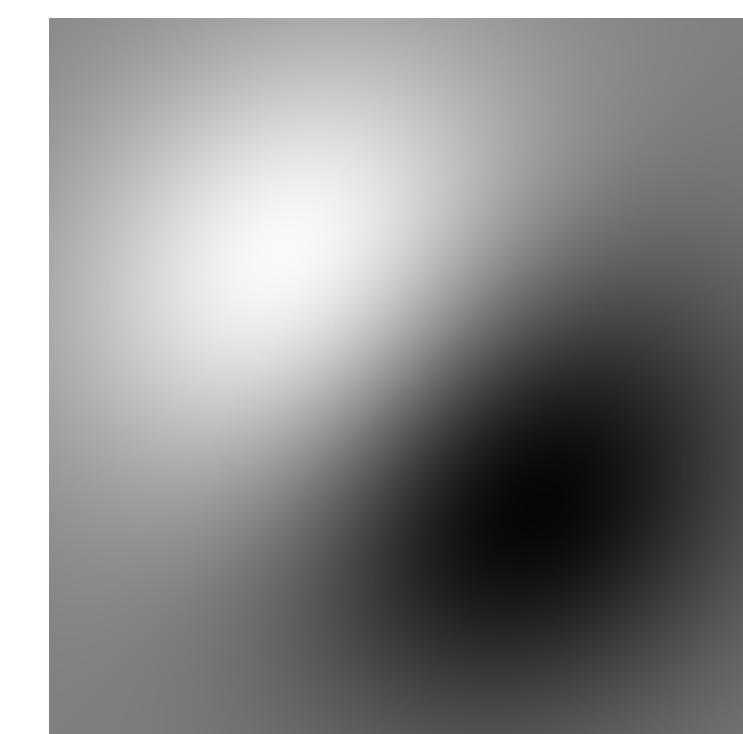


$$\cos(\theta)$$

$$+ \sin(\theta)$$



$$=$$



The Sobel operator

- Common approximation to derivative of Gaussian
- Where does this come from?

$$\frac{1}{8} \begin{array}{|c|c|c|}\hline -1 & 0 & 1 \\ \hline -2 & 0 & 2 \\ \hline -1 & 0 & 1 \\ \hline \end{array}$$

s_x

$$\frac{1}{8} \begin{array}{|c|c|c|}\hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$

s_y

An approximation to the Gaussian

- Apply filter to itself repeatedly.
- Converges to Gaussian, due to Central Limit Theorem

$$b_1 = [1 \ 1]$$

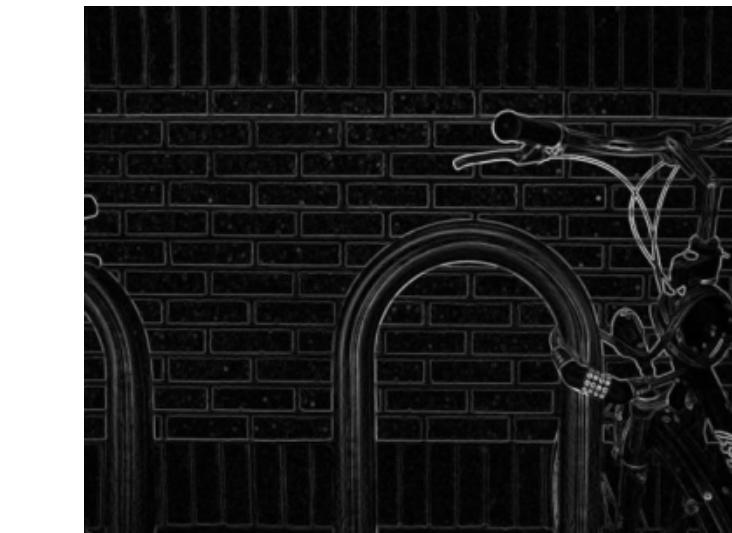
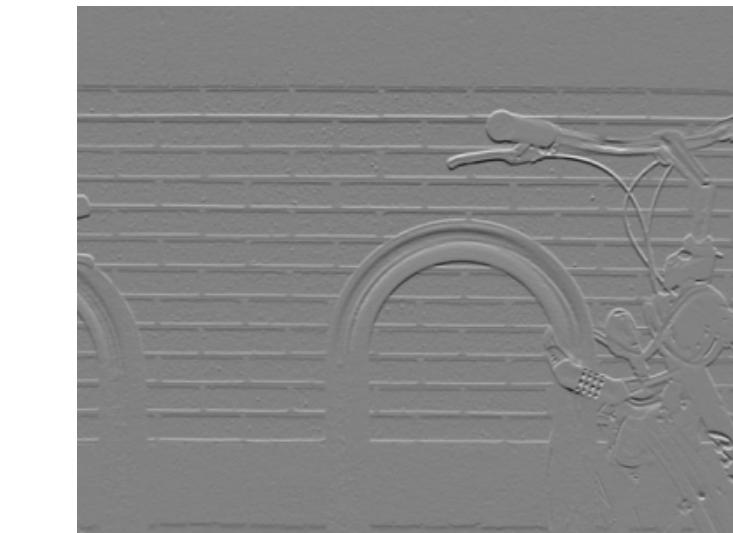
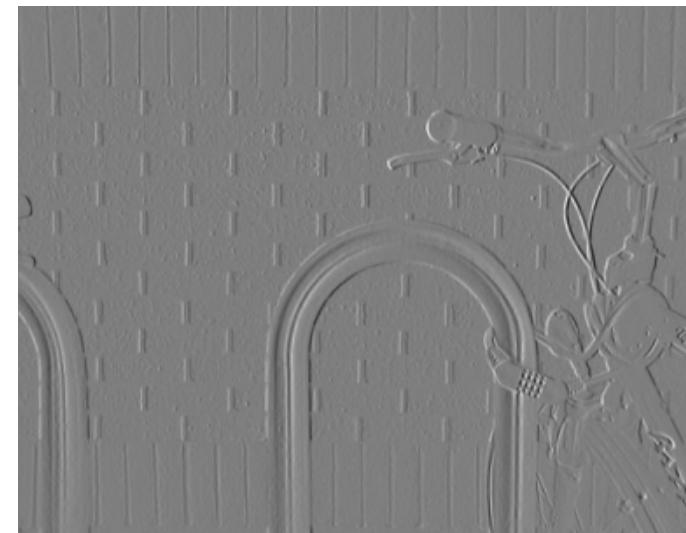
$$b_2 = [1 \ 1] \circ [1 \ 1] = [1 \ 2 \ 1]$$

$$b_3 = [1 \ 1] \circ [1 \ 1] \circ [1 \ 1] = [1 \ 3 \ 3 \ 1]$$

Binomial filter

b_1	1	1	1		$\sigma_1^2 = 1/4$					
b_2	1	2	1		$\sigma_2^2 = 1/2$					
b_3	1	3	3	1	$\sigma_3^2 = 3/4$					
b_4	1	4	6	4	1	$\sigma_4^2 = 1$				
b_5	1	5	10	10	5	1	$\sigma_5^2 = 5/4$			
b_6	1	6	15	20	15	6	1	$\sigma_6^2 = 3/2$		
b_7	1	7	21	35	35	21	7	1	$\sigma_7^2 = 7/4$	
b_8	1	8	28	56	70	56	28	8	1	$\sigma_8^2 = 2$

Sobel operator: example



Source: N. Snavely / Wikipedia

Next class: frequency